

# Chapter 3: Scalability Fundamentals

## Introduction: What is Scalability?

**Scalability** is the ability of a system to handle growing amounts of work by adding resources.

Think of it like a restaurant:

- **Day 1:** 20 customers → 1 chef, 2 tables ✓
- **Month 6:** 100 customers → Still 1 chef, 2 tables ✗ (chaos!)
- **Solution:** Scale the restaurant!

A scalable system maintains or improves performance as demand increases.

---

## 1. Vertical Scaling vs Horizontal Scaling

The two fundamental approaches to scaling any system.

### Vertical Scaling (Scale Up) - Making the Server Bigger

**Definition:** Adding more power to your existing machine (more CPU, RAM, storage)

**Analogy:** Making your car bigger and more powerful

Before Scaling:

Web Server	
4 CPU cores	Handles 1,000 requests/sec
8 GB RAM	
100 GB SSD	

After Vertical Scaling:

Web Server	
16 CPU cores	Handles 4,000 requests/sec
64 GB RAM	
1 TB SSD	

### Real-World Example: Database Server

#### Current Server:

- Handling 1,000 queries/second
- Response time: 50ms
- CPU at 80% usage
- Problem: Getting slow!

#### Vertical Scaling Solution:

Old: 8-core CPU, 32 GB RAM → \$200/month  
New: 32-core CPU, 128 GB RAM → \$800/month

#### Result:

- Handles 4,000 queries/second
- Response time: 20ms
- CPU at 40% usage

## Pros and Cons

### Advantages ✓

- **Simple:** Just upgrade, no code changes
- **No complexity:** Single machine, easier to manage
- **Data consistency:** Everything in one place
- **Fast communication:** No network latency between components

### Disadvantages ✗

- **Hardware limits:** Can't scale infinitely (physical limits)
- **Expensive:** High-end servers cost exponentially more
- **Single point of failure:** If server dies, entire system goes down
- **Downtime required:** Need to shut down for upgrades
- **Diminishing returns:**  $2x$  cost  $\neq$   $2x$  performance

### Cost Reality:

Server Specs	Monthly Cost	Performance Gain
4 cores, 16 GB RAM	\$100	Baseline
8 cores, 32 GB RAM	\$200	2x
16 cores, 64 GB RAM	\$400	3.5x (not 4x!)
32 cores, 128 GB RAM	\$800	6x (not 8x!)
64 cores, 256 GB RAM	\$2,000	9x (not 16x!)

Notice: **Cost grows faster than performance!**

---

## Horizontal Scaling (Scale Out) - Adding More Servers

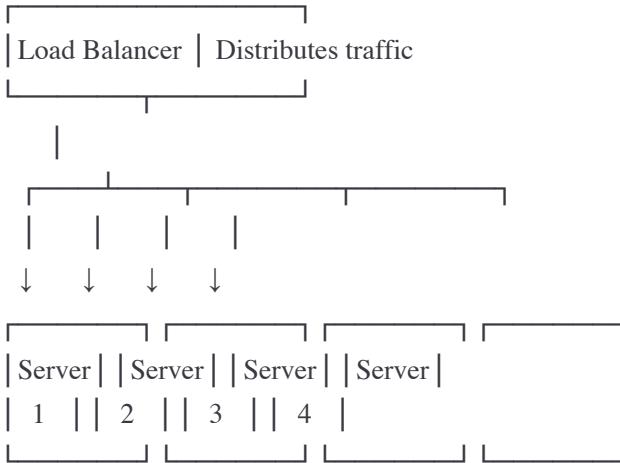
**Definition:** Adding more machines to handle the load

**Analogy:** Instead of one huge truck, use multiple smaller trucks

Before Scaling:

Web Server	Handles 1,000 requests/sec
4 CPU cores	
8 GB RAM	

After Horizontal Scaling:



Each handles 250 requests/sec

Total: 1,000 requests/sec

If load increases to 2,000:

Just add 4 more servers!

## Real-World Example: E-Commerce Website

## Black Friday Sale Scenario:

Normal Day:

- 10,000 concurrent users
- 2 web servers
- Works perfectly

Black Friday:

- 100,000 concurrent users (10x traffic!)
- Vertical scaling: Impossible to get 10x more powerful server
- Horizontal scaling: Add 18 more servers → 20 total
- Result: Handles load successfully!

## Pros and Cons

### Advantages ✓

- **Unlimited scaling:** Add as many servers as needed
- **Cost-effective:** Linear cost growth
- **Fault tolerance:** If one server fails, others continue
- **No downtime:** Add servers without stopping system
- **Geographic distribution:** Servers worldwide reduce latency

### Disadvantages ✗

- **Complexity:** Need load balancers, coordination
- **Data consistency:** Harder to keep data in sync
- **Code changes:** Application must support distributed architecture
- **Network overhead:** Communication between servers adds latency

## Side-by-Side Comparison

Aspect	Vertical Scaling	Horizontal Scaling
Implementation	Easy	Complex
Cost at scale	Very expensive	Linear growth
Maximum capacity	Hardware limited	Nearly unlimited
Single point fail	Yes	No (redundant)
Data consistency	Easy	Challenging
Setup time	Minutes	Hours/Days
Typical use case	Databases	Web servers, APIs

Example

Upgrade to 64-core    Add 10 more servers

---

## When to Use Each?

### Use Vertical Scaling When:

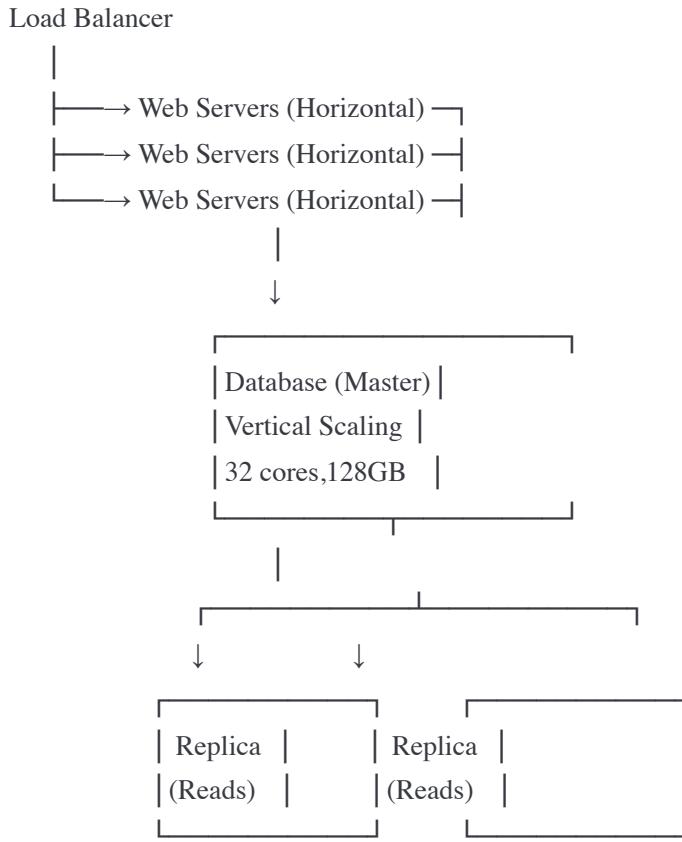
- ✓ Application can't be distributed (legacy monolith)
- ✓ Database that needs ACID transactions
- ✓ Quick fix needed (just upgrade hardware)
- ✓ Small to medium scale (< 10,000 users)
- ✓ Budget available for expensive hardware

### Use Horizontal Scaling When:

- ✓ Need to handle massive traffic
- ✓ Require high availability (no downtime)
- ✓ Want geographic distribution
- ✓ Cost-conscious growth
- ✓ Modern cloud-native application

### Best Practice: Hybrid Approach

Typical Modern Architecture:



Web layer: Horizontal (stateless, easy to scale)

DB layer: Vertical + replication (stateful, harder to scale)

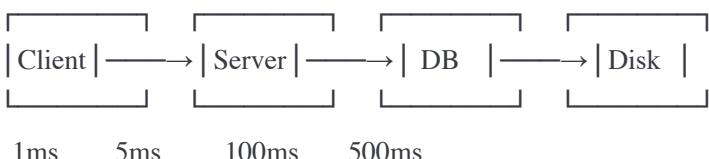
## 2. Understanding Bottlenecks

A **bottleneck** is the component that limits overall system performance.

**Analogy:** Traffic jam - doesn't matter how wide the highway is if there's a one-lane bridge in the middle!

### Identifying Bottlenecks

System Performance Chain:



Bottleneck: Disk! (slowest component)

Total response time: 606ms

Even if you make Client 10x faster → only saves 0.9ms!

But if you make Disk 2x faster → saves 250ms!

### Common Bottlenecks

## 1. CPU Bottleneck

### Symptoms:

- CPU usage consistently > 80%
- Slow response times for compute-heavy tasks
- Server gets hot/loud

### Example:

```
python

# CPU-intensive task
import time

def cpu_bottleneck_example():
    start = time.time()

    # Heavy computation
    total = 0
    for i in range(10_000_000):
        total += i ** 2

    elapsed = time.time() - start
    print(f"Took {elapsed:.2f} seconds")
    # Output: Took 3.45 seconds

# CPU at 100% during this time!

# Solutions:
# 1. Optimize algorithm
# 2. Add caching for repeated calculations
# 3. Add more CPU cores (vertical scaling)
# 4. Distribute work across multiple servers (horizontal)
```

### Real Example: Video Encoding Service

Scenario: Users upload videos to be encoded

Single Server (4 cores):

- Can encode 4 videos simultaneously
- Each video takes 10 minutes
- Throughput: 24 videos/hour
- Problem: 100 users uploading → 4+ hour wait!

Solution 1 - Vertical Scaling:

- Upgrade to 16 cores
- Throughput: 96 videos/hour
- Better, but still limited

Solution 2 - Horizontal Scaling:

- Add 10 servers (4 cores each)
- Throughput: 240 videos/hour
- Much better! And can add more as needed

## 2. Memory (RAM) Bottleneck

**Symptoms:**

- Frequent "out of memory" errors
- System using swap space (disk as RAM)
- Performance suddenly drops 100x
- Server becomes unresponsive

**Example:**

```
python
```

```

import sys

def memory_bottleneck_example():
    # Each user session stored in memory
    user_sessions = {}

    # Simulate 100,000 users
    for user_id in range(100_000):
        user_sessions[user_id] = {
            'data': 'x' * 50_000, # 50 KB per user
            'timestamp': time.time(),
            'preferences': [i for i in range(1000)]
        }

    # Total memory: 100,000 users × 50 KB = 5 GB
    print(f"Memory usage: {sys.getsizeof(user_sessions) / 1024 / 1024:.2f} MB")

    # If server only has 4 GB RAM → CRASH!

# Solutions:
# 1. Use Redis for session storage (separate server)
# 2. Implement session cleanup (delete old sessions)
# 3. Vertical scaling (more RAM)
# 4. Horizontal scaling (distribute users across servers)

```

## Real Example: Social Media Feed

Problem: Loading user's feed

Bad Approach (Memory bottleneck):

- Load all 10,000 posts into memory
- 10,000 posts × 2 KB each = 20 MB per user
- 1,000 concurrent users = 20 GB RAM needed!

Good Approach:

- Load only 20 posts at a time (pagination)
- 20 posts × 2 KB = 40 KB per user
- 1,000 concurrent users = 40 MB RAM needed!
- 500x less memory!

## 3. Database Bottleneck

### Symptoms:

- Slow query response times
- Database CPU/memory maxed out
- Queries timing out
- Connection pool exhausted

### Example Scenario:

sql

```
-- Slow query (bottleneck)
SELECT * FROM users
WHERE email LIKE '%@gmail.com%'
ORDER BY created_at DESC;
```

-- Scans entire table (millions of rows)

-- Takes 5+ seconds!

-- Performance:

Query time: 5,000ms

Throughput: 0.2 queries/second

With 100 concurrent users → Queue backs up!

-- Solution: Add index

```
CREATE INDEX idx_users_email ON users(email);
```

```
CREATE INDEX idx_users_created ON users(created_at);
```

-- Now query takes 10ms

-- Throughput: 100 queries/second

-- 500x improvement!

### Common Database Bottleneck Patterns:

## 1. Missing Indexes

Problem: Full table scans

Solution: Add appropriate indexes

## 2. N+1 Query Problem

Problem: Making 1000 queries instead of 1

Bad:

```
posts = get_all_posts()      # 1 query
for post in posts:
    author = get_author(post.author_id) # 1000 queries!
```

Good:

```
posts = get_all_posts_with_authors() # 1 query with JOIN
```

## 3. Expensive Joins

Problem: Joining 5 large tables

Solution: Denormalize data, use caching

## 4. Too Many Connections

Problem: 10,000 clients trying to connect

Solution: Connection pooling, read replicas

## 4. Network Bottleneck

### Symptoms:

- High latency between services
- Timeouts
- Packet loss
- Bandwidth saturation

### Example:

## Scenario: Microservices Architecture

Service A (API) → Service B (User Service) → Service C (Database)

Each arrow adds latency:

- A → B: 50ms (network)
- B → C: 100ms (network)
- Total: 150ms just for network!

If user request needs 10 service calls:

$$150\text{ms} \times 10 = 1,500\text{ms} = 1.5 \text{ seconds!}$$

Solutions:

1. Reduce service calls (batch requests)
2. Use caching (avoid repeated calls)
3. Deploy services closer together
4. Use faster network (upgrade bandwidth)
5. Async processing (don't wait for all responses)

## Real Example: Image Loading

Problem: Website loads 50 images per page

Naive approach:

- Each image: 2 MB
- User bandwidth: 10 Mbps
- Time to load ONE image:  $(2 \text{ MB} \times 8) / 10 \text{ Mbps} = 1.6 \text{ seconds}$
- Time to load 50 images: 80 seconds!

Optimized approach:

1. Compress images: 2 MB → 200 KB (10x smaller)
2. Use CDN: Serve from nearby server (300ms → 30ms)
3. Lazy loading: Load only visible images first
4. Result: Page loads in 3 seconds instead of 80!

## 5. Disk I/O Bottleneck

### Symptoms:

- High disk queue length
- Slow read/write operations
- System feels sluggish

## Example:

```
python

import time

# Writing 100,000 records to disk
def disk_bottleneck_example():
    # Bad: Write one at a time
    start = time.time()
    with open('data.txt', 'w') as f:
        for i in range(100_000):
            f.write(f"Record {i}\n")
            f.flush() # Force write to disk each time

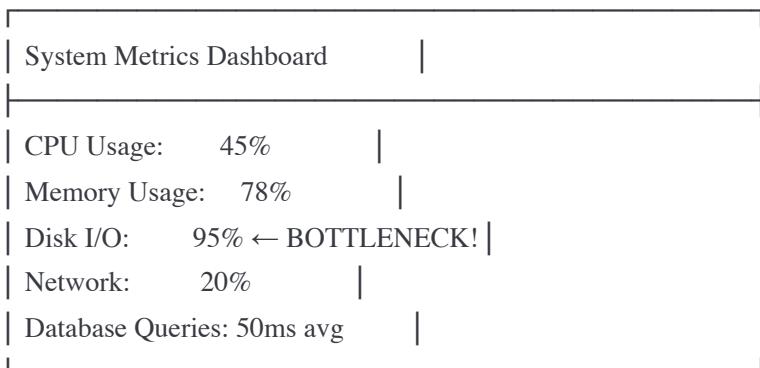
    elapsed = time.time() - start
    print(f"Unbuffered: {elapsed:.2f} seconds")
    # Output: 45 seconds (disk bottleneck!)

# Good: Batch writes (buffering)
start = time.time()
with open('data.txt', 'w') as f:
    for i in range(100_000):
        f.write(f"Record {i}\n")
    # Writes happen in batches automatically

elapsed = time.time() - start
print(f"Buffered: {elapsed:.2f} seconds")
# Output: 0.5 seconds (90x faster!)
```

## Bottleneck Identification Strategy

### Step 1: Monitor Everything



### Step 2: Trace a Request

User Request → Load Balancer (2ms)

```
→ Web Server (10ms)  
→ Cache Miss  
→ Database Query (500ms) ← BOTTLENECK!  
← Return Response
```

Total: 512ms

#### Step 3: Profile the Code

Function	Time	% of Total
----------	------	------------

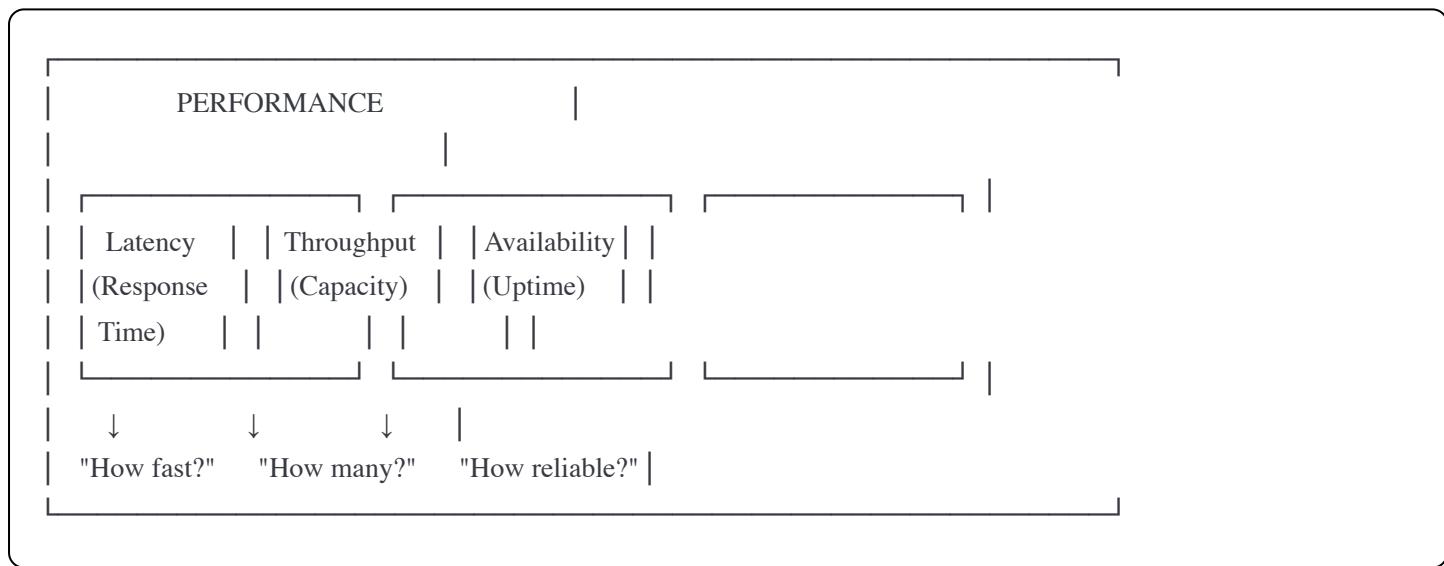
load_user_data()	450ms	88% ← BOTTLENECK!
render_template()	30ms	6%
send_response()	32ms	6%

#### Step 4: Optimize the Bottleneck

- Add caching: 450ms → 5ms
- Overall: 512ms → 67ms (7.6x faster!)

## 3. Performance Metrics

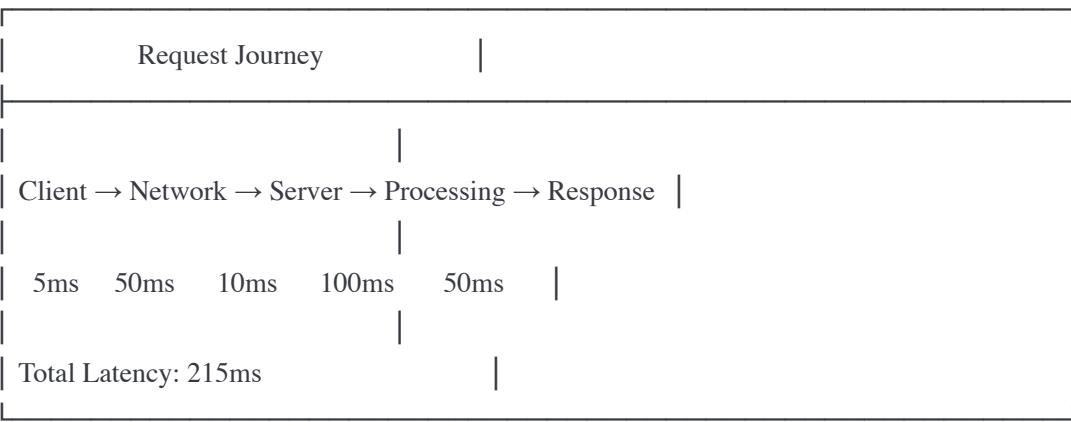
### The Three Pillars of System Performance



#### 1. Latency (Response Time)

**Definition:** Time from request to response

**Types of Latency:**



Breakdown:

- Network Latency:  $50\text{ms} + 50\text{ms} = 100\text{ms}$  (46%)
- Processing Latency:  $100\text{ms}$  (47%)
- Client Overhead:  $5\text{ms} + 10\text{ms} = 15\text{ms}$  (7%)

## Latency Percentiles (Critical Concept!):

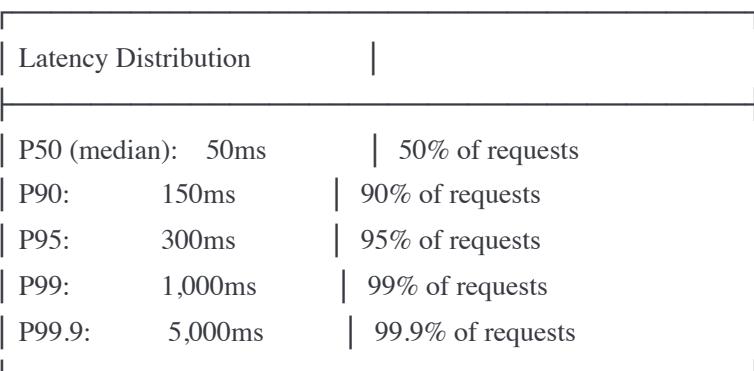
Why Percentiles Matter:

Average latency:  $100\text{ms}$  ← Misleading!

Could mean:

- Scenario A: All requests take  $100\text{ms}$  (good!)
- Scenario B: 99% take  $10\text{ms}$ , 1% take  $9\text{ seconds}$  (bad!)

Better Metrics:



Real Example:

If you have 1 million requests/day:

- P99.9 = 5s means 1,000 users experience 5+ second delays!

## Measuring Latency:

python

```

import time
import statistics

def measure_latency():
    latencies = []

    for i in range(1000):
        start = time.time()

        # Simulate operation (e.g., API call)
        response = make_api_call()

        latency = (time.time() - start) * 1000 # Convert to ms
        latencies.append(latency)

    # Calculate percentiles
    latencies.sort()

    print(f"Average latency: {statistics.mean(latencies):.2f}ms")
    print(f"P50 (median): {latencies[499]:.2f}ms")
    print(f"P90: {latencies[899]:.2f}ms")
    print(f"P95: {latencies[949]:.2f}ms")
    print(f"P99: {latencies[989]:.2f}ms")
    print(f"P99.9: {latencies[998]:.2f}ms")

# Example output:
# Average latency: 45.23ms
# P50 (median): 40.15ms
# P90: 75.30ms
# P95: 120.45ms
# P99: 450.80ms
# P99.9: 2,100.25ms

```

## Real-World Latency Targets:

Application Type	P50	P95	P99	User Experience
Search Engine	<50ms	<100ms	<200ms	Feels instant
Social Media Feed	<100ms	<300ms	<800ms	Smooth scrolling
E-commerce Checkout	<200ms	<500ms	<1s	Acceptable
Video Streaming	<2s	<5s	<10s	Tolerable
Batch Processing	N/A	N/A	N/A	Doesn't matter

## 2. Throughput (Capacity)

**Definition:** Number of operations completed per unit of time

### Measuring Throughput:

Requests per Second (RPS):

Time Period: 1 second	
Requests Completed: 1,000	
Throughput: 1,000 RPS	

Queries per Second (QPS):

Database	
Queries Completed: 5,000/sec	
Throughput: 5,000 QPS	

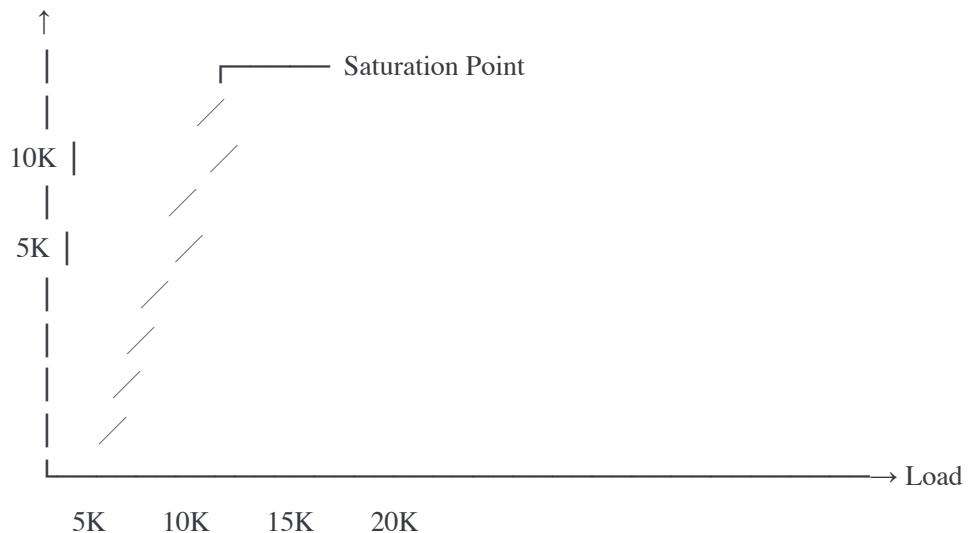
Bits per Second (bps):

Network	
Data Transferred: 100 Megabits/sec	
Throughput: 100 Mbps	

### Throughput vs Load:

## System Capacity Curve:

Throughput



Observations:

- Up to 10K: Throughput increases linearly
- 10K-15K: System approaching limits
- Beyond 15K: Throughput plateaus (saturated!)
- May even decrease due to overhead

## Calculating Maximum Throughput:

python

```

def calculate_max_throughput():
    """
    Calculate theoretical maximum throughput
    """

    # Server specs
    num_servers = 5
    cores_per_server = 8
    total_cores = num_servers * cores_per_server # 40 cores

    # Request handling
    avg_request_time_ms = 100 # 100ms per request
    requests_per_second_per_core = 1000 / avg_request_time_ms # 10 RPS

    # Maximum throughput
    max_throughput = total_cores * requests_per_second_per_core

    print(f"Theoretical Max Throughput: {max_throughput} RPS")
    # Output: 400 RPS

    # Real-world (accounting for overhead)
    efficiency = 0.7 # 70% efficiency (realistic)
    actual_max_throughput = max_throughput * efficiency

    print(f"Actual Max Throughput: {actual_max_throughput} RPS")
    # Output: 280 RPS

    return actual_max_throughput

```

## Load Testing Example:

```
python
```

```
import concurrent.futures
import requests
import time

def load_test(url, num_requests, concurrent_users):
    """
    Simulate load and measure throughput
    """

    start_time = time.time()
    successful_requests = 0
    failed_requests = 0

    def make_request(request_num):
        try:
            response = requests.get(url, timeout=5)
            return response.status_code == 200
        except:
            return False

    # Send concurrent requests
    with concurrent.futures.ThreadPoolExecutor(max_workers=concurrent_users) as executor:
        results = executor.map(make_request, range(num_requests))

        for success in results:
            if success:
                successful_requests += 1
            else:
                failed_requests += 1

    elapsed_time = time.time() - start_time
    throughput = successful_requests / elapsed_time

    print(f"Total Requests: {num_requests}")
    print(f"Successful: {successful_requests}")
    print(f"Failed: {failed_requests}")
    print(f"Time: {elapsed_time:.2f}s")
    print(f"Throughput: {throughput:.2f} RPS")
    print(f"Success Rate: {((successful_requests/num_requests)*100:.2f}%)"

# Example usage:
# load_test('http://localhost:8000/api/test', 1000, 50)

# Example output:
# Total Requests: 1000
# Successful: 980
```

```
# Failed: 20
# Time: 12.34s
# Throughput: 79.42 RPS
# Success Rate: 98.00%
```

### 3. Availability (Uptime)

**Definition:** Percentage of time a system is operational

#### The Nines:

Availability	Uptime/Year	Downtime/Year	Downtime/Month	Use Case
90%	328.5 days	36.5 days	3 days	Unacceptable
95%	346.6 days	18.3 days	1.5 days	Poor
99%	360.4 days	3.65 days	7.2 hours	Acceptable
99.9%	364.6 days	8.76 hours	43.2 minutes	Good
99.99%	365.1 days	52.6 minutes	4.32 minutes	Great
99.999%	365.2 days	5.26 minutes	26 seconds	Excellent
99.9999%	365.24 days	31.5 seconds	2.6 seconds	Extreme

Note: Each additional nine gets exponentially harder and more expensive!

#### Calculating Availability:

Formula:

$$\text{Availability} = (\text{Total Time} - \text{Downtime}) / \text{Total Time} \times 100\%$$

Example 1: Website down 1 hour in a month

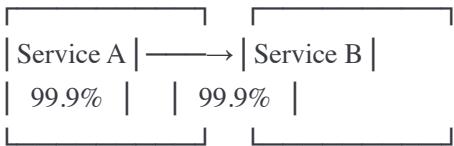
$$\begin{aligned}\text{Availability} &= (720 \text{ hours} - 1 \text{ hour}) / 720 \text{ hours} \times 100\% \\ &= 99.86\%\end{aligned}$$

Example 2: Service down 5 minutes in a year

$$\begin{aligned}\text{Availability} &= (525,600 \text{ min} - 5 \text{ min}) / 525,600 \text{ min} \times 100\% \\ &= 99.999\% \text{ (Five nines!)}\end{aligned}$$

#### Serial vs Parallel Components:

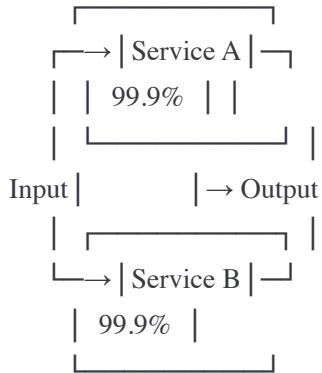
Serial System (both must work):



$$\text{Overall Availability} = 99.9\% \times 99.9\% = 99.8\%$$

(Worse than individual components!)

Parallel System (failover):

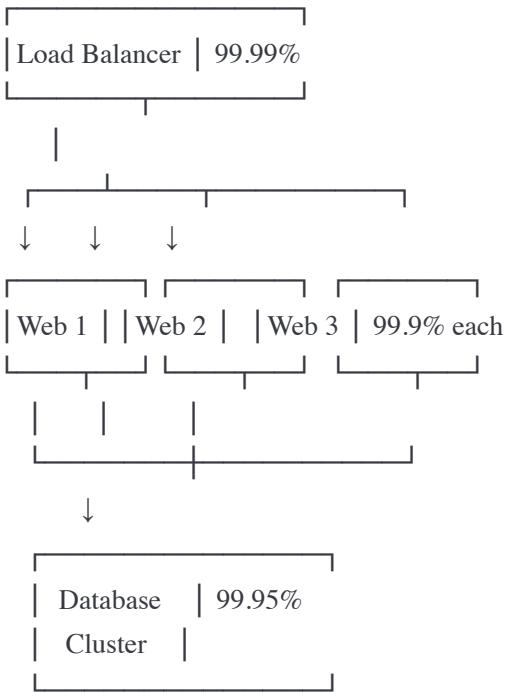


$$\text{Overall Availability} = 1 - (0.001 \times 0.001) = 99.9999\%$$

(Much better! Redundancy improves availability)

## Real-World Example: E-commerce Platform

Architecture:



Calculation:

- Web Servers (parallel):  $1 - (0.001)^3 = 99.9999\%$
- Load Balancer: 99.99%
- Database: 99.95%

$$\text{Overall} = 99.99\% \times 99.9999\% \times 99.95\% = 99.94\%$$

$$\text{Downtime per year} = 365 \text{ days} \times (1 - 0.9994)$$

$$\begin{aligned} &= 0.219 \text{ days} \\ &= 5.25 \text{ hours/year} \end{aligned}$$

## Performance Monitoring Dashboard Example

python

```
class PerformanceMonitor:
    def __init__(self):
        self.request_times = []
        self.successful_requests = 0
        self.failed_requests = 0
        self.start_time = time.time()

    def record_request(self, duration_ms, success=True):
        self.request_times.append(duration_ms)
        if success:
            self.successful_requests += 1
        else:
            self.failed_requests += 1

    def get_metrics(self):
        if not self.request_times:
            return None

        sorted_times = sorted(self.request_times)
        total_requests = len(self.request_times)
        elapsed_time = time.time() - self.start_time

        return {
            # Latency metrics
            'avg_latency': statistics.mean(sorted_times),
            'p50_latency': sorted_times[int(total_requests * 0.50)],
            'p95_latency': sorted_times[int(total_requests * 0.95)],
            'p99_latency': sorted_times[int(total_requests * 0.99)],

            # Throughput metrics
            'total_requests': total_requests,
            'throughput_rps': total_requests / elapsed_time,

            # Availability metrics
            'success_rate': (self.successful_requests / total_requests) * 100,
            'failure_rate': (self.failed_requests / total_requests) * 100,
        }

    def print_dashboard(self):
        metrics = self.get_metrics()
        if not metrics:
            print("No data yet")
            return

        print("\n" + "="*50)
        print("    PERFORMANCE DASHBOARD")
```

```

print("=*50)
print(f"\nLATENCY:")
print(f" Average: {metrics['avg_latency']:.2f}ms")
print(f" P50:    {metrics['p50_latency']:.2f}ms")
print(f" P95:    {metrics['p95_latency']:.2f}ms")
print(f" P99:    {metrics['p99_latency']:.2f}ms")
print(f"\nTHROUGHPUT:")
print(f" Total:  {metrics['total_requests']} requests")
print(f" RPS:    {metrics['throughput_rps']:.2f}")
print(f"\nAVAILABILITY:")
print(f" Success: {metrics['success_rate']:.2f}%")
print(f" Failure: {metrics['failure_rate']:.2f}%")
print("=*50 + "\n")

```

## 4. Back-of-the-Envelope Calculations

Essential skill for system design interviews and capacity planning!

### Key Numbers to Memorize

#### LATENCY NUMBERS (APPROXIMATE)

L1 cache reference	0.5 ns	
L2 cache reference	7 ns	
Main memory reference	100 ns	
Read 1 MB from memory	250 $\mu$ s	
SSD random read	150 $\mu$ s	
Read 1 MB from SSD	1 ms	
Disk seek	10 ms	
Read 1 MB from disk	20 ms	
Send packet CA→Netherlands→CA	150 ms	

#### DATA SIZE CONVERSIONS

1 Byte = 8 bits	
1 KB = 1,000 Bytes (or 1,024 in binary)	
1 MB = 1,000 KB	
1 GB = 1,000 MB	
1 TB = 1,000 GB	
1 PB = 1,000 TB	

## TIME CONVERSIONS

1 second = 1,000 milliseconds (ms)

1 ms = 1,000 microseconds ( $\mu$ s)

1  $\mu$ s = 1,000 nanoseconds (ns)

1 day = 86,400 seconds ( $\approx$  100,000 for quick calc)

1 month  $\approx$  2.5 million seconds

1 year  $\approx$  30 million seconds

## Calculation Examples

### Example 1: Twitter-like System

**Question:** Design Twitter. Estimate storage requirements.

## Step 1: Define Assumptions

---

- 300 million daily active users (DAU)
- Each user tweets 2 times per day on average
- 10% of tweets contain an image
- Average tweet size: 300 bytes (text)
- Average image size: 2 MB

## Step 2: Calculate Daily Data

---

Tweets per day:

$$300M \text{ users} \times 2 \text{ tweets} = 600M \text{ tweets/day}$$

Text storage:

$$600M \text{ tweets} \times 300 \text{ bytes} = 180 \text{ GB/day}$$

Images:

$$600M \text{ tweets} \times 10\% = 60M \text{ images/day}$$

$$60M \text{ images} \times 2 \text{ MB} = 120 \text{ TB/day}$$

Total daily storage:

$$180 \text{ GB} + 120 \text{ TB} \approx 120 \text{ TB/day}$$

## Step 3: Calculate 5-Year Storage

---

$$5 \text{ years} = 5 \times 365 \text{ days} = 1,825 \text{ days}$$

$$\text{Total storage} = 120 \text{ TB/day} \times 1,825 \text{ days}$$

$$= 219,000 \text{ TB}$$

$$= 219 \text{ PB}$$

With replication (3 copies for reliability):

$$219 \text{ PB} \times 3 = 657 \text{ PB}$$

## Step 4: Bandwidth Requirements

---

Upload bandwidth:

$$120 \text{ TB per day} / 86,400 \text{ seconds}$$

$$= 120 \times 10^{12} \text{ bytes} / 86,400 \text{ sec}$$

$$= 1.4 \times 10^9 \text{ bytes/sec}$$

$$= 11.1 \text{ Gbps}$$

Read:write ratio is typically 10:1 for social media

Download bandwidth:

$$11.1 \text{ Gbps} \times 10 = 111 \text{ Gbps}$$

## Step 5: Servers Needed

---

Assume each server can handle:

- 10,000 requests/second
- 1 Gbps network

Peak requests (assume 2x daily average):

$$600M \text{ tweets/day} / 86,400 \text{ sec} \times 2 = 14,000 \text{ requests/sec}$$

Servers for request handling:

$$14,000 / 10,000 = 2 \text{ servers (minimum)}$$

Add redundancy: 4-6 servers

Servers for bandwidth:

$$111 \text{ Gbps} / 1 \text{ Gbps} = 111 \text{ servers}$$

Therefore, we need ~120 servers for bandwidth alone!

---

## **Example 2: URL Shortener (like bit.ly)**

**Question:** How many unique short URLs can we generate?

## Step 1: Choose Character Set

---

Options:

- Numbers only (0-9): 10 characters
- Lowercase (a-z): 26 characters
- Alphanumeric (a-z, A-Z, 0-9): 62 characters

Let's use alphanumeric: 62 characters

## Step 2: Choose URL Length

---

Length	Combinations	Readable?
--------	--------------	-----------

---

Length	Combinations	Readable?
3	$62^3 = 238K$	Too short!
4	$62^4 = 14.7M$	Too short
5	$62^5 = 916M$	Getting there
6	$62^6 = 56.8B$	Good!
7	$62^7 = 3.5T$	Plenty!

## Step 3: Calculate Years of Service

---

Assumptions:

- 1 million new URLs per day
- 6-character URLs
- 56.8 billion possible combinations

$$\begin{aligned}\text{Years of service} &= 56.8B / (1M \text{ per day} \times 365 \text{ days}) \\ &= 56.8B / 365M \\ &= 155 \text{ years}\end{aligned}$$

Conclusion: 6 characters is sufficient!

## Step 4: Storage Requirements

---

Per URL record:

- Short code: 6 bytes
- Original URL: 200 bytes (average)
- Created date: 8 bytes
- User ID: 8 bytes
- Metadata: 50 bytes

Total: ~272 bytes per record

For 56.8 billion URLs:

$$56.8B \times 272 \text{ bytes} = 15.4 \text{ TB}$$

With indexing overhead (2x):

$15.4 \text{ TB} \times 2 = 30.8 \text{ TB}$

Very manageable!

---

### **Example 3: YouTube Video Storage**

**Question:** How much storage for 1 million videos?

## Step 1: Video Size Assumptions

---

Average video:

- Length: 10 minutes
- Quality: 1080p
- Size: ~100 MB per minute
- Total:  $10 \text{ min} \times 100 \text{ MB} = 1 \text{ GB per video}$

Multiple quality versions:

- 4K: 10 GB
- 1080p: 1 GB
- 720p: 0.5 GB
- 480p: 0.2 GB
- 360p: 0.1 GB

Total per video: 11.8 GB

## Step 2: Calculate for 1M Videos

---

Storage = 1M videos  $\times$  11.8 GB

$$= 11.8 \text{ million GB}$$

$$= 11.8 \text{ PB}$$

## Step 3: Add Thumbnails and Metadata

---

Thumbnails: 1M  $\times$  100 KB = 100 GB

Metadata: 1M  $\times$  10 KB = 10 GB

Total extra: ~110 GB (negligible compared to videos)

## Step 4: Factor in Redundancy

---

3 copies for reliability:

$$11.8 \text{ PB} \times 3 = 35.4 \text{ PB}$$

## Step 5: Daily Upload Rate

---

Assume 10,000 videos uploaded daily:

$$10,000 \text{ videos} \times 11.8 \text{ GB} = 118 \text{ TB/day}$$

Bandwidth needed:

$$118 \text{ TB} / 86,400 \text{ sec} = 1.4 \text{ GB/sec} = 11 \text{ Gbps upload}$$

## Step 6: Viewing Bandwidth

---

Assume:

- 100 million views per day
- Average video watched: 5 minutes

- Average quality: 720p (0.5 GB per 10 min)

Data served:

$$100M \text{ views} \times (5 \text{ min} / 10 \text{ min}) \times 0.5 \text{ GB}$$

$$= 100M \times 0.5 \times 0.5 \text{ GB}$$

$$= 25 \text{ million GB}$$

$$= 25 \text{ PB per day}$$

Bandwidth:

$$25 \text{ PB} / 86,400 \text{ sec} = 289 \text{ GB/sec} = 2.3 \text{ Tbps}$$

This is why YouTube uses CDNs worldwide!

#### Example 4: Real-Time Chat Application

**Question:** WhatsApp-like system for 1 billion users.

## Step 1: Active Users

---

Total users: 1 billion

Daily active users (DAU): 500 million (50%)

Peak concurrent users: 50 million (10% of DAU)

## Step 2: Message Volume

---

Average messages per user per day: 50

Total daily messages:  $500M \times 50 = 25$  billion messages

Messages per second (average):

$$25B / 86,400 \text{ sec} = 289,000 \text{ messages/sec}$$

Peak (2x average):

$$578,000 \text{ messages/sec}$$

## Step 3: Message Size

---

Text message: 100 bytes (average)

Image: 2 MB (10% of messages)

Video: 10 MB (1% of messages)

Average message size:

$$(0.89 \times 100 \text{ bytes}) + (0.10 \times 2 \text{ MB}) + (0.01 \times 10 \text{ MB})$$

$$= 89 \text{ bytes} + 0.2 \text{ MB} + 0.1 \text{ MB}$$

$$\approx 0.3 \text{ MB}$$

## Step 4: Daily Storage

---

$$25B \text{ messages} \times 0.3 \text{ MB} = 7.5 \text{ PB/day}$$

Annual storage:

$$7.5 \text{ PB} \times 365 = 2,737 \text{ PB} \approx 2.7 \text{ EB}$$

With 5-year retention:

$$2.7 \text{ EB} \times 5 = 13.5 \text{ EB}$$

## Step 5: Connection Requirements

---

Peak concurrent users: 50 million

Each maintains WebSocket connection

Assume each server handles 10,000 connections:

$$50M / 10,000 = 5,000 \text{ servers}$$

## Step 6: Bandwidth

---

Peak message rate: 578,000/sec

Average message: 0.3 MB

$$\text{Bandwidth} = 578,000 \times 0.3 \text{ MB} = 173 \text{ GB/sec} = 1.4 \text{ Tbps}$$

## Step 7: Database Queries

---

Each message requires:

- 1 write (sender)
- 1 read (recipient)
- Plus metadata updates

$$\text{Query rate: } 578,000 \times 3 = 1.7\text{M queries/sec}$$

If one DB handles 10,000 queries/sec:

$$1.7\text{M} / 10,000 = 170 \text{ database servers}$$

With read replicas (10:1 read:write ratio):

- Write DBs: 17 servers
- Read DBs: 153 servers

---

## Quick Estimation Framework

For ANY system design problem, calculate:

### 1. USERS

- Total users
- Daily active users (DAU)
- Peak concurrent users

### 2. TRAFFIC

- Requests/actions per day
- Requests per second (RPS)
- Peak RPS (typically 2-3x average)

### 3. STORAGE

- Data per item/record
- Daily data generated
- Total storage needed (5-10 years)
- With replication (3x)

### 4. BANDWIDTH

- Upload bandwidth

- Download bandwidth
- Read:write ratio

## 5. SERVERS

- Based on RPS capacity
- Based on storage needs
- Based on connection limits

## 6. MEMORY/CACHE

- Cache hit ratio assumption (80-90%)
- Hot data size (20% of total)
- RAM needed per server

# Key Takeaways

## 1. Scaling Strategies:

- Vertical: Quick but limited and expensive
- Horizontal: Complex but unlimited and cost-effective
- Use hybrid: Horizontal for stateless, vertical for stateful

## 2. Bottlenecks:

- Always monitor: CPU, Memory, Disk, Network, Database
- Optimize the slowest component first
- Measure before and after optimization

## 3. Performance Metrics:

- Latency: Use percentiles (P95, P99), not just averages
- Throughput: Plan for 2-3x peak load
- Availability: Each nine gets exponentially harder

## 4. Calculations:

- Memorize key numbers
- Break down into steps
- Round for simplicity
- State your assumptions clearly

# Practice Problems

**Problem 1:** Instagram-like app with 500M DAU, users upload 2 photos/day. Calculate: Daily storage, bandwidth, servers needed.

**Problem 2:** Design a cache for a system getting 10,000 RPS. Calculate: Cache size if 80% requests should hit cache.

**Problem 3:** Video conferencing app, 1M concurrent calls. Calculate: Bandwidth requirements.

---

## Next Chapter Preview

In Chapter 4, we'll dive into **Load Balancing**:

- How load balancers distribute traffic
- Different algorithms (Round Robin, Least Connections, etc.)
- Layer 4 vs Layer 7 load balancing
- Health checks and failover

Ready to continue? Let me know!