

Chapter 7: Database Scaling

Introduction: Why Database Scaling Matters

When your application grows, the database becomes the bottleneck.

Growth Timeline:

Month 1: 1,000 users

- Single database handles everything fine ✓

Month 6: 100,000 users

- Database CPU at 80%
- Queries slowing down
- Backup takes hours

Year 1: 10,000,000 users

- Database overwhelmed!
- Writes queuing up
- Reads timing out
- Users experiencing errors
- Revenue lost!

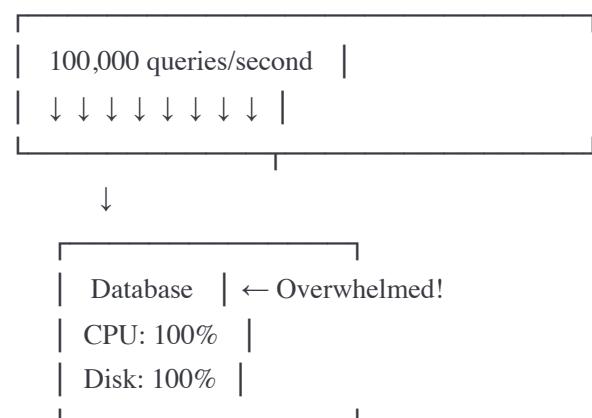
Solution: Database Scaling

1. Read Replicas and Replication Lag

The Read Scaling Problem

Single Database Problem:

All queries → One database



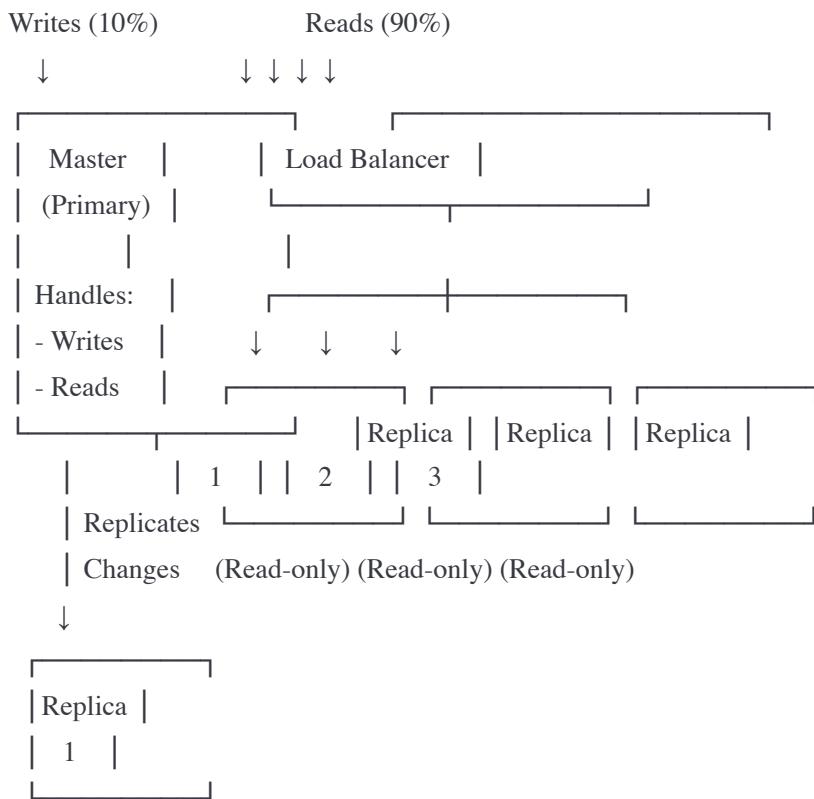
Symptoms:

- Slow queries (500ms → 5 seconds)
- Timeouts
- Connection pool exhausted
- System crashes

Read Replicas Solution

Concept: Copy data to multiple read-only databases.

With Read Replicas:



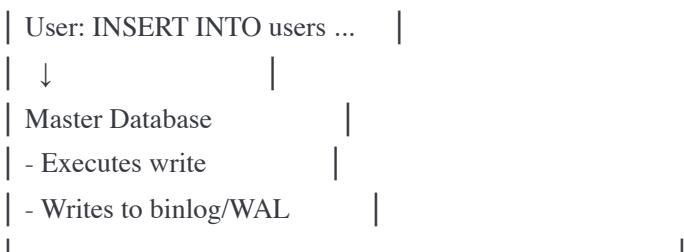
Benefits:

- Master: 10,000 writes/sec
- Each replica: 30,000 reads/sec
- Total read capacity: 90,000 reads/sec
- 9x read scaling!

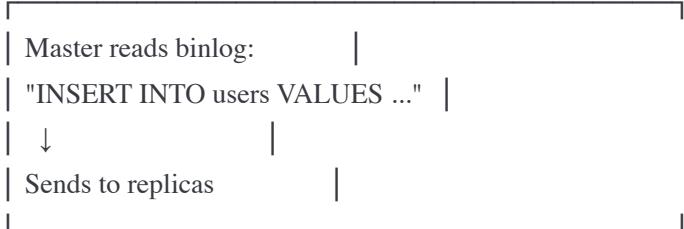
How Replication Works

Step-by-Step:

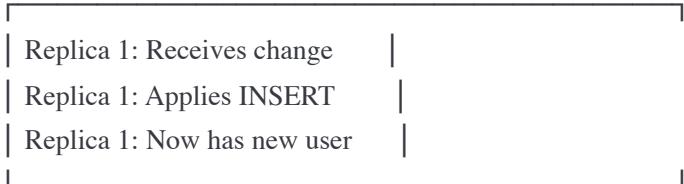
1. Write to Master



2. Master Ships Changes



3. Replicas Apply Changes



Timeline:

T=0ms: Write to master (complete)

T=5ms: Change shipped to replica

T=10ms: Replica applies change

Replication Lag

The Problem: Replicas are slightly behind master.

Replication Lag Example:

Time	Master	Replica
10:00	User creates post POST #123 created	(no data yet)
10:00	Response: "Post created!" User redirected to view post	
10:00	App reads from replica	← POST NOT THERE YET! SELECT * FROM posts WHERE id=123 Result: Not found ✗
10:00.05	Replication completes	POST #123 appears (50ms lag)
		Result: User sees "Post not found" error! This is called "read-your-own-writes" problem.

Measuring Replication Lag:

sql

```
-- PostgreSQL: Check replication lag
SELECT
    client_addr,
    state,
    sent_lsn,
    write_lsn,
    flush_lsn,
    replay_lsn,
    sync_state,
    -- Calculate lag in bytes
    pg_wal_lsn_diff(sent_lsn, replay_lsn) AS lag_bytes,
    -- Calculate lag in seconds
    EXTRACT(EPOCH FROM (now() - pg_last_xact_replay_timestamp())) AS lag_seconds
FROM pg_stat_replication;
```

-- Example output:

client_addr	state	lag_bytes	lag_seconds
10.0.1.10	streaming	4096	0.05
10.0.1.11	streaming	16384	0.15
10.0.1.12	streaming	32768	0.30

Interpretation:

- Replica 1: 50ms lag (excellent)
- Replica 2: 150ms lag (good)
- Replica 3: 300ms lag (acceptable)

Alert if lag > 1 second!

Handling Replication Lag

Strategy 1: Read from Master After Write

python

```
class DatabaseRouter:
    def __init__(self):
        self.master = connect_to_master()
        self.replicas = [
            connect_to_replica('replica1'),
            connect_to_replica('replica2'),
            connect_to_replica('replica3')
        ]
        self.recent_writes = {} # Track recent writes per user

    def write(self, user_id, query):
        """Execute write on master"""
        result = self.master.execute(query)

        # Remember this user just wrote
        self.recent_writes[user_id] = time.time()

        return result

    def read(self, user_id, query):
        """Smart read routing"""
        # Check if user wrote recently (last 1 second)
        if user_id in self.recent_writes:
            last_write = self.recent_writes[user_id]
            if time.time() - last_write < 1.0:
                # Read from master (has latest data)
                print(f'Reading from master for user {user_id} (recent write)')
                return self.master.execute(query)

        # Normal case: read from replica
        replica = random.choice(self.replicas)
        print(f'Reading from replica for user {user_id}')
        return replica.execute(query)

    # Usage
    router = DatabaseRouter()

    # User creates post
    router.write(user_id=123, query="INSERT INTO posts ...")

    # User immediately views their post
    # Reads from master (avoids replication lag!)
    post = router.read(user_id=123, query="SELECT * FROM posts WHERE id=456")

    # Another user views the post
```

```
# Can read from replica (doesn't need immediate consistency)
post = router.read(user_id=789, query="SELECT * FROM posts WHERE id=456")
```

Strategy 2: Monotonic Reads (Session Consistency)

```
python

class SessionConsistentRouter:
    """Ensure user always reads from same replica (or later)"""

    def __init__(self):
        self.master = connect_to_master()
        self.replicas = [
            connect_to_replica('replica1'),
            connect_to_replica('replica2'),
            connect_to_replica('replica3')
        ]
        self.user_replica_map = {} # Pin users to replicas

    def get_replica_for_user(self, user_id):
        """Pin user to a specific replica"""
        if user_id not in self.user_replica_map:
            # Assign user to replica based on hash
            replica_idx = hash(user_id) % len(self.replicas)
            self.user_replica_map[user_id] = replica_idx

        return self.replicas[self.user_replica_map[user_id]]

    def read(self, user_id, query):
        """Always read from same replica for this user"""
        replica = self.get_replica_for_user(user_id)
        return replica.execute(query)

# Benefit: User never sees data "go backwards"
# - First read: sees posts 1, 2, 3
# - Second read: will see 1, 2, 3, 4, 5 (never less!)
```

Strategy 3: Causal Consistency (Version Tracking)

```
python
```

```

class CausalConsistencyRouter:
    """Track data versions to ensure causality"""

    def write(self, query):
        """Write returns version number"""
        result = self.master.execute(query)
        version = self.master.get_current_log_position()
        return result, version

    def read(self, query, required_version=None):
        """Read from replica that has required version"""
        if required_version is None:
            # No version requirement, any replica is fine
            return random.choice(self.replicas).execute(query)

        # Find replica that has caught up to required version
        for replica in self.replicas:
            replica_version = replica.get_current_log_position()
            if replica_version >= required_version:
                return replica.execute(query)

        # No replica caught up yet, read from master
        return self.master.execute(query)

    # Usage
    router = CausalConsistencyRouter()

    # User A creates post
    result, version = router.write("INSERT INTO posts ...")

    # User A shares link with User B
    # Include version in URL: /posts/123?v=1234567

    # User B views post
    # Ensures they see the post (waits for replication if needed)
    post = router.read("SELECT * FROM posts WHERE id=123", required_version=version)

```

Replication Lag Trade-offs

Approach	Consistency	Performance
Always read master (no replicas used)	Perfect	Poor (1x)

Read from master after write (1s)	Good	Good (5x)	
Session consistency (sticky replicas)	Acceptable	Great (10x)	
Eventual consistency (any replica)	Weak	Excellent(10x)	

Choose based on use case:

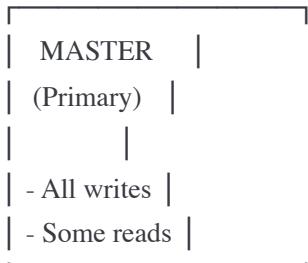
- Bank transactions: Read from master
- Social media: Eventual consistency fine
- User's own data: Read from master after write

2. Master-Slave vs Master-Master Replication

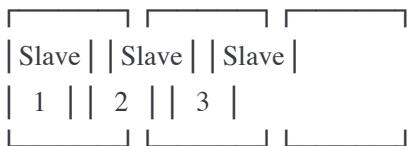
Master-Slave (Primary-Replica)

Architecture: One master, multiple read-only replicas.

Writes



Replicates



(Read) (Read) (Read)

Characteristics:

- Single source of truth (master)
- Simple to reason about
- No write conflicts
- Scales reads only

Configuration Example (MySQL):

sql

```

-- Master Configuration (my.cnf)
[mysqld]
server-id = 1
log-bin = /var/log/mysql/mysql-bin.log
binlog-format = ROW
binlog-do-db = myapp_db

-- Create replication user
CREATE USER 'replicator'@'%' IDENTIFIED BY 'password';
GRANT REPLICATION SLAVE ON *.* TO 'replicator'@'%';
FLUSH PRIVILEGES;

-- Get master status
SHOW MASTER STATUS;
-- Note: File and Position values needed for replica setup

-- Replica Configuration (my.cnf)
[mysqld]
server-id = 2
relay-log = /var/log/mysql/mysql-relay-bin
read-only = 1 -- Prevent writes to replica

-- Set up replication
CHANGE MASTER TO
MASTER_HOST='master-db.example.com',
MASTER_USER='replicator',
MASTER_PASSWORD='password',
MASTER_LOG_FILE='mysql-bin.000001',
MASTER_LOG_POS=12345;

START SLAVE;

-- Check replication status
SHOW SLAVE STATUS\G
-- Look for: Slave_IO_Running: Yes
--           Slave_SQL_Running: Yes
--           Seconds_Behind_Master: 0

```

Application Code:

```
python
```

```
import pymysql
from random import choice

class MasterSlaveDB:
    def __init__(self):
        self.master = pymysql.connect(
            host='master-db.example.com',
            user='app',
            password='password',
            database='myapp_db'
        )

        self.slaves = [
            pymysql.connect(host='slave1-db.example.com', ...),
            pymysql.connect(host='slave2-db.example.com', ...),
            pymysql.connect(host='slave3-db.example.com', ...)
        ]

    def execute_write(self, query, params):
        """All writes go to master"""
        cursor = self.master.cursor()
        cursor.execute(query, params)
        self.master.commit()
        return cursor.lastrowid

    def execute_read(self, query, params):
        """Reads go to random slave"""
        slave = choice(self.slaves)
        cursor = slave.cursor()
        cursor.execute(query, params)
        return cursor.fetchall()

# Usage
db = MasterSlaveDB()

# Write to master
user_id = db.execute_write(
    "INSERT INTO users (name, email) VALUES (%s, %s)",
    ("John", "john@example.com")
)

# Read from slave
users = db.execute_read(
    "SELECT * FROM users WHERE active = %s",

```

(True,)

)

Pros and Cons:

✓ ADVANTAGES:

- Simple to set up and manage
- No write conflicts
- Clear data flow
- Easy to reason about consistency
- Most common architecture

✗ DISADVANTAGES:

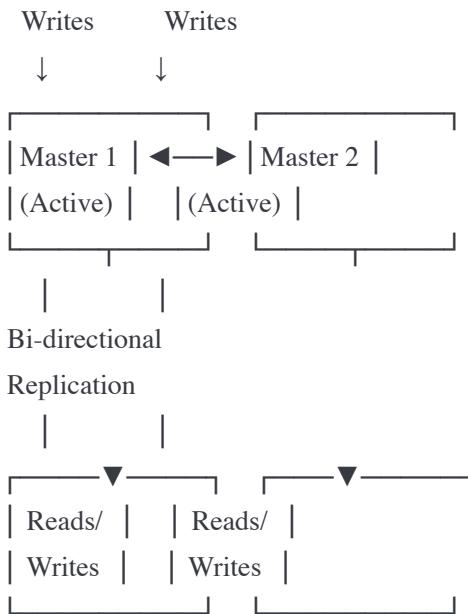
- Single point of failure (master)
- Master can become write bottleneck
- Doesn't scale writes
- Failover requires manual/automated promotion

When to use:

- Most applications (default choice)
- Read-heavy workloads (90% reads, 10% writes)
- When strong consistency for writes is needed

Master-Master (Multi-Master)

Architecture: Multiple masters, all accept writes.



Characteristics:

- Both accept writes
- Replicate to each other
- Complex conflict resolution
- Scales writes (to a point)
- High availability

Conflict Example:

The Write Conflict Problem:

Time Master 1 Master 2

10:00 User balance: \$1000 User balance: \$1000

10:00 User withdraws \$500

Balance = \$500

10:00 User withdraws \$600

Balance = \$400

10:01 Replication arrives Replication arrives

Sees Master 2's update Sees Master 1's update

Which is correct? Which is correct?

Options:

1. Last Write Wins (LWW): Use timestamp

- Problem: Both think they won!

2. Version Vectors: Track all changes

- Detect conflict, require manual resolution

3. Application Logic: Custom resolution

- Example: Sum all withdrawals, check if valid

Conflict Resolution Strategies:

sql

-- Strategy 1: Last Write Wins (LWW)

-- Use timestamp to decide winner

```
CREATE TABLE accounts (
    id INTEGER PRIMARY KEY,
    balance DECIMAL(10, 2),
    updated_at TIMESTAMP,
    updated_by VARCHAR(50) -- Which master made the change
);
```

-- On conflict, keep the one with latest timestamp

-- Problem: Can lose data!

-- Strategy 2: Version Vectors

```
CREATE TABLE accounts (
    id INTEGER PRIMARY KEY,
    balance DECIMAL(10, 2),
    version_m1 INTEGER, -- Version from master 1
    version_m2 INTEGER -- Version from master 2
);
```

-- Master 1 write: version_m1++

-- Master 2 write: version_m2++

-- On replication, check if versions compatible

-- If conflict detected, flag for manual review

-- Strategy 3: CRDTs (Conflict-Free Replicated Data Types)

-- Use data structures that merge automatically

-- Example: Counter that only increments (never decreases)

-- Both masters can increment, sum is always correct

```
CREATE TABLE page_views (
    page_id INTEGER PRIMARY KEY,
    count_m1 INTEGER, -- Increments from master 1
    count_m2 INTEGER -- Increments from master 2
);
```

-- Total views = count_m1 + count_m2

-- No conflicts! Both can increment freely

Configuration Example (PostgreSQL BDR):

```
sql
```

```

-- Master 1 Configuration
CREATE EXTENSION pglogical;

-- Create node
SELECT pglogical.create_node(
    node_name := 'master1',
    dsn := 'host=master1.example.com dbname=myapp'
);

-- Create replication set
SELECT pglogical.create_replication_set('default');

-- Add all tables to replication
SELECT pglogical.replication_set_add_all_tables(
    'default',
    ARRAY['public']
);

-- Master 2 Configuration
CREATE EXTENSION pglogical;

SELECT pglogical.create_node(
    node_name := 'master2',
    dsn := 'host=master2.example.com dbname=myapp'
);

-- Subscribe to Master 1
SELECT pglogical.create_subscription(
    subscription_name := 'sub_from_master1',
    provider_dsn := 'host=master1.example.com dbname=myapp',
    replication_sets := ARRAY['default']
);

-- On Master 1, subscribe to Master 2
-- (Creates bi-directional replication)
SELECT pglogical.create_subscription(
    subscription_name := 'sub_from_master2',
    provider_dsn := 'host=master2.example.com dbname=myapp',
    replication_sets := ARRAY['default']
);

```

Pros and Cons:

✓ ADVANTAGES:

- High availability (both can handle writes)
- Scales writes across multiple masters
- Geographic distribution possible
- No single point of failure

✗ DISADVANTAGES:

- Complex conflict resolution
- Increased latency (cross-master replication)
- Can lose data on conflicts
- Harder to reason about
- More operational complexity

When to use:

- Multi-region deployments (master per region)
- High write throughput needed
- Can tolerate eventual consistency
- Application can handle conflicts

Failover Strategies

Automatic Failover (Master-Slave)

python

```
import time
import threading

class AutoFailover:
    def __init__(self):
        self.master = "master1.db.com"
        self.replicas = ["replica1.db.com", "replica2.db.com"]
        self.current_master = self.master

        # Start health check thread
        self.monitoring = True
        self.health_thread = threading.Thread(target=self._monitor_health)
        self.health_thread.start()

    def _monitor_health(self):
        """Continuously check master health"""
        failures = 0

        while self.monitoring:
            if not self._check_health(self.current_master):
                failures += 1
                print(f"Master health check failed ({failures}/3)")

            if failures >= 3:
                # Master is down! Promote replica
                self._failover()
                failures = 0
            else:
                failures = 0

        time.sleep(5) # Check every 5 seconds

    def _check_health(self, host):
        """Check if database is responsive"""
        try:
            conn = connect(host)
            conn.execute("SELECT 1")
            return True
        except:
            return False

    def _failover(self):
        """Promote replica to master"""
        print("🔴 Master failed! Initiating failover...")

        # Choose most up-to-date replica
```

```

best_replica = self._choose_best_replica()

print(f"Promoting {best_replica} to master")

# 1. Stop replication on chosen replica
self._stop_replication(best_replica)

# 2. Make it writable
self._make_writable(best_replica)

# 3. Point other replicas to new master
for replica in self.replicas:
    if replica != best_replica:
        self._repaint_replica(replica, best_replica)

# 4. Update application connection
self.current_master = best_replica

print(f"✓ Failover complete. New master: {best_replica}")

# 5. Alert operations team
self._send_alert(f"Database failover: {best_replica} is now master")

def _choose_best_replica(self):
    """Choose replica with least replication lag"""
    best_replica = None
    min_lag = float('inf')

    for replica in self.replicas:
        lag = self._get_replication_lag(replica)
        if lag < min_lag:
            min_lag = lag
            best_replica = replica

    return best_replica

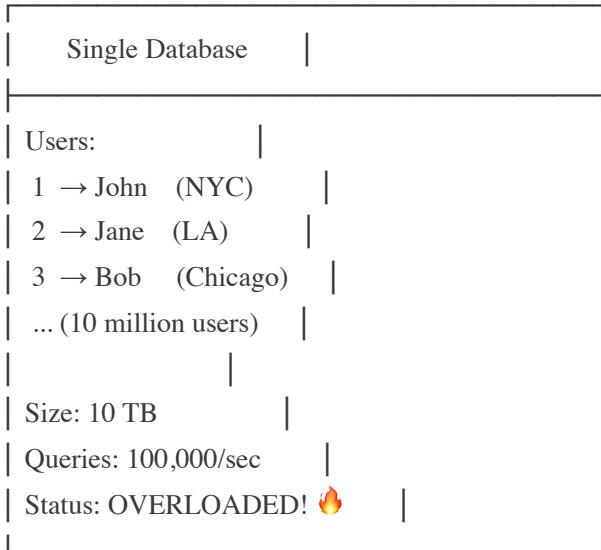
```

3. Database Sharding (Horizontal Partitioning)

What is Sharding?

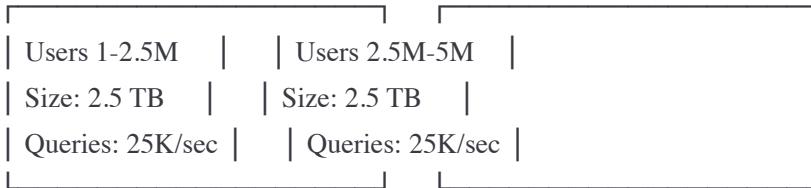
Sharding: Split database into multiple independent databases (shards).

Before Sharding (Single Database):

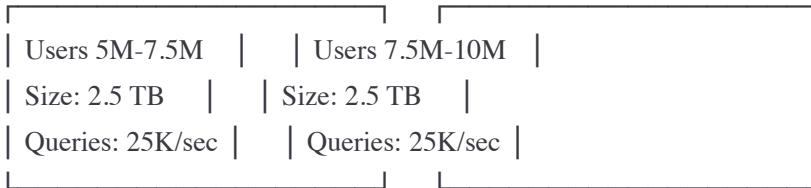


After Sharding (4 Shards):

Shard 1 (Users 0-25%) Shard 2 (Users 25-50%)



Shard 3 (Users 50-75%) Shard 4 (Users 75-100%)



Benefits:

- 4x smaller databases (manageable)
- 4x write throughput
- 4x read throughput
- Can add more shards as you grow

Why Shard?

Problem: Instagram's Growth

2010: 1 million users

- Single database: Works fine ✓

2012: 100 million users

- Single database: Struggling
- Queries slow (500ms)
- Backups take 12 hours
- Failover takes 1 hour

2015: 500 million users

- Single database: Impossible!
- Database size: 50+ TB
- Can't fit on one machine
- Can't handle write load

Solution: Shard by user_id

- Shard 1: Users 0-100M
- Shard 2: Users 100-200M
- Shard 3: Users 200-300M
- etc.

Result:

- Each shard: manageable size
- Linear scalability
- Can support billions of users

Sharding Challenges

1. CROSS-SHARD QUERIES

Problem: User A (Shard 1) follows User B (Shard 2)

Query: "Get all posts from people I follow"

Solution: Query both shards, merge results

2. CROSS-SHARD JOINS

Problem: Can't JOIN across shards

Solution: Denormalize data or query multiple shards

3. DISTRIBUTED TRANSACTIONS

Problem: Transfer money between users on different shards

Solution: Two-phase commit or eventual consistency

4. RESHARDING

Problem: Adding new shards requires moving data

Solution: Consistent hashing (covered later)

5. HOTSPOTS

Problem: One shard gets more traffic than others

Solution: Better shard key selection

4. Sharding Strategies

Strategy 1: Hash-Based Sharding

Concept: Hash the shard key to determine shard.

Algorithm:

```
shard_number = hash(user_id) % number_of_shards
```

Example: 4 shards

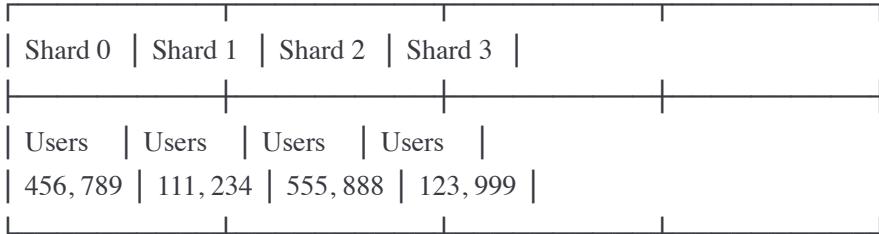
user_id = 123

shard = hash(123) % 4 = 3

user_id = 456

shard = hash(456) % 4 = 0

Distribution:



Hash function distributes users evenly

Implementation:

```
python
```

```
import hashlib

class HashBasedSharding:
    def __init__(self, shard_configs):
        """
        shard_configs: [
            {'host': 'shard0.db.com', 'db': 'shard_0'},
            {'host': 'shard1.db.com', 'db': 'shard_1'},
            {'host': 'shard2.db.com', 'db': 'shard_2'},
            {'host': 'shard3.db.com', 'db': 'shard_3'}
        ]
        """

        self.shards = [
            connect_to_db(config) for config in shard_configs
        ]
        self.num_shards = len(self.shards)

    def get_shard(self, user_id):
        """Determine which shard contains this user"""
        # Hash the user_id
        hash_value = int(hashlib.md5(str(user_id).encode()).hexdigest(), 16)

        # Modulo to get shard number
        shard_num = hash_value % self.num_shards

        return self.shards[shard_num]

    def insert_user(self, user_id, name, email):
        """Insert user into appropriate shard"""
        shard = self.get_shard(user_id)

        shard.execute(
            "INSERT INTO users (id, name, email) VALUES (%s, %s, %s)",
            (user_id, name, email)
        )

        print(f"Inserted user {user_id} into shard {shard_num}")

    def get_user(self, user_id):
        """Retrieve user from appropriate shard"""
        shard = self.get_shard(user_id)

        result = shard.execute(
            "SELECT * FROM users WHERE id = %s",
            (user_id,)
        )
```

```
return result.fetchone()

def get_user_posts(self, user_id):
    """Get posts for a user (same shard)"""
    shard = self.get_shard(user_id)

    result = shard.execute(
        "SELECT * FROM posts WHERE user_id = %s",
        (user_id,),
    )

    return result.fetchall()

def get_feed(self, user_id, following_ids):
    """Get feed from multiple users (cross-shard query!)"""

    # Group users by shard
    shard_to_users = {}

    for followed_id in following_ids:
        shard_num = self.get_shard_number(followed_id)
        if shard_num not in shard_to_users:
            shard_to_users[shard_num] = []
        shard_to_users[shard_num].append(followed_id)

    # Query each shard
    all_posts = []
    for shard_num, user_ids in shard_to_users.items():
        shard = self.shards[shard_num]

        # Get posts from all users on this shard
        placeholders = ','.join(['%s'] * len(user_ids))
        result = shard.execute(
            f"SELECT * FROM posts WHERE user_id IN ({placeholders})",
            user_ids
        )

        all_posts.extend(result.fetchall())

    # Merge and sort
    all_posts.sort(key=lambda p: p['created_at'], reverse=True)

    return all_posts[:20] # Top 20 posts

# Usage
sharding = HashBasedSharding([
    {'host': 'shard0.db.com'},
    {'host': 'shard1.db.com'},
```

```

{'host': 'shard2.db.com'},
{'host': 'shard3.db.com'}
)

# Insert user
sharding.insert_user(123, "John", "john@example.com")

# Get user (goes to correct shard automatically)
user = sharding.get_user(123)

# Get feed (queries multiple shards)
feed = sharding.get_feed(
    user_id=123,
    following_ids=[111, 222, 333, 444, 555]
)

```

Pros and Cons:

✓ ADVANTAGES:

- Even distribution (if good hash function)
- Simple to implement
- No hotspots (assuming uniform access pattern)
- Deterministic (same key always goes to same shard)

✗ DISADVANTAGES:

- Adding shards requires rehashing ALL data!
- Example: 4 shards → 5 shards
 - $\text{hash}(123) \% 4 = 3$
 - $\text{hash}(123) \% 5 = 3$ (lucky, same shard)
 - $\text{hash}(456) \% 4 = 0$
 - $\text{hash}(456) \% 5 = 1$ (different shard! Must move data)
- Approximately 80% of data needs to move!
- Can't do range queries

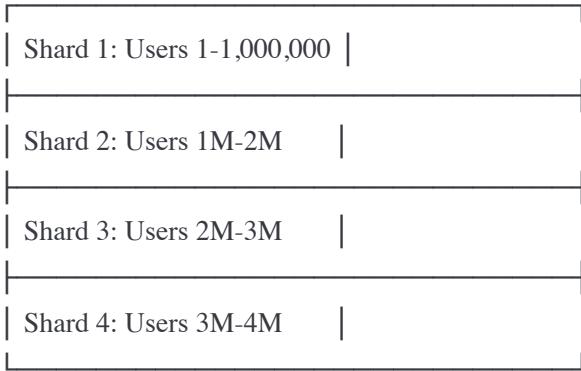
When to use:

- Stable number of shards
- Don't need range queries
- Access pattern is uniform

Strategy 2: Range-Based Sharding

Concept: Shard based on ranges of values.

Shard by user_id ranges:



Or shard by date:



Or shard by geography:



Implementation:

```
python
```

```

class RangeBasedSharding:
    def __init__(self):
        self.shard_map = [
            {'range': (0, 1_000_000), 'shard': connect_to_db('shard0.db.com')},
            {'range': (1_000_000, 2_000_000), 'shard': connect_to_db('shard1.db.com')},
            {'range': (2_000_000, 3_000_000), 'shard': connect_to_db('shard2.db.com')},
            {'range': (3_000_000, 4_000_000), 'shard': connect_to_db('shard3.db.com')}]
    ]

    def get_shard(self, user_id):
        """Find shard containing this user_id"""
        for mapping in self.shard_map:
            start, end = mapping['range']
            if start <= user_id < end:
                return mapping['shard']

        raise ValueError(f"No shard found for user_id {user_id}")

    def get_users_in_range(self, start_id, end_id):
        """Efficiently query range (may span multiple shards)"""
        results = []

        for mapping in self.shard_map:
            shard_start, shard_end = mapping['range']

            # Check if this shard overlaps with query range
            if shard_start < end_id and shard_end > start_id:
                # Calculate overlap
                query_start = max(start_id, shard_start)
                query_end = min(end_id, shard_end)

                # Query this shard
                shard = mapping['shard']
                result = shard.execute(
                    "SELECT * FROM users WHERE id >= %s AND id < %s",
                    (query_start, query_end)
                )

                results.extend(result.fetchall())

        return results

```

Example: Get users 1,500,000 to 2,500,000
Queries Shard 2 and Shard 3 only!

```
sharding = RangeBasedSharding()  
users = sharding.get_users_in_range(1_500_000, 2_500_000)
```

Time-Based Sharding Example:

```
python
```

```
from datetime import datetime, timedelta

class TimeBasedSharding:
    """Shard logs by date - useful for time-series data"""

    def __init__(self):
        self.shards = {
            '2024-Q1': connect_to_db('logs_2024_q1.db.com'),
            '2024-Q2': connect_to_db('logs_2024_q2.db.com'),
            '2024-Q3': connect_to_db('logs_2024_q3.db.com'),
            '2024-Q4': connect_to_db('logs_2024_q4.db.com')
        }

    def get_shard_for_date(self, date):
        """Determine shard based on date"""
        quarter = (date.month - 1) // 3 + 1
        shard_key = f'{date.year}-Q{quarter}'
        return self.shards[shard_key]

    def insert_log(self, timestamp, message):
        """Insert log into appropriate shard"""
        date = datetime.fromtimestamp(timestamp)
        shard = self.get_shard_for_date(date)

        shard.execute(
            "INSERT INTO logs (timestamp, message) VALUES (%s, %s)",
            (timestamp, message)
        )

    def query_logs(self, start_date, end_date):
        """Query logs across date range"""
        results = []

        # Determine which shards to query
        current = start_date
        while current <= end_date:
            shard = self.get_shard_for_date(current)

            result = shard.execute(
                "SELECT * FROM logs WHERE timestamp >= %s AND timestamp < %s",
                (current.timestamp(), (current + timedelta(days=90)).timestamp())
            )

            results.extend(result.fetchall())
            current += timedelta(days=90)
```

```
return results
```

```
# Benefit: Old shards can be archived/deleted
```

```
# Query only relevant shards for date range
```

Pros and Cons:

✓ ADVANTAGES:

- Range queries are efficient
- Easy to add new shards (just add new range)
- Can archive old shards (for time-based)
- Intuitive and easy to understand

✗ DISADVANTAGES:

- Uneven distribution (hotspots!)
- Example: New users (high IDs) more active than old
- Recent data gets more traffic (for time-based)
- May need to rebalance ranges

When to use:

- Time-series data (logs, events)
- Need range queries
- Data has natural ranges (geography, dates)
- Can tolerate some imbalance

Strategy 3: Directory-Based Sharding

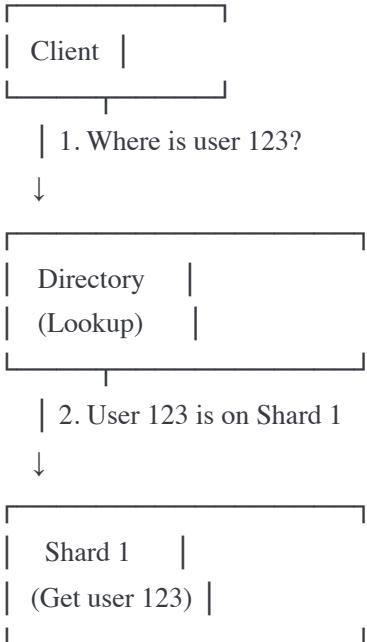
Concept: Maintain a lookup table (directory) that maps keys to shards.

Directory Table (Lookup Service):

User_ID	Shard
123	Shard 1
456	Shard 2
789	Shard 1
111	Shard 3
222	Shard 2

Process:

1. Query directory: "Where is user 123?"
2. Directory returns: "Shard 1"
3. Query Shard 1 for user data



Implementation:

```
python
```

```
class DirectoryBasedSharding:
    def __init__(self):
        # Directory: Maps user_id → shard_id
        self.directory = connect_to_db('directory.db.com')

        # Shards
        self.shards = {
            'shard_1': connect_to_db('shard1.db.com'),
            'shard_2': connect_to_db('shard2.db.com'),
            'shard_3': connect_to_db('shard3.db.com'),
            'shard_4': connect_to_db('shard4.db.com')
        }

    def get_shard_for_user(self, user_id):
        """Lookup which shard has this user"""
        result = self.directory.execute(
            "SELECT shard_id FROM user_shard_map WHERE user_id = %s",
            (user_id,)
        )

        row = result.fetchone()
        if not row:
            # User not found, assign to shard with most capacity
            shard_id = self._choose_shard_for_new_user()
            self._register_user(user_id, shard_id)
            return self.shards[shard_id]

        return self.shards[row['shard_id']]

    def _choose_shard_for_new_user(self):
        """Choose shard with most available capacity"""
        # Query each shard for size
        shard_sizes = {}
        for shard_id, shard in self.shards.items():
            result = shard.execute("SELECT COUNT(*) FROM users")
            shard_sizes[shard_id] = result.fetchone()[0]

        # Return shard with fewest users
        return min(shard_sizes, key=shard_sizes.get)

    def _register_user(self, user_id, shard_id):
        """Record user's shard in directory"""
        self.directory.execute(
            "INSERT INTO user_shard_map (user_id, shard_id) VALUES (%s, %s)",
            (user_id, shard_id)
        )
```

```
def insert_user(self, user_id, name, email):
    """Insert user"""
    shard = self.get_shard_for_user(user_id)

    shard.execute(
        "INSERT INTO users (id, name, email) VALUES (%s, %s, %s)",
        (user_id, name, email)
    )

def get_user(self, user_id):
    """Get user from appropriate shard"""
    shard = self.get_shard_for_user(user_id)

    result = shard.execute(
        "SELECT * FROM users WHERE id = %s",
        (user_id,)
    )

    return result.fetchone()

def move_user_to_shard(self, user_id, new_shard_id):
    """Rebalancing: Move user to different shard"""
    # Get user data from old shard
    old_shard = self.get_shard_for_user(user_id)
    user_data = old_shard.execute(
        "SELECT * FROM users WHERE id = %s",
        (user_id,)
    ).fetchone()

    # Insert into new shard
    new_shard = self.shards[new_shard_id]
    new_shard.execute(
        "INSERT INTO users (id, name, email) VALUES (%s, %s, %s)",
        (user_data['id'], user_data['name'], user_data['email'])
    )

    # Delete from old shard
    old_shard.execute("DELETE FROM users WHERE id = %s", (user_id,))

    # Update directory
    self.directory.execute(
        "UPDATE user_shard_map SET shard_id = %s WHERE user_id = %s",
        (new_shard_id, user_id)
    )

    print(f"Moved user {user_id} to {new_shard_id}")
```

```

# Usage
sharding = DirectoryBasedSharding()

# Insert user (automatically assigns to least-loaded shard)
sharding.insert_user(123, "John", "john@example.com")

# Later, can rebalance
sharding.move_user_to_shard(123, 'shard_2')

```

Pros and Cons:

✓ ADVANTAGES:

- Flexible (can place users anywhere)
- Easy to rebalance (just update directory)
- No rehashing when adding shards
- Can optimize placement (e.g., by geography, tenant size)
- Supports complex sharding logic

✗ DISADVANTAGES:

- Directory is single point of failure
- Directory can become bottleneck
- Extra lookup on every query (latency)
- Directory must be highly available

Solutions for directory scaling:

- Cache directory lookups
- Replicate directory
- Use distributed directory (e.g., ZooKeeper)

When to use:

- Need flexibility in shard assignment
- Plan to rebalance frequently
- Can tolerate extra lookup latency
- Complex sharding requirements

5. Consistent Hashing

Solves the "adding shards rehashes everything" problem.

The Problem with Simple Hashing

Simple Hash: $\text{shard} = \text{hash}(\text{key}) \% N$

4 shards:

User 123 → hash(123) % 4 = 3 (Shard 3)

User 456 → hash(456) % 4 = 0 (Shard 0)

User 789 → hash(789) % 4 = 1 (Shard 1)

Add 1 shard (now 5 shards):

User 123 → hash(123) % 5 = 3 (Shard 3) ✓ Same!

User 456 → hash(456) % 5 = 1 (Shard 1) ✗ Moved!

User 789 → hash(789) % 5 = 4 (Shard 4) ✗ Moved!

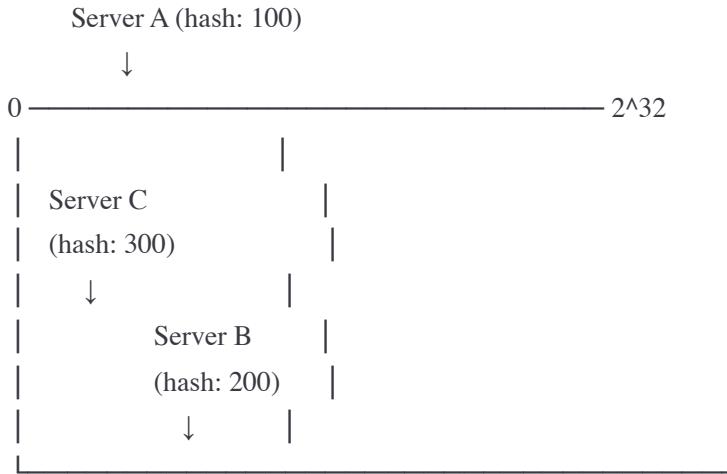
Result: ~80% of keys move to different shards!

Must migrate 80% of data!

Consistent Hashing Solution

Concept: Hash keys and servers onto a ring. Key goes to next server clockwise.

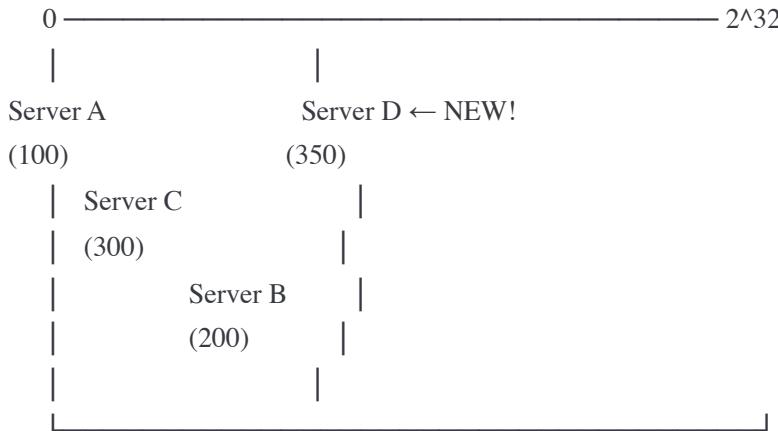
Hash Ring (0 to $2^{32}-1$):



Keys:

- Key X (hash: 50) → Server A (next clockwise)
- Key Y (hash: 150) → Server B (next clockwise)
- Key Z (hash: 250) → Server C (next clockwise)

Add Server D (hash: 350):



Keys affected:

- Key X (50) → Still Server A ✓
- Key Y (150) → Still Server B ✓
- Key Z (250) → Still Server C ✓
- NEW keys (325-350) → Server D

Only ~25% of keys moved! (1/4 servers)

Implementation

python

```
import hashlib
import bisect

class ConsistentHash:
    def __init__(self, nodes=None, virtual_nodes=150):
        """
        nodes: List of server addresses
        virtual_nodes: Number of virtual nodes per physical node
            (helps with even distribution)
        """
        self.virtual_nodes = virtual_nodes
        self.ring = {} # hash_value → node
        self.sorted_keys = [] # Sorted list of hash values
        self.nodes = set()

    if nodes:
        for node in nodes:
            self.add_node(node)

    def _hash(self, key):
        """Hash function: string → integer"""
        return int(hashlib.md5(key.encode()).hexdigest(), 16)

    def add_node(self, node):
        """Add a server to the ring"""
        self.nodes.add(node)

        # Add virtual nodes for better distribution
        for i in range(self.virtual_nodes):
            virtual_key = f"{node}:{i}"
            hash_value = self._hash(virtual_key)

            self.ring[hash_value] = node
            bisect.insort(self.sorted_keys, hash_value)

        print(f"Added node {node} with {self.virtual_nodes} virtual nodes")

    def remove_node(self, node):
        """Remove a server from the ring"""
        self.nodes.discard(node)

        # Remove all virtual nodes
        for i in range(self.virtual_nodes):
            virtual_key = f"{node}:{i}"
            hash_value = self._hash(virtual_key)
```

```
del self.ring[hash_value]
self.sorted_keys.remove(hash_value)

print(f"Removed node {node}")

def get_node(self, key):
    """Find which server should handle this key"""
    if not self.ring:
        return None

    hash_value = self._hash(str(key))

    # Find first node clockwise on ring
    # Binary search for first hash_value >= key's hash
    idx = bisect.bisect_right(self.sorted_keys, hash_value)

    # Wrap around if necessary
    if idx == len(self.sorted_keys):
        idx = 0

    return self.ring[self.sorted_keys[idx]]

def get_distribution(self, num_keys=10000):
    """Analyze how evenly keys are distributed"""
    distribution = {node: 0 for node in self.nodes}

    for i in range(num_keys):
        node = self.get_node(f"key_{i}")
        distribution[node] += 1

    return distribution

# Usage
ch = ConsistentHash([
    'server1.db.com',
    'server2.db.com',
    'server3.db.com'
])

# Get node for keys
print(ch.get_node('user_123')) # → server2.db.com
print(ch.get_node('user_456')) # → server1.db.com
print(ch.get_node('user_789')) # → server3.db.com

# Check distribution
print(ch.get_distribution())
# {
```

```
# 'server1.db.com': 3341, (~33%)
# 'server2.db.com': 3329, (~33%)
# 'server3.db.com': 3330 (~33%)
# }
```

```
# Add new server
ch.add_node('server4.db.com')
```

```
# Check distribution again
print(ch.get_distribution())
# {
# 'server1.db.com': 2501, (~25%)
# 'server2.db.com': 2499, (~25%)
# 'server3.db.com': 2500, (~25%)
# 'server4.db.com': 2500 (~25%)
# }
```

```
# Only ~25% of keys moved to new server!
```

```
# Rest stayed on same servers
```

Virtual Nodes

Problem: With few servers, distribution may be uneven.

Without Virtual Nodes (3 servers):

Server A might get 40% of keys

Server B might get 35% of keys

Server C might get 25% of keys

With Virtual Nodes (3 servers \times 150 vnodes = 450 points on ring):

Server A gets ~33.3% of keys

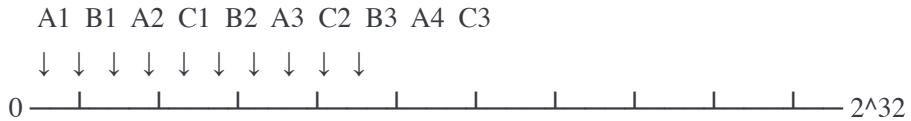
Server B gets ~33.3% of keys

Server C gets ~33.4% of keys

More points = better distribution!

Visual:

Ring with Virtual Nodes:



Each physical server has multiple positions on ring

Keys distributed more evenly

6. Database Partitioning Techniques

Partitioning: Splitting a single table into multiple physical tables.

Vertical Partitioning

Split table by columns

Original Users Table:

ID	Name	Email	City	Bio	Settings
1	John	j@x.com	NYC	Long...	JSON...
2	Jane	ja@x.com	LA	Long...	JSON...

After Vertical Partitioning:

Users_Core (Frequently accessed):

ID	Name	Email	City
1	John	j@x.com	NYC
2	Jane	ja@x.com	LA

Users_Extended (Rarely accessed):

ID	Bio	Settings
1	Long...	JSON...
2	Long...	JSON...

Benefits:

- Smaller, faster queries on core data
- Can optimize storage differently
- Reduce I/O for common queries

SQL Example:

```
sql
-- Original query scans all columns
SELECT name, email FROM users WHERE city = 'NYC';
-- Scans: ID, name, email, city, bio, settings

-- After partitioning
SELECT name, email FROM users_core WHERE city = 'NYC';
-- Scans: ID, name, email, city (less I/O!)

-- Only fetch extended data when needed
SELECT bio FROM users_extended WHERE id = 123;
```

Horizontal Partitioning (Same as Sharding)

Split table by rows

Already covered in sharding section!

Users Table → Split into 4 partitions

Partition 1: Users 0-250K

Partition 2: Users 250K-500K

Partition 3: Users 500K-750K

Partition 4: Users 750K-1M

List Partitioning

Partition by discrete values

sql

-- Partition by country

```
CREATE TABLE users (
    id SERIAL,
    name VARCHAR(100),
    email VARCHAR(255),
    country VARCHAR(2)
) PARTITION BY LIST (country);
```

-- Create partitions

```
CREATE TABLE users_us PARTITION OF users FOR VALUES IN ('US');
```

```
CREATE TABLE users_uk PARTITION OF users FOR VALUES IN ('UK');
```

```
CREATE TABLE users_ca PARTITION OF users FOR VALUES IN ('CA');
```

```
CREATE TABLE users_other PARTITION OF users DEFAULT;
```

-- Queries automatically route to correct partition

```
SELECT * FROM users WHERE country = 'US';
```

-- Only scans users_us partition!

-- Benefit: Can optimize each partition differently

-- Example: US partition on fast SSD, others on slower storage

Range Partitioning (PostgreSQL)

sql

```

-- Partition by date
CREATE TABLE logs (
    id SERIAL,
    timestamp TIMESTAMP,
    message TEXT
) PARTITION BY RANGE (timestamp);

-- Create partitions
CREATE TABLE logs_2024_q1 PARTITION OF logs
    FOR VALUES FROM ('2024-01-01') TO ('2024-04-01');

CREATE TABLE logs_2024_q2 PARTITION OF logs
    FOR VALUES FROM ('2024-04-01') TO ('2024-07-01');

CREATE TABLE logs_2024_q3 PARTITION OF logs
    FOR VALUES FROM ('2024-07-01') TO ('2024-10-01');

CREATE TABLE logs_2024_q4 PARTITION OF logs
    FOR VALUES FROM ('2024-10-01') TO ('2025-01-01');

-- Query automatically prunes partitions
SELECT * FROM logs
WHERE timestamp >= '2024-06-01'
    AND timestamp < '2024-07-01';
-- Only scans logs_2024_q2!

-- Easy to drop old data
DROP TABLE logs_2024_q1; -- Removes old logs instantly

```

Hash Partitioning (PostgreSQL)

sql

```

-- Partition by hash of user_id
CREATE TABLE users (
    id INTEGER,
    name VARCHAR(100),
    email VARCHAR(255)
) PARTITION BY HASH (id);

-- Create partitions
CREATE TABLE users_p0 PARTITION OF users
    FOR VALUES WITH (MODULUS 4, REMAINDER 0);

CREATE TABLE users_p1 PARTITION OF users
    FOR VALUES WITH (MODULUS 4, REMAINDER 1);

CREATE TABLE users_p2 PARTITION OF users
    FOR VALUES WITH (MODULUS 4, REMAINDER 2);

CREATE TABLE users_p3 PARTITION OF users
    FOR VALUES WITH (MODULUS 4, REMAINDER 3);

-- Distributes data evenly across 4 partitions
-- hash(id) % 4 determines partition

```

Key Takeaways

1. Read Replicas:

- Scale reads by adding replicas
- Handle replication lag with smart routing
- Read from master after write for consistency

2. Replication Types:

- Master-Slave: Simple, one source of truth, scales reads
- Master-Master: Complex, scales writes, conflict resolution needed

3. Sharding:

- Split database horizontally for massive scale
- Hash-based: Even distribution, hard to add shards
- Range-based: Easy range queries, potential hotspots
- Directory-based: Flexible, extra lookup overhead

4. Consistent Hashing:

- Solves resharding problem
- Only K/N keys move when adding/removing nodes
- Use virtual nodes for better distribution

5. Partitioning:

- Vertical: Split by columns (hot/cold data)
- Horizontal: Split by rows (sharding)
- List/Range/Hash: Different access patterns

Practice Problems

1. Design a sharding strategy for Instagram with 1 billion users and 500 billion photos.
2. Calculate how much data moves when adding a 5th shard to a 4-shard hash-based system with 1TB total data.
3. Design a multi-region database architecture for a global social network with users in US, Europe, and Asia.

Next Chapter Preview

In Chapter 8, we'll explore **Storage Systems**:

- Block vs Object vs File storage
- Distributed file systems (HDFS, GFS)
- Blob storage (S3, Azure Blob)
- Different database types and when to use them

Ready to continue?