


Chapter 14: Monitoring and Observability

Introduction: The Three Pillars of Observability

Observability: Understanding what's happening inside your system by examining its outputs.



Together they answer:

- What went wrong?
- When did it happen?
- Where did it happen?
- Why did it happen?

Monitoring vs Observability:

MONITORING (Known Unknowns):

- "Is the server up?"
- "Is CPU usage high?"
- "Are we getting errors?"
- Dashboard with predefined metrics

OBSERVABILITY (Unknown Unknowns):

- "Why is this specific user's request slow?"
- "What's the path this request took?"
- "What happened between 10:23:45 and 10:23:46?"
- Ability to ask arbitrary questions

1. Logging Best Practices

Log Levels

Level	When to Use
TRACE	Very detailed (function entry/exit) Usually disabled in production
DEBUG	Detailed diagnostic info Enable when troubleshooting
INFO	General informational messages "User logged in", "Order created"
WARN	Warning: potential problem "Cache miss", "Retry attempted"
ERROR	Error: something failed "Payment failed", "DB connection"
FATAL	Critical: system unusable "Out of memory", "DB down"

Structured Logging

Bad (Unstructured):

```
javascript

// ✗ Hard to parse and search
console.log('User John logged in from 192.168.1.1 at 10:30');
console.log('Order 123 created with total $99.99');
console.log('Payment processing failed - insufficient funds');
```

Good (Structured):

```
javascript
```

```

// ✓ JSON format, easy to parse and query
const logger = require('winston');

logger.info('user_login', {
  userId: 123,
  username: 'john',
  ip: '192.168.1.1',
  timestamp: '2024-01-20T10:30:00Z',
  userAgent: 'Mozilla/5.0...'
});

logger.info('order_created', {
  orderId: 123,
  userId: 456,
  total: 99.99,
  itemCount: 3,
  timestamp: '2024-01-20T10:31:00Z'
});

logger.error('payment_failed', {
  orderId: 123,
  userId: 456,
  amount: 99.99,
  errorCode: 'INSUFFICIENT_FUNDS',
  timestamp: '2024-01-20T10:32:00Z'
});

// Output (JSON):
//{
//  "level": "info",
//  "message": "user_login",
//  "userId": 123,
//  "username": "john",
//  "ip": "192.168.1.1",
//  "timestamp": "2024-01-20T10:30:00Z"
//}

```

Complete Logging Implementation

javascript

```
const winston = require('winston');

// Create logger with multiple transports
const logger = winston.createLogger({
  level: process.env.LOG_LEVEL || 'info',
  format: winston.format.combine(
    winston.format.timestamp(),
    winston.format.errors({ stack: true }),
    winston.format.json()
  ),
  defaultMeta: {
    service: 'order-service',
    version: '1.0.0',
    environment: process.env.NODE_ENV
  },
  transports: [
    // Console output (for development)
    new winston.transports.Console({
      format: winston.format.combine(
        winston.format.colorize(),
        winston.format.simple()
      )
    }),
    // File output (all logs)
    new winston.transports.File({
      filename: 'logs/combined.log'
    }),
    // File output (errors only)
    new winston.transports.File({
      filename: 'logs/error.log',
      level: 'error'
    })
  ]
});

// Enhanced logging with context
class ContextLogger {
  constructor(baseLogger) {
    this.logger = baseLogger;
  }

  child(context) {
    """Create logger with additional context"""
    return this.logger.child(context);
  }
}
```

```
}
```

```
logRequest(req, res, duration) {
    """Log HTTP request"""
    this.logger.info('http_request', {
        method: req.method,
        path: req.path,
        statusCode: res.statusCode,
        duration: duration,
        ip: req.ip,
        userAgent: req.get('user-agent'),
        userId: req.user?.id,
        requestId: req.id,
        query: req.query,
        body: this._sanitize(req.body)
    });
}
```

```
logDatabaseQuery(query, duration, rowCount) {
    """Log database query"""
    this.logger.debug('database_query', {
        query: query,
        duration: duration,
        rowCount: rowCount,
        slow: duration > 100 // Flag slow queries
    });
}
```

```
logError(error, context = {}) {
    """Log error with full stack trace"""
    this.logger.error('error_occurred', {
        errorMessage: error.message,
        errorCode: error.code,
        stack: error.stack,
        ...context
    });
}
```

```
logBusinessEvent(eventType, data) {
    """Log business events"""
    this.logger.info(`business_event:${eventType}`, {
        eventType: eventType,
        ...data
    });
}
```

```
_sanitize(data) {
```

```
"""Remove sensitive data from logs"""
if (!data) return data;

const sanitized = { ...data };

// Remove sensitive fields
const sensitiveFields = ['password', 'creditCard', 'ssn', 'token'];
sensitiveFields.forEach(field => {
  if (sanitized[field]) {
    sanitized[field] = '***REDACTED***';
  }
});

return sanitized;
}

// Usage
const contextLogger = new ContextLogger(logger);

// Request logging middleware
app.use((req, res, next) => {
  const startTime = Date.now();

  // Add request ID for tracing
  req.id = generateRequestId();

  // Log after response
  res.on('finish', () => {
    const duration = Date.now() - startTime;
    contextLogger.logRequest(req, res, duration);
  });

  next();
});

// Business event logging
app.post('/orders', async (req, res) => {
  const order = await createOrder(req.body);

  contextLogger.logBusinessEvent('order_created', {
    orderId: order.id,
    userId: req.user.id,
    total: order.total,
    itemCount: order.items.length
  });
});
```

```

res.json(order);
});

// Error logging
app.use((err, req, res, next) => {
  contextLogger.logError(err, {
    requestId: req.id,
    path: req.path,
    method: req.method,
    userId: req.user?.id
  });
  res.status(500).json({ error: 'Internal server error' });
});

```

Correlation IDs (Request Tracing)

Problem: Request goes through multiple services

How do you correlate logs?

Request Flow:

```

User → API Gateway → User Service → Database
      → Order Service → Database
      → Payment Service

```

Without Correlation ID:

- API Gateway logs: "Request received"
- User Service logs: "User fetched"
- Order Service logs: "Order created"

How do you know these are from same request?

With Correlation ID:

- API Gateway: [req-abc-123] "Request received"
- User Service: [req-abc-123] "User fetched"
- Order Service: [req-abc-123] "Order created"

Now can trace entire request flow!

Implementation:

javascript

```
const { v4: uuidv4 } = require('uuid');

// Generate or propagate correlation ID
function correlationMiddleware(req, res, next) {
  // Get correlation ID from header or generate new one
  req.correlationId = req.get('X-Correlation-ID') || uuidv4();

  // Add to response headers
  res.set('X-Correlation-ID', req.correlationId);

  // Create logger with correlation ID
  req.log = logger.child({
    correlationId: req.correlationId,
    path: req.path,
    method: req.method
  });

  next();
}

app.use(correlationMiddleware);

// All logs now include correlation ID
app.get('/users/:id', async (req, res) => {
  req.log.info('Fetching user');

  const user = await db.query('SELECT * FROM users WHERE id = ?', req.params.id);

  req.log.info('User fetched', { userId: user.id });

  res.json(user);
});

// When calling other services, propagate correlation ID
async function callOrderService(correlationId, data) {
  return await fetch('https://order-service.example.com/orders', {
    method: 'POST',
    headers: {
      'X-Correlation-ID': correlationId // Propagate!
    },
    body: JSON.stringify(data)
  });
}

// Now can search logs by correlation ID:
```

```
// grep "req-abc-123" *.log  
// Shows complete request flow across all services!
```

Logging Best Practices

1. STRUCTURED LOGS (JSON)

- ✓ {"level": "info", "msg": "user_login", "userId": 123}
- ✗ "User 123 logged in"

2. INCLUDE CONTEXT

- ✓ userId, orderId, correlationId, timestamp
- ✗ Just the error message

3. LOG LEVELS APPROPRIATELY

- ✓ INFO for business events
- ✓ ERROR for actual errors
- ✗ ERROR for expected conditions
- ✗ DEBUG in production (too verbose)

4. DON'T LOG SENSITIVE DATA

- ✗ Passwords, credit cards, SSNs, tokens
- ✓ Sanitize or hash sensitive fields

5. ADD CORRELATION IDS

- ✓ Trace requests across services
- ✓ Propagate in headers

6. LOG AT BOUNDARIES

- ✓ Request entry/exit
- ✓ External service calls
- ✓ Database queries (slow ones)

7. INCLUDE TIMING

- ✓ Duration of operations
- ✓ Timestamp of events

8. CENTRALIZE LOGS

- ✓ Send to central system (ELK, Datadog)
- ✗ Leave on individual servers

2. Metrics Collection and Monitoring

Types of Metrics

1. SYSTEM METRICS (Infrastructure)

- CPU usage (%)
- Memory usage (MB)
- Disk usage (GB, IOPS)
- Network I/O (Mbps)
- File descriptors

2. APPLICATION METRICS (Business)

- Requests per second
- Response time (p50, p95, p99)
- Error rate (%)
- Active users
- Orders per minute
- Revenue per hour

3. BUSINESS METRICS (KPIs)

- Conversion rate
- Customer acquisition cost
- Average order value
- Churn rate
- User engagement

The Four Golden Signals (Google SRE)

1. LATENCY

Time to serve a request

Track:

- p50 (median): 50% of requests faster than this
- p95: 95% of requests faster than this
- p99: 99% of requests faster than this

Example:

p50: 100ms (good)

p95: 500ms (acceptable)

p99: 2000ms (some users have bad experience)

2. TRAFFIC

Demand on your system

Track:

- Requests per second
- Transactions per second
- Concurrent users

Example:

Normal: 1,000 req/s

Peak: 10,000 req/s

Alert if > 15,000 req/s (capacity limit)

3. ERRORS

Rate of failed requests

Track:

- Error rate (%)
- Errors per second
- By error type (4xx vs 5xx)

Example:

Normal: 0.1% error rate

Alert if > 1% (degraded)

Alert if > 5% (critical)

4. SATURATION

How "full" your service is

Track:

- CPU usage (%)
- Memory usage (%)
- Disk I/O (%)
- Connection pool usage (%)

Example:

CPU: 60% (normal)


Alert if > 80% (warning)

Alert if > 90% (critical)

Metrics with Prometheus

Prometheus: Pull-based monitoring system

Architecture:



Metrics Implementation (Node.js):

javascript

```
const express = require('express');
const promClient = require('prom-client');

const app = express();

// Create a Registry
const register = new promClient.Registry();

// Add default metrics (CPU, memory, etc.)
promClient.collectDefaultMetrics({ register });

// Custom metrics

// 1. Counter (only goes up)
const httpRequestsTotal = new promClient.Counter({
  name: 'http_requests_total',
  help: 'Total number of HTTP requests',
  labelNames: ['method', 'path', 'status'],
  registers: [register]
});

// 2. Gauge (can go up and down)
const activeUsers = new promClient.Gauge({
  name: 'active_users',
  help: 'Number of currently active users',
  registers: [register]
});

// 3. Histogram (distribution of values)
const httpRequestDuration = new promClient.Histogram({
  name: 'http_request_duration_seconds',
  help: 'Duration of HTTP requests in seconds',
  labelNames: ['method', 'path', 'status'],
  buckets: [0.1, 0.5, 1, 2, 5, 10], // Buckets for histogram
  registers: [register]
});

// 4. Summary (similar to histogram, calculates percentiles)
const paymentProcessingTime = new promClient.Summary({
  name: 'payment_processing_seconds',
  help: 'Payment processing time',
  percentiles: [0.5, 0.9, 0.95, 0.99],
  registers: [register]
});

// Middleware to track metrics
```

```
app.use((req, res, next) => {
  const start = Date.now();

  res.on('finish', () => {
    const duration = (Date.now() - start) / 1000;

    // Increment request counter
    httpRequestsTotal.inc({
      method: req.method,
      path: req.route?.path || req.path,
      status: res.statusCode
    });
  });

  // Record request duration
  httpRequestDuration.observe({
    method: req.method,
    path: req.route?.path || req.path,
    status: res.statusCode
  }, duration);
});

next();
});

// Business metrics
app.post('/orders', async (req, res) => {
  const start = Date.now();

  try {
    const order = await createOrder(req.body);

    // Track successful order
    const orderCounter = new promClient.Counter({
      name: 'orders_created_total',
      help: 'Total orders created',
      registers: [register]
    });
    orderCounter.inc();

    // Track order value
    const orderValue = new promClient.Histogram({
      name: 'order_value_dollars',
      help: 'Order value in dollars',
      buckets: [10, 50, 100, 500, 1000],
      registers: [register]
    });
    orderValue.observe(order.total);
  }
});
```

```

res.json(order);

} catch (error) {
// Track error

const errorCounter = new promClient.Counter({
  name: 'orders_failed_total',
  help: 'Total failed orders',
  labelNames: ['reason'],
  registers: [register]
});

errorCounter.inc({ reason: error.code });

res.status(500).json({ error: error.message });
}

});

// Update active users gauge
setInterval(async () => {
  const count = await getActiveUserCount();
  activeUsers.set(count);
}, 10000); // Update every 10 seconds

// Expose metrics endpoint
app.get('/metrics', async (req, res) => {
  res.set('Content-Type', register.contentType);
  res.end(await register.metrics());
});

app.listen(3000);

// Metrics output (Prometheus format):
// # HELP http_requests_total Total number of HTTP requests
// # TYPE http_requests_total counter
// http_requests_total{method="GET",path="/users",status="200"} 1523
// http_requests_total{method="POST",path="/orders",status="201"} 342
//
// # HELP http_request_duration_seconds Duration of HTTP requests
// # TYPE http_request_duration_seconds histogram
// http_request_duration_seconds_bucket{method="GET",path="/users",le="0.1"} 1200
// http_request_duration_seconds_bucket{method="GET",path="/users",le="0.5"} 1500
// http_request_duration_seconds_sum{method="GET",path="/users"} 152.3
// http_request_duration_seconds_count{method="GET",path="/users"} 1523

```

Prometheus Configuration

yaml

```

# prometheus.yml

global:
  scrape_interval: 15s      # Scrape every 15 seconds
  evaluation_interval: 15s # Evaluate rules every 15 seconds

# Scrape configurations
scrape_configs:
  # Application servers
  - job_name: 'order-service'
    static_configs:
      - targets:
          - 'app-server-1:3000'
          - 'app-server-2:3000'
          - 'app-server-3:3000'
    scrape_interval: 10s

  # Database
  - job_name: 'postgres'
    static_configs:
      - targets: ['postgres-exporter:9187']

  # Redis
  - job_name: 'redis'
    static_configs:
      - targets: ['redis-exporter:9121']

  # Node exporters (system metrics)
  - job_name: 'node'
    static_configs:
      - targets:
          - 'server-1:9100'
          - 'server-2:9100'
          - 'server-3:9100'

# Alert rules
rule_files:
  - 'alerts.yml'

# Alertmanager
alerting:
  alertmanagers:
    - static_configs:
        - targets: ['alertmanager:9093']

```

Alert Rules (alerts.yml):

yaml

groups:

- name: application_alerts

interval: 30s

rules:

High error rate

- alert: HighErrorRate

expr: |

sum(rate(http_requests_total{status=~"5.."}[5m]))

/ sum(rate(http_requests_total[5m])) > 0.05

for: 5m

labels:

severity: critical

annotations:

summary: "High error rate detected"

description: "Error rate is {{ \$value | humanizePercentage }}"

High latency

- alert: HighLatency

expr: |

histogram_quantile(0.95,

sum(rate(http_request_duration_seconds_bucket[5m])) by (le

) > 1

for: 5m

labels:

severity: warning

annotations:

summary: "High latency detected"

description: "95th percentile latency is {{ \$value }}s"

Service down

- alert: ServiceDown

expr: up == 0

for: 1m

labels:

severity: critical

annotations:

summary: "Service {{ \$labels.instance }} is down"

High CPU

- alert: HighCPU

expr: |

100 - (avg by(instance) (rate(node_cpu_seconds_total{mode="idle"}[5m])) * 100) > 80

for: 5m

labels:

severity: warning

annotations:

```
summary: "High CPU on {{ $labels.instance }}"
description: "CPU usage is {{ $value }}%"

# Disk almost full
- alert: DiskSpaceLow
  expr: |
    (node_filesystem_avail_bytes / node_filesystem_size_bytes) < 0.1
  for: 5m
  labels:
    severity: warning
  annotations:
    summary: "Disk space low on {{ $labels.instance }}"
    description: "Only {{ $value | humanizePercentage }} free"
```

PromQL Queries

```
promql
```

```

# Request rate (requests per second)
rate(http_requests_total[5m])

# Error rate (%)
sum(rate(http_requests_total{status=~"5.."}[5m]))
/ sum(rate(http_requests_total[5m])) * 100

# 95th percentile latency
histogram_quantile(0.95,
rate(http_request_duration_seconds_bucket[5m]))
)

# Average response time by endpoint
avg(rate(http_request_duration_seconds_sum[5m]))
by (path)
/ avg(rate(http_request_duration_seconds_count[5m]))
by (path)

# Memory usage
node_memory_MemTotal_bytes - node_memory_MemAvailable_bytes

# Disk usage (%)
(1 - node_filesystem_avail_bytes / node_filesystem_size_bytes) * 100

# Active connections
sum(pg_stat_database_numbackends) by (datname)

# Queue depth
rabbitmq_queue_messages_ready

# Top 5 slowest endpoints
topk(5,
avg(rate(http_request_duration_seconds_sum[5m])) by (path)
)

```

3. Distributed Tracing


What is Distributed Tracing?

Problem: Request goes through multiple services

Without Tracing:

User reports: "Checkout is slow"


Where's the problem?



Total: 2000ms (slow!)

But which service caused it?


With Tracing:



Now we know exactly where to optimize!

Trace Structure

Trace: Complete request flow



Each span contains:

- Span ID (unique)
- Trace ID (same for all spans in trace)
- Parent Span ID (to build hierarchy)
- Service name
- Operation name
- Start timestamp
- Duration
- Tags (metadata)
- Logs (events during span)

Tracing Implementation (OpenTelemetry)

javascript

```
const { NodeTracerProvider } = require('@opentelemetry/sdk-trace-node');
const { registerInstrumentations } = require('@opentelemetry/instrumentation');
const { HttpInstrumentation } = require('@opentelemetry/instrumentation-http');
const { ExpressInstrumentation } = require('@opentelemetry/instrumentation-express');
const { JaegerExporter } = require('@opentelemetry/exporter-jaeger');
const { Resource } = require('@opentelemetry/resources');
const { SemanticResourceAttributes } = require('@opentelemetry/semantic-conventions');

// Initialize tracer
const provider = new NodeTracerProvider({
  resource: new Resource({
    [SemanticResourceAttributes.SERVICE_NAME]: 'order-service',
    [SemanticResourceAttributes.SERVICE_VERSION]: '1.0.0'
  })
});

// Configure exporter (send to Jaeger)
const exporter = new JaegerExporter({
  endpoint: 'http://jaeger:14268/api/traces'
});

provider.addSpanProcessor(
  new SimpleSpanProcessor(exporter)
);

// Register provider
provider.register();

// Auto-instrument HTTP and Express
registerInstrumentations({
  instrumentations: [
    new HttpInstrumentation(),
    new ExpressInstrumentation()
  ]
});

// Get tracer for manual instrumentation
const tracer = provider.getTracer('order-service');

// Application code with tracing
app.post('/orders', async (req, res) => {
  // Create span for this operation
  const span = tracer.startSpan('create_order');

  // Add attributes (tags)
  span.setAttribute('user.id', req.user.id);
});
```

```
span.setAttribute('order.total', req.body.total);

try {
  // Step 1: Validate order
  const validateSpan = tracer.startSpan('validate_order', {
    parent: span
  });

  const order = await validateOrder(req.body);
  validateSpan.end();

  // Step 2: Check inventory
  const inventorySpan = tracer.startSpan('check_inventory', {
    parent: span
  });

  const available = await checkInventory(order.items);
  inventorySpan.setAttribute('inventory.available', available);
  inventorySpan.end();

  if (!available) {
    span.setAttribute('error', true);
    span.setAttribute('error.reason', 'out_of_stock');
    throw new Error('Out of stock');
  }

  // Step 3: Process payment
  const paymentSpan = tracer.startSpan('process_payment', {
    parent: span
  });

  const payment = await processPayment(order.total);
  paymentSpan.setAttribute('payment.id', payment.id);
  paymentSpan.end();

  // Step 4: Create order
  const dbSpan = tracer.startSpan('database_insert', {
    parent: span
  });

  const savedOrder = await db.createOrder(order);
  dbSpan.end();

  span.setAttribute('order.id', savedOrder.id);
  span.setStatus({ code: SpanStatusCode.OK });

  res.json(savedOrder);
```

```

} catch (error) {
  // Record error in span
  span.recordException(error);
  span.setStatus({
    code: SpanStatusCode.ERROR,
    message: error.message
  });

  res.status(500).json({ error: error.message });

} finally {
  span.end();
}
};

// Propagate trace context to downstream services
async function callUserService(userId, traceContext) {
  const span = tracer.startSpan('call_user_service');

  try {
    const response = await fetch(`https://user-service/users/${userId}`, {
      headers: {
        // Propagate trace context
        'traceparent': traceContext.traceparent,
        'tracestate': traceContext.tracestate
      }
    });

    const user = await response.json();
    span.setAttribute('user.name', user.name);

    return user;
  } finally {
    span.end();
  }
}

```

Trace Analysis

Analyzing slow request:

Trace ID: abc-123-def-456

Total Duration: 2.5 seconds

Breakdown:

Service	Operation	Duration	
API Gateway	authenticate	50ms	
API Gateway	route	10ms	
User Service	get_user	80ms	
Order Service	create_order	1900ms !	
	validate_items	100ms	
	check_inventory	1500ms ! !	
	database_query	1450ms ! !	
	calculate_total	50ms	
Payment Service	charge	400ms	
API Gateway	format_response	60ms	

Root cause found:

Order Service → check_inventory → database_query (1450ms)

Action items:

1. Add database index on inventory table
2. Add caching layer
3. Optimize query (was doing full table scan)

After optimization:

inventory check: 1500ms → 50ms

Total request: 2500ms → 600ms (4x faster!)

4. Alerting Strategies

Alert Design Principles

GOOD ALERTS:

- ✓ Actionable (can do something about it)
- ✓ Relevant (impacts users or will soon)
- ✓ Timely (detected quickly)
- ✓ Context-rich (enough info to debug)

BAD ALERTS:

- ✗ "Server CPU is high" (so what? is it affecting users?)
- ✗ Too many alerts (alert fatigue)
- ✗ Too sensitive (false positives)
- ✗ Not sensitive enough (miss real issues)

Alert Severity Levels

Severity	Condition	Response	Example
P0	User-facing outage	Wake up anyone immediately	Site down failing
P1	Degraded service	Page on-call during work hours	Slow response times
P2	Warning sign	Create ticket	Disk 80% full
P3	Informational	Review next day	Cache hit rate low

Alert Implementation

```
javascript
```

```
class AlertManager {
  constructor() {
    this.alerts = new Map();
    this.alertChannels = {
      pagerduty: new PagerDutyClient(),
      slack: new SlackClient(),
      email: new EmailClient()
    };
  }

  async checkAndAlert(metricName, value, threshold, severity) {
    const alertKey = `${metricName}_${severity}`;

    // Check if threshold exceeded
    if (value > threshold) {
      // Check if already alerted (prevent spam)
      if (this.alerts.has(alertKey)) {
        const lastAlert = this.alerts.get(alertKey);
        const timeSinceLastAlert = Date.now() - lastAlert;

        // Don't re-alert within 15 minutes
        if (timeSinceLastAlert < 15 * 60 * 1000) {
          return;
        }
      }
    }

    // Send alert
    await this.sendAlert({
      severity,
      metric: metricName,
      value,
      threshold,
      message: `${metricName} is ${value} (threshold: ${threshold})`
    });

    // Record alert time
    this.alerts.set(alertKey, Date.now());
  }

} else {
  // Metric back to normal
  if (this.alerts.has(alertKey)) {
    // Send recovery notification
    await this.sendRecovery({
      metric: metricName,
      message: `${metricName} is back to normal (${value})`
    });
  }
}
```

```
        this.alerts.delete(alertKey);
    }
}

}

async sendAlert(alert) {
    console.log(`🚨 ALERT [${alert.severity}]: ${alert.message}`);

    // Route based on severity
    switch (alert.severity) {
        case 'critical':
            // Page on-call engineer
            await this.alertChannels.pagerduty.trigger({
                title: alert.message,
                severity: 'critical',
                details: alert
            });

            // Also send to Slack
            await this.alertChannels.slack.send({
                channel: '#incidents',
                text: `🚨 CRITICAL: ${alert.message}`,
                color: 'danger'
            });
            break;

        case 'warning':
            // Send to Slack only
            await this.alertChannels.slack.send({
                channel: '#alerts',
                text: `⚠️ WARNING: ${alert.message}`,
                color: 'warning'
            });
            break;

        case 'info':
            // Just log it
            console.log(`ℹ️ INFO: ${alert.message}`);
            break;
    }
}

async sendRecovery(recovery) {
    console.log(`✅ RECOVERED: ${recovery.message}`);

    await this.alertChannels.slack.send({
```

```

channel: '#alerts',
text: `✓ RECOVERED: ${recovery.message}` ,
color: 'good'
});
}

}

// Usage with Prometheus metrics
const alertManager = new AlertManager();

setInterval(async () => {
// Query Prometheus for metrics
const errorRate = await getErrorRate();
const latency = await getP95Latency();
const cpu = await getCPUUsage();

// Check thresholds and alert
await alertManager.checkAndAlert('error_rate', errorRate, 5, 'critical');
await alertManager.checkAndAlert('p95_latency', latency, 1000, 'warning');
await alertManager.checkAndAlert('cpu_usage', cpu, 80, 'warning');

}, 30000); // Check every 30 seconds

```

Alert Escalation

javascript

```
class AlertEscalation {
  async handleAlert(alert) {
    const severity = alert.severity;
    const metric = alert.metric;

    // Escalation levels
    const escalation = {
      critical: [
        { delay: 0, action: () => this.pageOnCall(alert) },
        { delay: 5 * 60, action: () => this.pageBackup(alert) },
        { delay: 15 * 60, action: () => this.pageManager(alert) }
      ],
      warning: [
        { delay: 0, action: () => this.slackAlert(alert) },
        { delay: 30 * 60, action: () => this.emailTeam(alert) }
      ]
    };
  }

  const levels = escalation[severity] || [];

  for (const level of levels) {
    setTimeout(async () => {
      // Check if alert still active
      if (await this.isAlertStillActive(metric)) {
        console.log(`Escalating alert: ${metric} (${level.delay}s elapsed}`);
        await level.action();
      } else {
        console.log(`Alert resolved before escalation: ${metric}`);
      }
    }, level.delay * 1000);
  }
}

async pageOnCall(alert) {
  console.log(`PagerDuty API call`);
  // PagerDuty API call
}

async pageBackup(alert) {
  console.log(`Paging backup on-call`);
  // Alert hasn't been acknowledged, escalate
}

async pageManager(alert) {
  console.log(`Paging engineering manager`);
  // Still not resolved after 15 minutes!
}
```

```

}

async slackAlert(alert) {
  console.log('💬 Sending Slack alert');
}

async emailTeam(alert) {
  console.log('✉️ Emailing team');
}

async isAlertStillActive(metric) {
  // Check if metric still above threshold
  return true; // Simplified
}

```

SLIs, SLOs, and SLAs

SLI (Service Level Indicator):

- Measurement of service behavior
- Examples: latency, error rate, availability

SLO (Service Level Objective):

- Target for SLI
- Examples: p95 latency < 200ms, 99.9% availability

SLA (Service Level Agreement):

- Contract with customers
- What happens if SLO not met (refund, credits)

Example:

SLI: Request success rate

Measurement: successful_requests / total_requests

SLO: 99.9% of requests succeed

Target: success_rate >= 0.999

SLA: If success rate < 99.9% for a month

Penalty: 10% service credit to customer

Error Budget:

SLO: 99.9% availability

Allowed downtime: $0.1\% = 43.2$ minutes per month

Error Budget: 43.2 minutes/month

Usage:

Week 1: 10 minutes downtime → 33.2 minutes remaining

Week 2: 5 minutes downtime → 28.2 minutes remaining

Week 3: 30 minutes downtime → BUDGET EXCEEDED! ⚡

Action when budget exceeded:

- Freeze feature development
- Focus on reliability
- No risky deployments
- Fix issues causing downtime

When budget healthy:

- Can take more risks
- Deploy faster
- Try new features

Implementation:

javascript

```
class ErrorBudget {
  constructor(slo, periodDays) {
    this.slo = slo; // e.g., 0.999 for 99.9%
    this.periodDays = periodDays;
    this.incidents = [];
  }

  recordIncident(durationMinutes) {
    this.incidents.push({
      duration: durationMinutes,
      timestamp: new Date()
    });

    // Remove incidents outside the period
    const cutoff = new Date();
    cutoff.setDate(cutoff.getDate() - this.periodDays);

    this.incidents = this.incidents.filter(
      incident => incident.timestamp > cutoff
    );
  }

  getStatus() {
    const totalMinutesInPeriod = this.periodDays * 24 * 60;
    const allowedDowntime = totalMinutesInPeriod * (1 - this.slo);

    const actualDowntime = this.incidents.reduce(
      (sum, incident) => sum + incident.duration,
      0
    );

    const remaining = allowedDowntime - actualDowntime;
    const percentUsed = (actualDowntime / allowedDowntime) * 100;

    return {
      slo: this.slo * 100,
      period: `${this.periodDays} days`,
      allowedDowntime: allowedDowntime.toFixed(2),
      actualDowntime: actualDowntime.toFixed(2),
      remaining: remaining.toFixed(2),
      percentUsed: percentUsed.toFixed(2),
      status: percentUsed > 100 ? 'EXCEEDED' : 'OK'
    };
  }

  canDeploy() {

```

```
const status = this.getStatus();

// If used > 80% of budget, be cautious
if (status.percentUsed > 80) {
  console.log('⚠️ Error budget almost exhausted!');
  console.log('Consider delaying risky deployments');
  return false;
}

return true;
}

// Usage
const errorBudget = new ErrorBudget(0.999, 30); // 99.9% over 30 days

// Record incidents
errorBudget.recordIncident(10); // 10 minute outage
errorBudget.recordIncident(5); // 5 minute outage

// Check status
const status = errorBudget.getStatus();
console.log('Error Budget Status:', status);

// Output:
//{
//  slo: 99.9,
//  period: '30 days',
//  allowedDowntime: '43.20',
//  actualDowntime: '15.00',
//  remaining: '28.20',
//  percentUsed: '34.72',
//  status: 'OK'
//}

// Before deployment
if (errorBudget.canDeploy()) {
  console.log('✓ Safe to deploy');
  deploy();
} else {
  console.log('✗ Focus on reliability first');
}
```

5. Popular Tools

ELK Stack (Elasticsearch, Logstash, Kibana)

Architecture:



Logstash Configuration:

```
ruby
```

```
# logstash.conf

input {
    # Read from application log files
    file {
        path => "/var/log/app/*.log"
        start_position => "beginning"
        type => "application"
    }
}

# Read from syslog
syslog {
    port => 5000
    type => "syslog"
}

# Read from beats (lightweight shippers)
beats {
    port => 5044
}
}

filter {
    # Parse JSON logs
    if [type] == "application" {
        json {
            source => "message"
        }
    }

    # Add hostname
    mutate {
        add_field => { "hostname" => "%{[host][name]}" }
    }
}

# Parse timestamp
date {
    match => [ "timestamp", "ISO8601" ]
    target => "@timestamp"
}
}

# Grok parsing for unstructured logs
grok {
    match => {
        "message" => "%{TIMESTAMP_ISO8601:timestamp} %{LOGLEVEL:level} %{GREEDYDATA:message}"
    }
}
```

```

}

# GeoIP lookup
if [ip] {
  geoip {
    source => "ip"
    target => "geoip"
  }
}

# Drop debug logs in production
if [level] == "debug" {
  drop { }
}

output {
  # Send to Elasticsearch
  elasticsearch {
    hosts => ["elasticsearch:9200"]
    index => "logs-%{+YYYY.MM.dd}"
  }

  # Also output to console (for debugging)
  stdout {
    codec => rubydebug
  }

  # Send errors to Slack
  if [level] == "error" {
    slack {
      url => "https://hooks.slack.com/services/..."
      channel => "#errors"
      username => "Logstash"
    }
  }
}

```

Elasticsearch Queries

```

# Search for errors in last hour
GET /logs-*/_search
{
  "query": {
    "bool": {

```

```
"must": [
  { "match": { "level": "error" } },
  { "range": { "@timestamp": { "gte": "now-1h" } } }
],
},
},
},
"sort": [ { "@timestamp": { "order": "desc" } } ]
}
```

Count errors by service

```
GET /logs-*/_search
{
  "query": {
    "match": { "level": "error" }
  },
  "aggs": {
    "by_service": {
      "terms": { "field": "service.keyword" }
    }
  }
}
```

Find slow requests (> 1 second)

```
GET /logs-*/_search
{
  "query": {
    "bool": {
      "must": [
        { "match": { "message": "http_request" } },
        { "range": { "duration": { "gte": 1000 } } }
      ]
    }
  }
}
```

Trace request by correlation ID

```
GET /logs-*/_search
{
  "query": {
    "match": { "correlationId": "req-abc-123" }
  },
  "sort": [ { "@timestamp": { "order": "asc" } } ]
}
```

Grafana Dashboard

Dashboard Configuration (JSON):

```
json
```

```
{  
  "dashboard": {  
    "title": "Application Performance",  
    "panels": [  
      {  
        "title": "Request Rate",  
        "targets": [{  
          "expr": "sum(rate(http_requests_total[5m]))"  
        }],  
        "type": "graph"  
      },  
      {  
        "title": "Error Rate",  
        "targets": [{  
          "expr": "sum(rate(http_requests_total{status=~\"5..\"}[5m])) / sum(rate(http_requests_total[5m])) * 100"  
        }],  
        "type": "graph",  
        "alert": {  
          "conditions": [{  
            "evaluator": {  
              "params": [5],  
              "type": "gt"  
            }  
          }]  
        }]  
      },  
      {  
        "title": "Response Time (p95)",  
        "targets": [{  
          "expr": "histogram_quantile(0.95, sum(rate(http_request_duration_seconds_bucket[5m])) by (le))"  
        }],  
        "type": "graph"  
      },  
      {  
        "title": "Active Users",  
        "targets": [{  
          "expr": "active_users"  
        }],  
        "type": "stat"  
      }  
    ]  
  }  
}
```

Datadog Implementation

javascript

```
const StatsD = require('node-dogstatsd').StatsD;

// Initialize Datadog client
const dogstatsd = new StatsD('localhost', 8125);

// Add default tags
dogstatsd.globalTags = [
  `service:order-service`,
  `env:${process.env.NODE_ENV}`,
  `version:1.0.0`
];

class DatadogMetrics {
  constructor(client) {
    this.client = client;
  }

  // Increment counter
  incrementCounter(metric, value = 1, tags = []) {
    this.client.increment(metric, value, tags);
  }

  // Set gauge (current value)
  gauge(metric, value, tags = []) {
    this.client.gauge(metric, value, tags);
  }

  // Record histogram
  histogram(metric, value, tags = []) {
    this.client.histogram(metric, value, tags);
  }

  // Timing (convenience for duration)
  timing(metric, duration, tags = []) {
    this.client.timing(metric, duration, tags);
  }

  // Set (count unique values)
  set(metric, value, tags = []) {
    this.client.set(metric, value, tags);
  }
}

const metrics = new DatadogMetrics(dogstatsd);

// Track HTTP requests
```

```
app.use((req, res, next) => {
  const start = Date.now();

  res.on('finish', () => {
    const duration = Date.now() - start;

    const tags = [
      `method:${req.method}`,
      `path:${req.route?.path || 'unknown'}`,
      `status:${res.statusCode}`
    ];

    // Increment request counter
    metrics.incrementCounter('http.requests', 1, tags);

    // Record duration
    metrics.histogram('http.request.duration', duration, tags);

    // Track errors separately
    if (res.statusCode >= 500) {
      metrics.incrementCounter('http.errors', 1, tags);
    }
  });

  next();
});

// Business metrics
app.post('/orders', async (req, res) => {
  try {
    const order = await createOrder(req.body);

    // Track order creation
    metrics.incrementCounter('orders.created', 1, [
      `user_tier:${req.user.tier}`
    ]);

    // Track order value
    metrics.histogram('orders.value', order.total);

    // Track active orders
    const activeOrders = await getActiveOrderCount();
    metrics.gauge('orders.active', activeOrders);

    res.json(order);
  } catch (error) {
```

```
metrics.incrementCounter('orders.failed', 1, [
  `reason:${error.code}`
]);

res.status(500).json({ error: error.message });
}

});

// Custom business metrics
setInterval(async () => {
  // Revenue metrics
  const revenue = await getHourlyRevenue();
  metrics.gauge('revenue.hourly', revenue);

  // User metrics
  const activeUsers = await getActiveUserCount();
  metrics.gauge('users.active', activeUsers);

  // Database metrics
  const dbConnections = await getDBConnectionCount();
  metrics.gauge('database.connections', dbConnections);
}, 60000); // Every minute
```

Complete Monitoring Setup

javascript

```
class MonitoringService {
  constructor() {
    // Metrics
    this.prometheus = this.initPrometheus();
    this.datadog = this.initDatadog();

    // Logging
    this.logger = this.initLogger();

    // Tracing
    this.tracer = this.initTracer();
  }

  initPrometheus() {
    const register = new promClient.Registry();
    promClient.collectDefaultMetrics({ register });

    return {
      register,
      counter: (name, help, labels) =>
        new promClient.Counter({ name, help, labelNames: labels, registers: [register] }),
      gauge: (name, help, labels) =>
        new promClient.Gauge({ name, help, labelNames: labels, registers: [register] }),
      histogram: (name, help, labels, buckets) =>
        new promClient.Histogram({ name, help, labelNames: labels, buckets, registers: [register] })
    };
  }

  initDatadog() {
    return new StatsD('localhost', 8125);
  }

  initLogger() {
    return winston.createLogger({
      level: 'info',
      format: winston.format.json(),
      transports: [
        new winston.transports.File({ filename: 'error.log', level: 'error' }),
        new winston.transports.File({ filename: 'combined.log' })
      ]
    });
  }

  initTracer() {
    // OpenTelemetry setup
    return provider.getTracer('my-service');
  }
}
```

```
}

// Unified interface
trackRequest(req, res, duration) {
  const tags = {
    method: req.method,
    path: req.path,
    status: res.statusCode
  };

  // Log
  this.logger.info('http_request', {
    ...tags,
    duration,
    correlationId: req.correlationId
  });

  // Metrics (both Prometheus and Datadog)
  this.prometheus.counter('http_requests_total', 'HTTP requests', ['method', 'path', 'status'])
    .inc(tags);

  this.prometheus.histogram(
    'http_request_duration_seconds',
    'Request duration',
    ['method', 'path'],
    [0.1, 0.5, 1, 2, 5]
  ).observe({ method: tags.method, path: tags.path }, duration / 1000);

  this.datadog.increment('http.requests', 1, [
    `method:${tags.method}`,
    `path:${tags.path}`,
    `status:${tags.status}`
  ]);

  this.datadog.histogram('http.duration', duration, [
    `method:${tags.method}`,
    `path:${tags.path}`
  ]);
}

trackError(error, context) {
  // Log with full context
  this.logger.error('error', {
    message: error.message,
    stack: error.stack,
    ...context
  });
}
```

```
// Increment error counter
this.prometheus.counter('errors_total', 'Errors', ['type'])
  .inc({ type: error.code || 'unknown' });

this.datadog.increment('errors', 1, [
  `type:${error.code || 'unknown'}`
]);
}

trackBusinessEvent(eventType, data) {
  // Log
  this.logger.info(eventType, data);

  // Metrics
  this.prometheus.counter('business_events_total', 'Business events', ['type'])
    .inc({ type: eventType });

  this.datadog.increment('business.events', 1, [`type:${eventType}`]);
}

async getMetrics() {
  return await this.prometheus.register.metrics();
}

// Global monitoring instance
const monitoring = new MonitoringService();

// Use throughout application
app.use((req, res, next) => {
  const start = Date.now();

  res.on('finish', () => {
    monitoring.trackRequest(req, res, Date.now() - start);
  });

  next();
});

app.post('/orders', async (req, res) => {
  try {
    const order = await createOrder(req.body);

    monitoring.trackBusinessEvent('order_created', {
      orderId: order.id,
      total: order.total
    });
  }
});
```

```

    });

    res.json(order);

} catch (error) {
  monitoring.trackError(error, {
    userId: req.user.id,
    path: req.path
  });

  res.status(500).json({ error: error.message });
}

});

// Expose metrics
app.get('/metrics', async (req, res) => {
  res.set('Content-Type', 'text/plain');
  res.send(await monitoring.getMetrics());
});

```

Key Takeaways

1. Three Pillars:

- Logs: What happened (events)
- Metrics: How much (numbers)
- Traces: Where (request path)
- Use all three together!

2. Logging:

- Structured (JSON)
- Correlation IDs
- Appropriate levels
- Sanitize sensitive data

3. Metrics:

- Four Golden Signals (latency, traffic, errors, saturation)
- Counter, Gauge, Histogram
- Prometheus or Datadog

4. Tracing:

- Distributed tracing across services
- Find bottlenecks
- OpenTelemetry standard

5. Alerting:

- SLIs, SLOs, SLAs
- Error budgets
- Severity levels
- Escalation policies

6. Tools:

- Prometheus + Grafana: Open source, powerful
- ELK Stack: Log aggregation and search
- Datadog: All-in-one, managed

Practice Problems

1. Design a monitoring strategy for a microservices architecture with 10 services. What metrics would you track?
2. You have 99.9% SLO. You've had 30 minutes downtime this month. Can you deploy a risky feature?
3. Design an alerting policy for an e-commerce site. What alerts and thresholds?

Ready for the next chapter or want to explore monitoring deeper?