

Chapter 17: Data Processing at Scale

Introduction: Big Data Challenges

When you have massive amounts of data, traditional processing doesn't work.

Small Data (Traditional):

- | | |
|---------------------------|--|
| 1 GB of data | |
| Process on single machine | |
| Time: 10 minutes | |
| ✓ Simple | |

Big Data (Scale):

- | | |
|--------------------------|--|
| 1 PB (1,000,000 GB) | |
| Single machine: | |
| • Would take 10M minutes | |
| • = 19 years! | |
| • Won't fit in memory | |
| ✗ Impossible | |

Solution: Distributed Processing

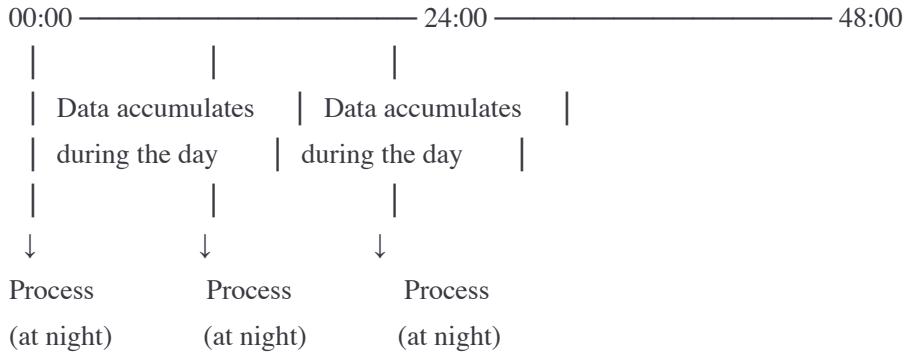
- | | |
|----------------------------|--|
| 1 PB of data | |
| 1,000 machines in parallel | |
| Time: 10 minutes | |
| ✓ Practical | |

1. Batch Processing vs Stream Processing

Batch Processing

Concept: Process large volumes of data at once, periodically.

Batch Processing Timeline:



Characteristics:

- Processes data in large chunks
- Run periodically (hourly, daily, weekly)
- High latency (hours to days)
- High throughput (TB to PB per job)
- Complex analytics possible

Example: E-commerce Daily Sales Report

Daily Batch Job (runs at midnight):

Input:

- 10 million orders from past 24 hours
- 50 GB of transaction data

Process:

1. Load all orders (6 PM - 12 AM)

2. Calculate metrics:

- Total revenue
- Revenue by category
- Top products
- Customer segments
- Inventory needs

3. Generate reports

Output:

- Sales dashboard
- Email to management
- Inventory recommendations

Duration: 2 hours

Freshness: Up to 26 hours old (24h + 2h processing)

Batch Processing Code:

```
python
```

```
from pyspark.sql import SparkSession

# Initialize Spark
spark = SparkSession.builder \
    .appName("DailySalesReport") \
    .getOrCreate()

# Read data (24 hours of orders)
orders = spark.read.parquet("s3://data/orders/2024-01-20/*.parquet")

# Batch processing
daily_report = orders \
    .filter(orders.status == 'completed') \
    .groupBy('category') \
    .agg(
        sum('total').alias('revenue'),
        count('order_id').alias('order_count'),
        avg('total').alias('avg_order_value')
    ) \
    .orderBy('revenue', ascending=False)

# Save results
daily_report.write.parquet("s3://reports/daily-sales/2024-01-20/")

# Show top categories
daily_report.show(10)

# Output:
# +-----+-----+-----+-----+
# | category | revenue | order_count | avg_order_value |
# +-----+-----+-----+-----+
# | Electronics | 1250000 | 5234 | 238.79 |
# | Clothing | 890000 | 12456 | 71.45 |
# | Books | 456000 | 23123 | 19.72 |
# +-----+-----+-----+-----+
```

Stream Processing

Concept: Process data continuously as it arrives.

Stream Processing Timeline:

Event 1 → Process → Result (0.1s)

Event 2 → Process → Result (0.1s)

Event 3 → Process → Result (0.1s)

Event 4 → Process → Result (0.1s)

...continuous...

Characteristics:

- Processes data in real-time
- Event-by-event or micro-batches
- Low latency (milliseconds to seconds)
- Lower throughput per instance
- Simpler per-event operations

Example: Real-time Fraud Detection

Stream Processing (processes every transaction):

Event arrives → Process → Alert (within 100ms)

Transaction 1: \$50 at Starbucks NYC

↓ Process in 50ms

✓ Normal (no alert)

Transaction 2: \$5000 at Electronics Store LA

↓ Process in 80ms

⚠ Unusual amount + unusual location

→ Alert sent immediately!

Transaction 3: \$10 at Gas Station LA

↓ Process in 45ms

⚠ After flagged transaction, same location

→ Confirm fraud pattern

→ Block card immediately!

Latency: 50-100ms per transaction

Freshness: Real-time (immediate)

Stream Processing Code:

javascript

```
const { Kafka } = require('kafkajs');

const kafka = new Kafka({ brokers: ['kafka:9092'] });
const consumer = kafka.consumer({ groupId: 'fraud-detector' });

// Real-time fraud detection
async function streamProcessing() {
  await consumer.connect();
  await consumer.subscribe({ topic: 'transactions' });

  const userHistory = new Map();

  await consumer.run({
    eachMessage: async ({ message }) => {
      const transaction = JSON.parse(message.value);
      const startTime = Date.now();

      // Get user's recent transactions
      if (!userHistory.has(transaction.userId)) {
        userHistory.set(transaction.userId, []);
      }

      const history = userHistory.get(transaction.userId);

      // Detect fraud patterns (real-time!)
      const fraudScore = calculateFraudScore(transaction, history);

      if (fraudScore > 0.8) {
        // High fraud risk - alert immediately!
        await alertFraud({
          transactionId: transaction.id,
          userId: transaction.userId,
          amount: transaction.amount,
          fraudScore,
          processingTime: Date.now() - startTime
        });
      }

      // Block card
      await blockCard(transaction.userId);

      console.log(`⚠️ FRAUD DETECTED (${Date.now() - startTime}ms)`);
    }
  });

  // Update history (keep last 20)
  history.push(transaction);
  if (history.length > 20) {

```

```

    history.shift();
}

console.log(`Processed transaction in ${Date.now() - startTime}ms`);
}
});

}

function calculateFraudScore(transaction, history) {
let score = 0;

// Pattern 1: Large amount
const avgAmount = history.reduce((sum, t) => sum + t.amount, 0) / history.length;
if (transaction.amount > avgAmount * 5) {
  score += 0.4;
}

// Pattern 2: Geographic anomaly
if (history.length > 0) {
  const lastLoc = history[history.length - 1].location;
  const distance = calculateDistance(lastLoc, transaction.location);
  const timeDiff = transaction.timestamp - history[history.length - 1].timestamp;

  if (distance > 1000 && timeDiff < 3600000) { // 1000km in 1 hour
    score += 0.5;
  }
}

// Pattern 3: Multiple transactions quickly
const recent = history.filter(t =>
  transaction.timestamp - t.timestamp < 300000 // Last 5 minutes
);
if (recent.length >= 5) {
  score += 0.3;
}

return Math.min(score, 1.0);
}

streamProcessing();

```

Batch vs Stream Comparison

Feature	Batch	Stream
---------	-------	--------

Data Volume	Large (TB-PB)	Continuous	
Latency	Hours/Days	Milliseconds/Sec	
Processing	Periodic	Continuous	
Complexity	Complex queries	Simple per-event	
Throughput	Very high	Medium	
Use Case	Reports,ETL	Monitoring,fraud	
Examples			
	Daily sales	Click tracking	
	ML training	Real-time alerts	
	Data warehouse	Live dashboards	
	Backup/Archive	Fraud detection	

Latency vs Throughput Trade-off:

Batch: Low latency ✗, High throughput ✓

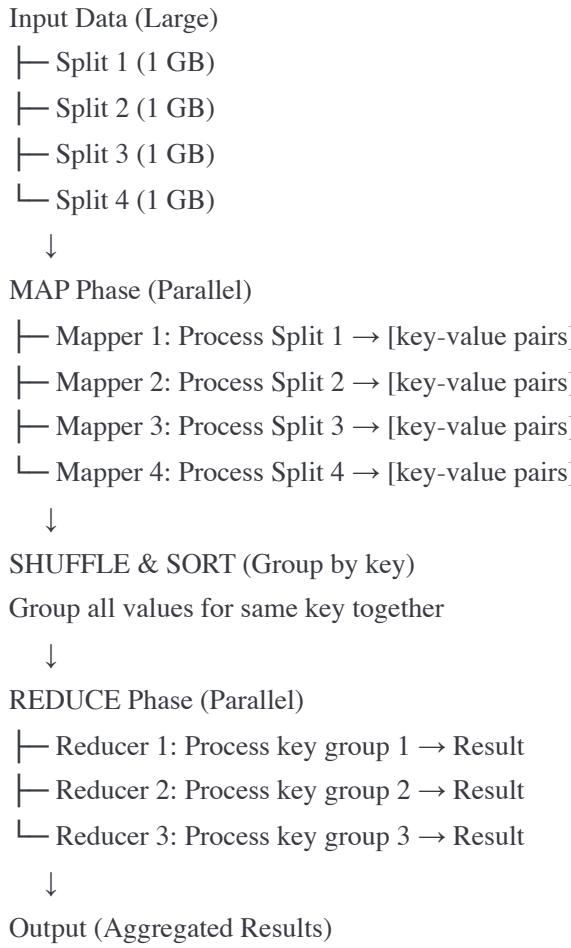
Stream: Low latency ✓, Lower throughput ✗

2. MapReduce Paradigm

What is MapReduce?

Concept: Process massive datasets by splitting into map and reduce phases.

MapReduce Workflow:



MapReduce Example: Word Count

Problem: Count word frequency in 1 TB of text files.

Input Files (4 files):

File 1: "hello world"

File 2: "hello hadoop"

File 3: "world of hadoop"

File 4: "hello world hadoop"

MAP Phase (4 mappers in parallel):

Mapper 1 (File 1):

Input: "hello world"

Output: [("hello", 1), ("world", 1)]

Mapper 2 (File 2):

Input: "hello hadoop"

Output: [("hello", 1), ("hadoop", 1)]

Mapper 3 (File 3):

Input: "world of hadoop"

Output: [("world", 1), ("of", 1), ("hadoop", 1)]

Mapper 4 (File 4):

Input: "hello world hadoop"

Output: [("hello", 1), ("world", 1), ("hadoop", 1)]

SHUFFLE & SORT:

Group by key:

"hello": [1, 1, 1] (from mappers 1, 2, 4)

"world": [1, 1, 1] (from mappers 1, 3, 4)

"hadoop": [1, 1, 1] (from mappers 2, 3, 4)

"of": [1] (from mapper 3)

REDUCE Phase (parallel):

Reducer 1:

Input: ("hello", [1, 1, 1])

Output: ("hello", 3)

Reducer 2:

Input: ("world", [1, 1, 1])

Output: ("world", 3)

Reducer 3:

Input: ("hadoop", [1, 1, 1])

Output: ("hadoop", 3)

Reducer 4:

Input: ("of", [1])

Output: ("of", 1)

Final Output:

hello: 3

world: 3

hadoop: 3

of: 1

MapReduce Code (Hadoop)

java

```

// Mapper
public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        // Input: Line of text
        String line = value.toString();

        // Split into words
        String[] words = line.split("\\s+");

        // Emit (word, 1) for each word
        for (String w : words) {
            word.set(w.toLowerCase());
            context.write(word, one);
        }
    }
}

// Reducer
public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        // Input: (word, [1, 1, 1, ...])
        int sum = 0;

        // Sum all occurrences
        for (IntWritable val : values) {
            sum += val.get();
        }

        // Emit (word, count)
        context.write(key, new IntWritable(sum));
    }
}

// Driver (Main)
public class WordCount {

```

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");

    job.setJarByClass(WordCount.class);
    job.setMapperClass(WordCountMapper.class);
    job.setReducerClass(WordCountReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}

// Run:
// hadoop jar wordcount.jar WordCount /input /output
```

MapReduce: More Complex Example

Problem: Calculate average order value per customer

python

```

from mrjob.job import MRJob
from mrjob.step import MRStep
import json

class CustomerAvgOrder(MRJob):

    def mapper(self, _, line):
        """
        Input: JSON line with order data
        Output: (customer_id, order_total)
        """
        order = json.loads(line)
        yield order['customer_id'], order['total']

    def reducer(self, customer_id, order_totals):
        """
        Input: (customer_id, [total1, total2, total3, ...])
        Output: (customer_id, average_order_value)
        """
        totals = list(order_totals)
        avg = sum(totals) / len(totals)

        yield customer_id, {
            'avg_order_value': avg,
            'total_orders': len(totals),
            'total_spent': sum(totals)
        }

    if __name__ == '__main__':
        CustomerAvgOrder.run()

# Input data:
# {"customer_id": "C1", "total": 100}
# {"customer_id": "C1", "total": 200}
# {"customer_id": "C2", "total": 50}
# {"customer_id": "C1", "total": 150}

# Map output:
# ("C1", 100), ("C1", 200), ("C2", 50), ("C1", 150)

# Shuffle & sort:
# ("C1", [100, 200, 150])
# ("C2", [50])

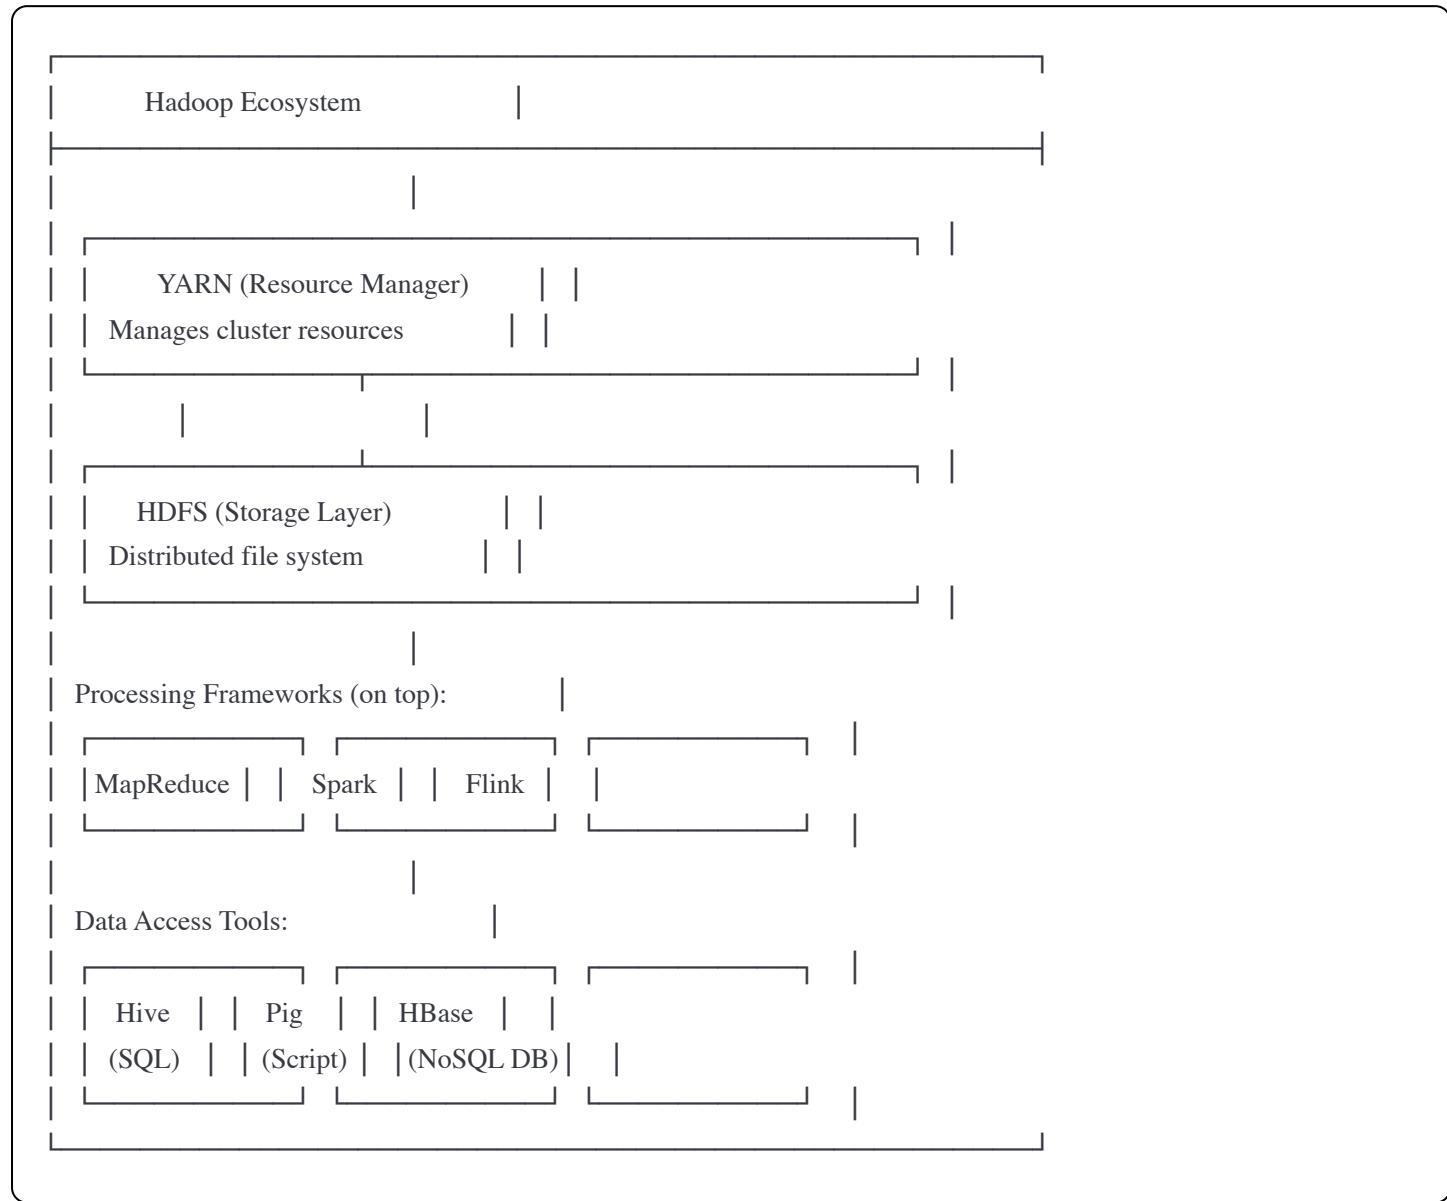
# Reduce output:

```

```
# C1: {"avg_order_value": 150, "total_orders": 3, "total_spent": 450}
# C2: {"avg_order_value": 50, "total_orders": 1, "total_spent": 50}
```

3. Apache Hadoop Ecosystem

Hadoop Architecture



HDFS (Hadoop Distributed File System)

Already covered in Chapter 8, quick recap:

HDFS Architecture:

NameNode (Master):

- └─ Manages metadata
- └─ Knows which blocks are where
- └─ Single point (but has backup)

DataNodes (Workers):

- └─ Store actual data blocks
- └─ 128 MB blocks
- └─ 3 replicas per block
- └─ Report health to NameNode

Example: Store 1 GB file

$1\text{ GB} / 128\text{ MB} = 8\text{ blocks}$

$8\text{ blocks} \times 3\text{ replicas} = 24\text{ total blocks}$

Distributed across DataNodes

Hive (SQL on Hadoop)

Concept: Write SQL, executes as MapReduce jobs.

sql

```

-- Create table
CREATE TABLE orders (
    order_id INT,
    customer_id INT,
    total DECIMAL(10,2),
    order_date DATE
)
STORED AS PARQUET
LOCATION '/data/orders';

-- Query (compiles to MapReduce)
SELECT
    customer_id,
    COUNT(*) as order_count,
    SUM(total) as total_spent,
    AVG(total) as avg_order
FROM orders
WHERE order_date >= '2024-01-01'
GROUP BY customer_id
HAVING SUM(total) > 1000
ORDER BY total_spent DESC
LIMIT 100;

-- Behind the scenes:
-- 1. Hive compiles SQL to MapReduce jobs
-- 2. Launches map tasks to scan data
-- 3. Shuffle groups by customer_id
-- 4. Reduce tasks aggregate
-- 5. Returns results

-- Can process petabytes!
-- Just write SQL!

```

4. Apache Spark

Why Spark? (Faster than MapReduce)

MapReduce Problem:

Every operation writes to disk!

Job 1: Read HDFS → Map → Write to HDFS

Job 2: Read HDFS → Reduce → Write to HDFS

Job 3: Read HDFS → Map → Write to HDFS

↑ Slow disk I/O between every step!

Spark Solution:

Keep data in memory!

Job: Read HDFS → Map → Reduce → Map → Output

↑ All in memory (RAM) ↑

100x faster for iterative jobs!

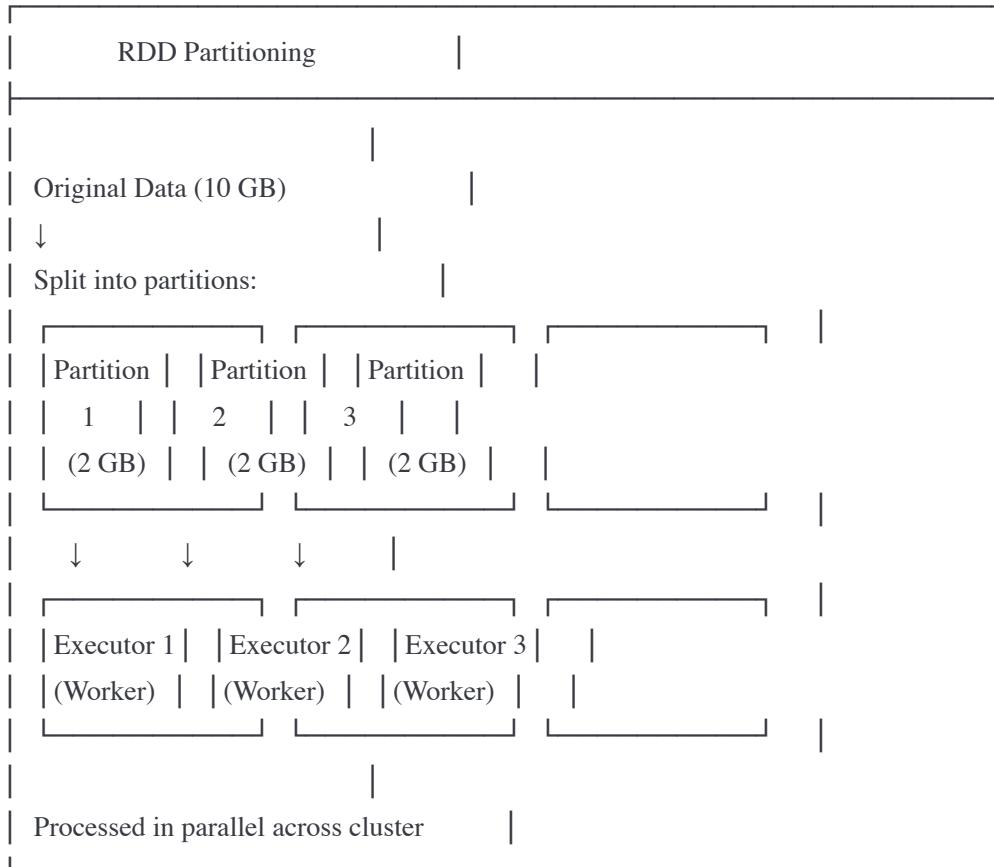
Performance Comparison:

MapReduce: 100 GB in 10 minutes (disk-bound)

Spark: 100 GB in 1 minute (memory-bound)

Spark RDD (Resilient Distributed Dataset)

RDD: Immutable distributed collection



Spark Code Examples

python

```
from pyspark.sql import SparkSession

# Initialize Spark
spark = SparkSession.builder \
    .appName("SparkExample") \
    .config("spark.executor.memory", "4g") \
    .getOrCreate()

# Example 1: Word Count (much simpler than Hadoop!)
text_file = spark.read.text("hdfs://data/input.txt")

word_counts = text_file \
    .rdd \
    .flatMap(lambda line: line.value.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)

word_counts.saveAsTextFile("hdfs://data/output")

# Example 2: DataFrame API (SQL-like)
df = spark.read.parquet("s3://data/orders/")

# Filter
high_value = df.filter(df.total > 1000)

# Aggregation
revenue_by_category = df \
    .groupBy("category") \
    .agg(
        sum("total").alias("revenue"),
        count("order_id").alias("order_count"),
        avg("total").alias("avg_order")
    )

revenue_by_category.show()

# Example 3: Join (across multiple datasets)
orders = spark.read.parquet("s3://data/orders/")
customers = spark.read.parquet("s3://data/customers/")

# Join orders with customer info
result = orders.join(
    customers,
    orders.customer_id == customers.id,
    "inner"
)
```

```
# Group by customer tier
tier_analysis = result \
    .groupBy(customers.tier) \
    .agg(
        sum(orders.total).alias("revenue"),
        count(orders.order_id).alias("orders")
    )
```

```
tier_analysis.show()
```

```
# Example 4: Window functions (complex analytics)
```

```
from pyspark.sql.window import Window
from pyspark.sql.functions import row_number, rank

# Rank products by revenue within each category
window_spec = Window.partitionBy("category").orderBy(desc("revenue"))
```

```
ranked = df \
    .withColumn("rank", rank().over(window_spec)) \
    .filter(col("rank") <= 10) # Top 10 per category
```

```
ranked.show()
```

```
# Example 5: Machine Learning (MLlib)
```

```
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.feature import VectorAssembler
```

```
# Prepare features
assembler = VectorAssembler(
    inputCols=["age", "income", "purchases"],
    outputCol="features"
)
```

```
data = assembler.transform(df)
```

```
# Train model
```

```
lr = LogisticRegression(
    featuresCol="features",
    labelCol="will_churn"
)
```

```
model = lr.fit(data)
```

```
# Predict
```

```
predictions = model.transform(test_data)
```

Spark Streaming (Micro-Batch)

python

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import window, col, count, avg

spark = SparkSession.builder \
    .appName("SparkStreaming") \
    .getOrCreate()

# Read from Kafka (streaming source)
stream = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("subscribe", "transactions") \
    .load()

# Parse JSON
from pyspark.sql.functions import from_json, col
from pyspark.sql.types import StructType, StringType, DoubleType

schema = StructType() \
    .add("transaction_id", StringType()) \
    .add("user_id", StringType()) \
    .add("amount", DoubleType()) \
    .add("timestamp", StringType())

transactions = stream \
    .selectExpr("CAST(value AS STRING)") \
    .select(from_json(col("value"), schema).alias("data")) \
    .select("data.*")

# Windowed aggregation (5-minute windows)
windowed_metrics = transactions \
    .withWatermark("timestamp", "10 minutes") \
    .groupBy(
        window(col("timestamp"), "5 minutes"),
        col("user_id")
    ) \
    .agg(
        count("transaction_id").alias("tx_count"),
        sum("amount").alias("total_amount"),
        avg("amount").alias("avg_amount")
    )

# Write to console (for demo)
query = windowed_metrics \
    .writeStream \

```

```

.outputMode("update") \
.format("console") \
.start()

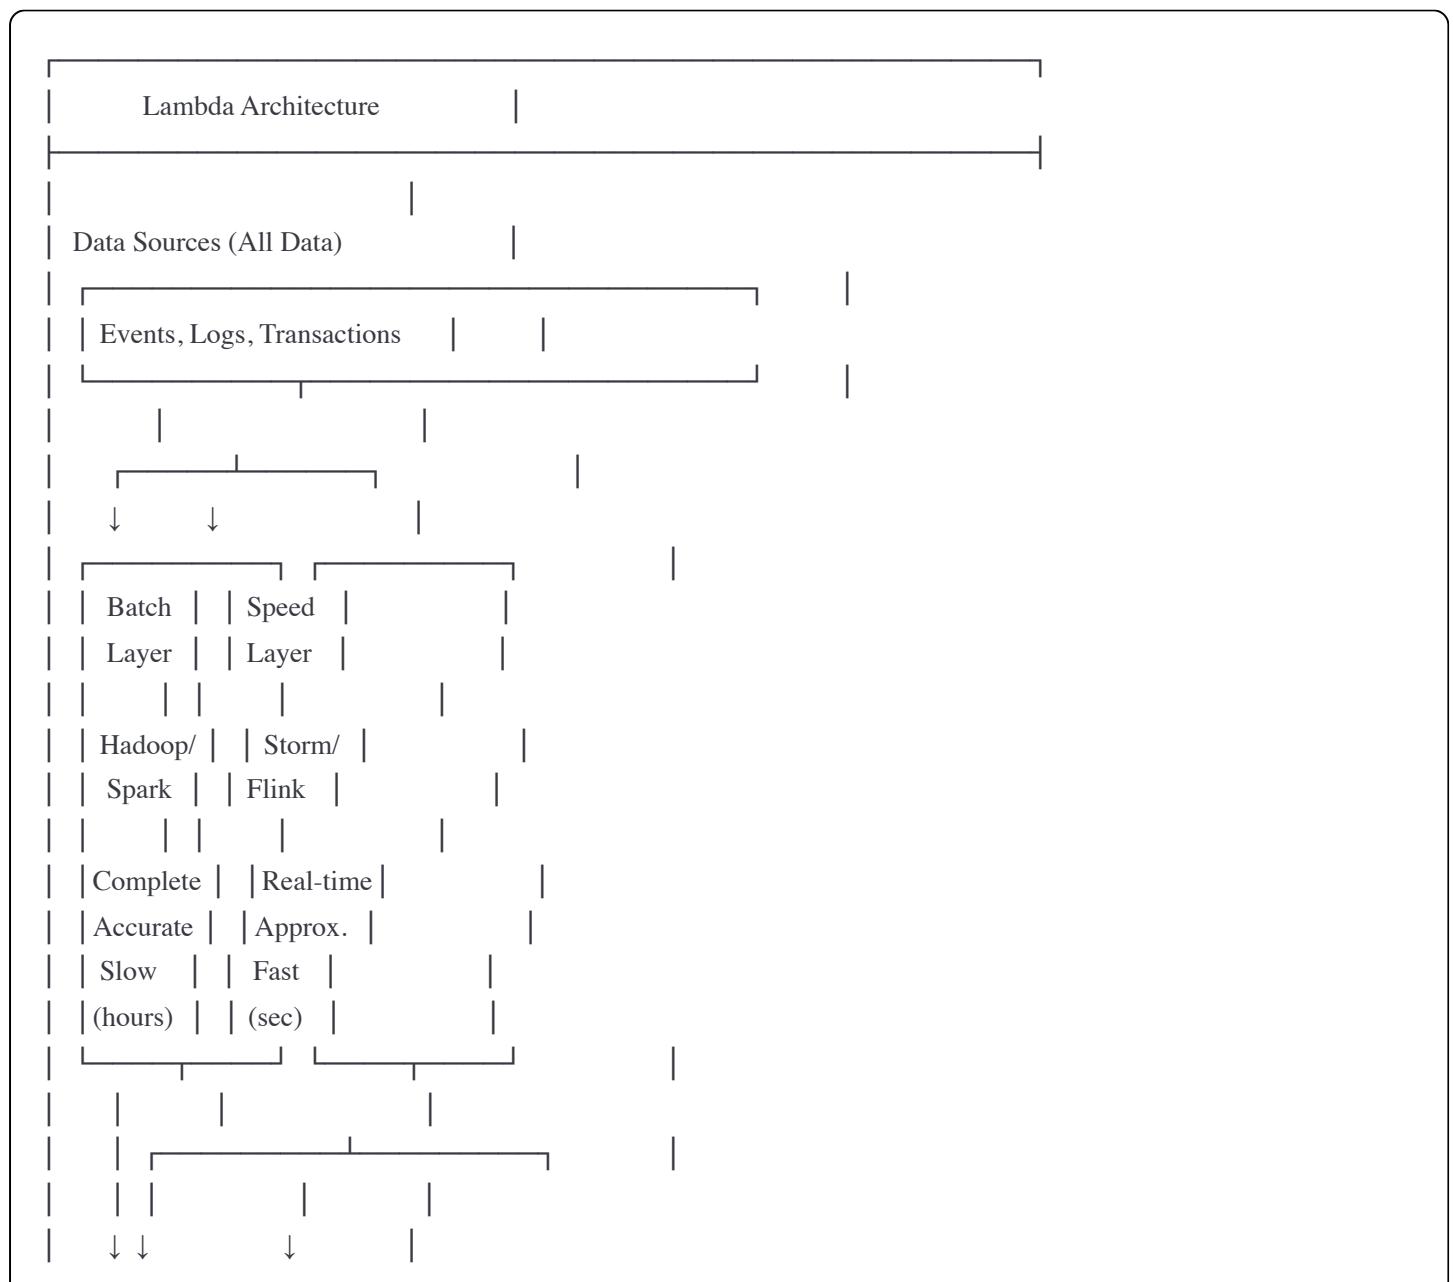
query.awaitTermination()

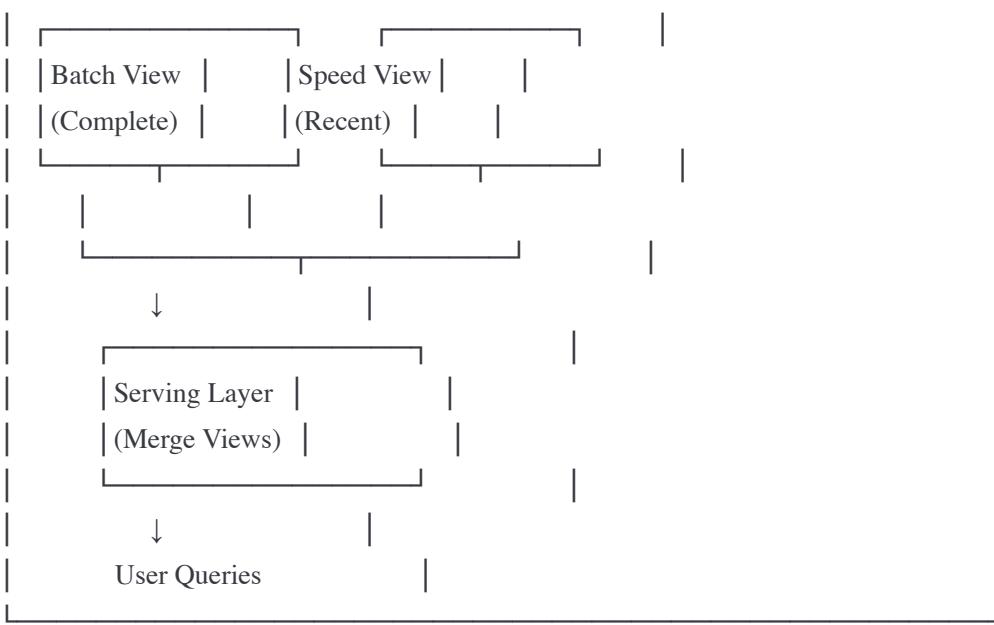
# Output every 5 minutes:
# +-----+-----+-----+-----+
# |window| user_id | tx_count | total_amount | avg_amount |
# +-----+-----+-----+-----+
# | 2024-01-20 10:00,... | user123 | 45 | 4532.50 | 100.72 |
# | 2024-01-20 10:00,... | user456 | 23 | 1234.00 | 53.65 |
# +-----+-----+-----+-----+

```

5. Lambda Architecture

Concept: Batch + Stream = Best of Both Worlds





Example: Page View Analytics

Batch Layer:

- Processes all historical data
- Runs daily at midnight
- Complete, accurate counts
- Available at 2 AM

Speed Layer:

- Processes events in real-time
- Updates every second
- Approximate counts
- Available immediately

Serving Layer:

- Query at 10 AM:
 - Batch view: Accurate until midnight (10 hours old)
 - Speed view: Last 10 hours (approximate)
 - Combined: Complete picture!

Lambda Architecture Implementation

python

```

# BATCH LAYER (Spark)
from pyspark.sql import SparkSession

def batch_processing():
    spark = SparkSession.builder.appName("BatchLayer").getOrCreate()

    # Process all historical data
    all_page_views = spark.read.parquet("s3://data/page-views/")

    # Calculate complete statistics
    page_stats = all_page_views \
        .groupBy("page_id", "date") \
        .agg(
            count("user_id").alias("total_views"),
            countDistinct("user_id").alias("unique_visitors")
        )

    # Save to serving layer
    page_stats.write \
        .mode("overwrite") \
        .parquet("s3://serving/batch-views/")

    print("Batch layer updated (complete and accurate)")

# SPEED LAYER (Kafka Streams)
from kafka import KafkaConsumer
import json

def speed_layer():
    consumer = KafkaConsumer(
        'page-views',
        bootstrap_servers=['kafka:9092']
    )

    # Real-time aggregation
    recent_views = {} # page_id -> count

    for message in consumer:
        event = json.loads(message.value)
        page_id = event['page_id']

        # Update count
        recent_views[page_id] = recent_views.get(page_id, 0) + 1

    # Store in fast database (Redis)
    redis_client.hincrby('page_views_realtime', page_id, 1)

```

```

# Publish to serving layer
if recent_views[page_id] % 100 == 0:
    print(f"Speed layer: Page {page_id} has {recent_views[page_id]} recent views")

# SERVING LAYER (API)
from flask import Flask, jsonify
import redis

app = Flask(__name__)
redis_client = redis.Redis(host='localhost', port=6379)
spark = SparkSession.builder.appName("ServingLayer").getOrCreate()

@app.route('/api/page-stats/<page_id>')
def get_page_stats(page_id):
    # Get batch view (complete historical data)
    batch_df = spark.read.parquet("s3://serving/batch-views/")
    batch_stats = batch_df \
        .filter(batch_df.page_id == page_id) \
        .agg(sum("total_views").alias("historical_views")) \
        .collect()

    historical_views = batch_stats[0]['historical_views'] if batch_stats else 0

    # Get speed view (recent data)
    recent_views = int(redis_client.hget('page_views_realtime', page_id) or 0)

    # Combine both
    total_views = historical_views + recent_views

    return jsonify({
        'page_id': page_id,
        'total_views': total_views,
        'batch_views': historical_views, # Complete, accurate
        'speed_views': recent_views, # Recent, approximate
        'batch_updated': '2024-01-20 02:00:00',
        'speed_updated': 'real-time'
    })

app.run()

# Example response:
# {
#   "page_id": "page-123",
#   "total_views": 15234,
#   "batch_views": 15000, (until midnight)
#   "speed_views": 234, (since midnight)

```

```
# "batch_updated": "2024-01-20 02:00:00",
# "speed_updated": "real-time"
# }
```

Lambda Architecture Pros and Cons

✓ ADVANTAGES:

- Complete and accurate (batch layer)
- Real-time updates (speed layer)
- Fault tolerant (can recompute from batch)
- Handle both historical and recent data

✗ DISADVANTAGES:

- Complex (two processing pipelines!)
- Code duplication (same logic in batch and speed)
- Operational overhead (maintain two systems)
- Eventual consistency (batch lags behind speed)

When to use:

- ✓ Need both accuracy and real-time
- ✓ Historical analytics + live dashboards
- ✓ Can handle complexity
- ✓ Have resources for two systems

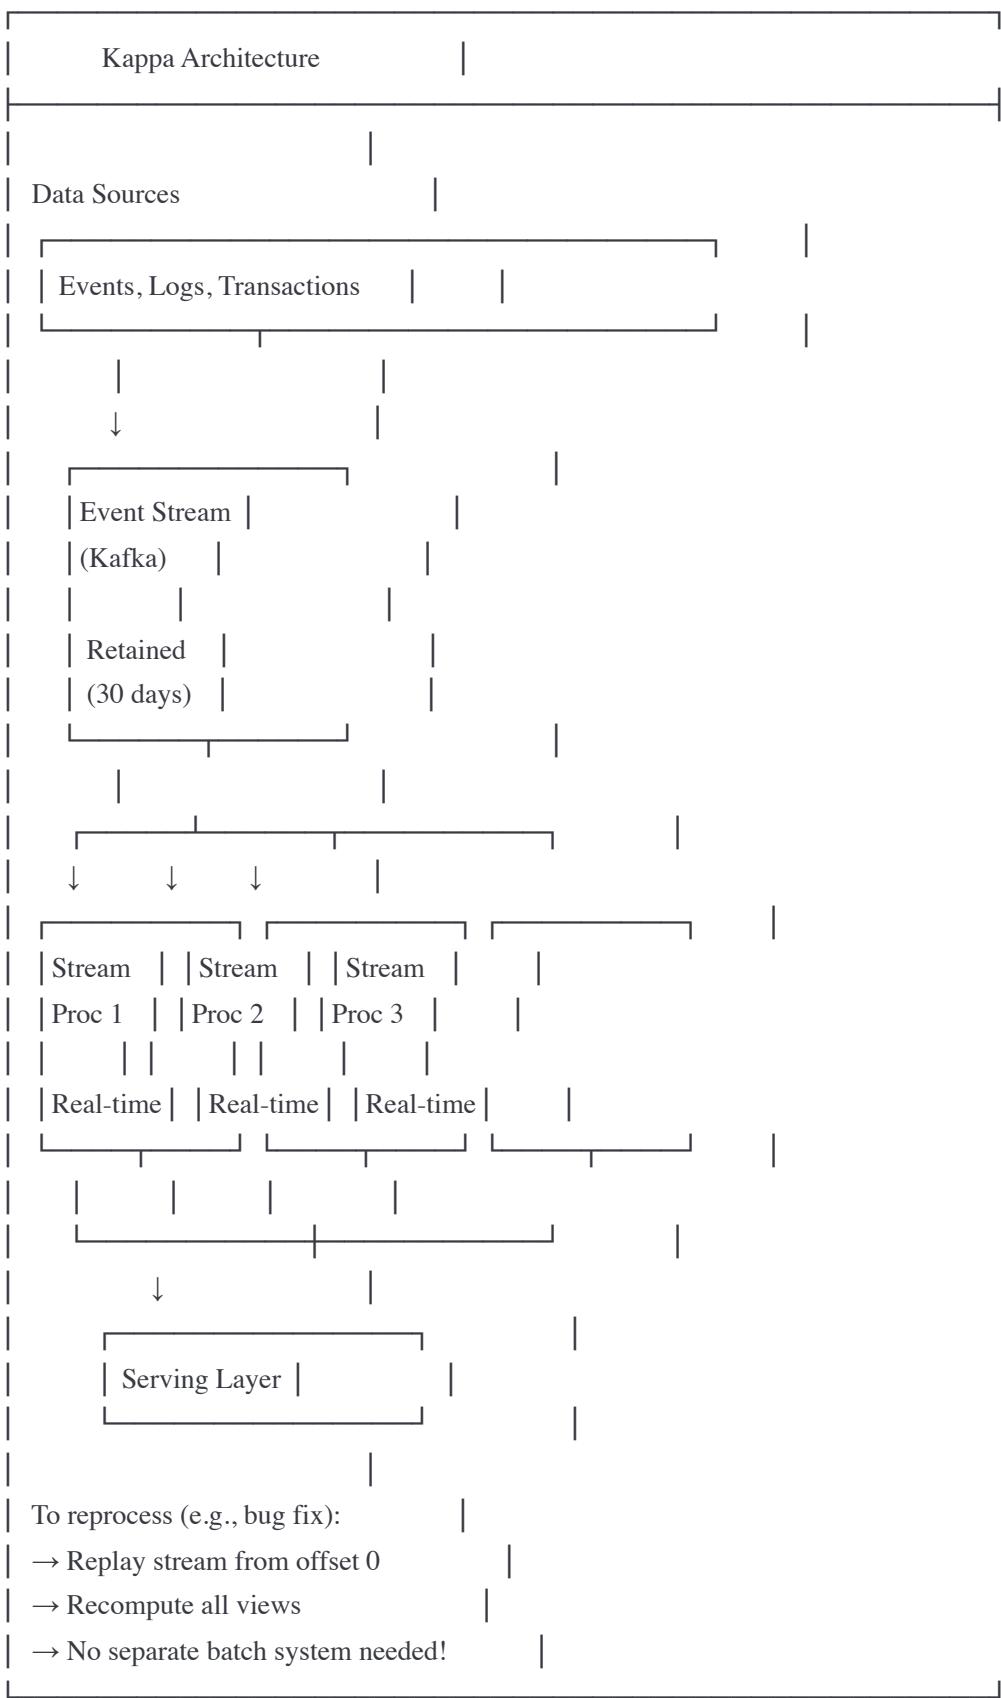
Examples:

- LinkedIn (uses Lambda for analytics)
- Twitter (for trending topics, analytics)

6. Kappa Architecture

Simplified Alternative to Lambda

Concept: Everything is a stream! No batch layer.



Key Difference from Lambda:

- Only ONE processing code path (stream)
- Reprocess by replaying stream
- Simpler!

Kappa Implementation

javascript

```
const { Kafka } = require('kafkajs');

class KappaProcessor {
  constructor() {
    this.kafka = new Kafka({ brokers: ['kafka:9092'] });
    this.consumer = this.kafka.consumer({ groupId: 'kappa-processor' });
    this.producer = this.kafka.producer();
  }

  async processStream() {
    await this.consumer.connect();
    await this.producer.connect();

    await this.consumer.subscribe({
      topic: 'page-views',
      fromBeginning: false // Or true to reprocess all
    });

    // State (in-memory or database)
    const state = {
      pageViews: new Map(), // page_id -> count
      userSessions: new Map(), // user_id -> session
      hourlyMetrics: new Map() // hour -> metrics
    };

    await this.consumer.run({
      eachMessage: async ({ message }) => {
        const event = JSON.parse(message.value);

        // Process event (update all views)
        this.updatePageViews(state, event);
        this.updateSessions(state, event);
        this.updateHourlyMetrics(state, event);

        // Publish derived events
        await this.publishDerivedEvents(event, state);
      }
    });

    updatePageViews(state, event) {
      const pageId = event.page_id;
      const current = state.pageViews.get(pageId) || 0;
      state.pageViews.set(pageId, current + 1);

      // Persist to database
    }
  }
}
```

```
db.update('page_views', { page_id: pageId }, { count: current + 1 });

}

updateSessions(state, event) {
  const userId = event.user_id;
  const now = new Date(event.timestamp);

  let session = state.userSessions.get(userId);

  if (!session || now - session.lastActivity > 30 * 60 * 1000) {
    // New session
    session = {
      sessionId: generateId(),
      startTime: now,
      pageViews: 0
    };
    state.userSessions.set(userId, session);
  }

  session.pageViews++;
  session.lastActivity = now;
}

updateHourlyMetrics(state, event) {
  const hour = new Date(event.timestamp).getHours();
  const key = `${event.date}_${hour}`;

  if (!state.hourlyMetrics.has(key)) {
    state.hourlyMetrics.set(key, {
      views: 0,
      uniqueUsers: new Set()
    });
  }

  const metrics = state.hourlyMetrics.get(key);
  metrics.views++;
  metrics.uniqueUsers.add(event.user_id);
}

async publishDerivedEvents(event, state) {
  // Publish aggregated metrics to output topic
  const pageViews = state.pageViews.get(event.page_id);

  if (pageViews % 1000 === 0) {
    await this.producer.send({
      topic: 'page-view-milestones',
      messages: [{
```

```

    value: JSON.stringify({
      page_id: event.page_id,
      milestone: pageViews,
      timestamp: new Date().toISOString()
    })
  }]
});
}
}

// Reprocess from beginning (e.g., after bug fix)
async reprocess() {
  console.log('Starting reprocessing from beginning...');

  // Create new consumer group
  const reprocessConsumer = this.kafka.consumer({
    groupId: `kappa-reprocess-${Date.now()}`,
  });

  await reprocessConsumer.connect();
  await reprocessConsumer.subscribe({
    topic: 'page-views',
    fromBeginning: true // Replay entire stream!
  });

  // Process all events from scratch
  // Rebuild all views
  await reprocessConsumer.run({
    eachMessage: async ({ message }) => {
      // Same processing logic
      // Rebuilds state from scratch
    }
  });

  console.log('Reprocessing complete');
}

const processor = new KappaProcessor();
processor.processStream();

```

Lambda vs Kappa

Feature		Lambda		Kappa		
---------	--	--------	--	-------	--	--

Layers	Batch + Speed	Stream only
Code	Duplicate	Single codebase
Complexity	High	Lower
Latency	Mixed	Consistent
Accuracy	Batch=perfect	Stream=good
Reprocessing	Batch layer	Replay stream
Operational	Complex	Simpler

When to use Lambda:

- ✓ Need perfect accuracy (batch)
- ✓ Complex batch analytics
- ✓ Different processing for batch vs real-time

When to use Kappa:

- ✓ All processing is stream-based
- ✓ Can replay stream (Kafka retention)
- ✓ Want simpler architecture
- ✓ Real-time is good enough

Modern Trend: Moving toward Kappa

Reason: Spark Streaming bridges the gap

Can do both batch and stream with same code!

7. Real-Time Analytics

Use Cases

1. FRAUD DETECTION

Process: Transaction → Analyze → Alert (< 100ms)

Pattern: Compare to user history, geo-location

2. RECOMMENDATION

Process: Click → Update model → Show recommendations

Pattern: Collaborative filtering in real-time

3. MONITORING

Process: Metric → Aggregate → Alert (< 1 second)

Pattern: Window aggregation, anomaly detection

4. REAL-TIME BIDDING (Ad Tech)

Process: Page load → Auction → Ad served (< 100ms)

Pattern: Complex matching in real-time

5. LIVE DASHBOARDS

Process: Event → Aggregate → Update UI (< 5 seconds)

Pattern: Windowed aggregation, streaming joins

Real-Time Analytics Pipeline

javascript

```

// Complete real-time analytics system

class RealTimeAnalytics {
  constructor() {
    this.kafka = new Kafka({ brokers: ['kafka:9092'] });
    this.redis = new Redis({ host: 'redis' });
    this.websockets = new Map(); // Active dashboard connections
  }

  // Stage 1: Collect events
  async collectEvents() {
    const producer = this.kafka.producer();
    await producer.connect();

    // Simulate events
    setInterval(async () => {
      const event = {
        type: 'page_view',
        userId: `user-${Math.floor(Math.random() * 1000)}`,
        pageId: `page-${Math.floor(Math.random() * 100)}`,
        timestamp: new Date().toISOString(),
        duration: Math.floor(Math.random() * 60000)
      };

      await producer.send({
        topic: 'events',
        messages: [{ value: JSON.stringify(event) }]
      });
    }, 10); // 100 events/second
  }

  // Stage 2: Real-time aggregation
  async aggregateRealTime() {
    const consumer = this.kafka.consumer({ groupId: 'aggregator' });
    await consumer.connect();
    await consumer.subscribe({ topic: 'events' });

    // Tumbling windows (1 minute)
    const windows = new Map();
    const WINDOW_SIZE = 60000; // 1 minute

    await consumer.run({
      eachMessage: async ({ message }) => {
        const event = JSON.parse(message.value);
        const timestamp = new Date(event.timestamp).getTime();
        ...
      }
    });
  }
}

```

```

// Determine window
const windowStart = Math.floor(timestamp / WINDOW_SIZE) * WINDOW_SIZE;
const windowKey = `${event.pageId}:${windowStart}`;

// Update window
if (!windows.has(windowKey)) {
  windows.set(windowKey, {
    pageId: event.pageId,
    windowStart: new Date(windowStart).toISOString(),
    viewCount: 0,
    uniqueUsers: new Set(),
    totalDuration: 0
  });
}

const window = windows.get(windowKey);
window.viewCount++;
window.uniqueUsers.add(event.userId);
window.totalDuration += event.duration;

// Publish window results to Redis
await this.redis.setex(
  `window:${windowKey}`,
  300, // 5 minute TTL
  JSON.stringify({
    pageId: window.pageId,
    windowStart: window.windowStart,
    viewCount: window.viewCount,
    uniqueUsers: window.uniqueUsers.size,
    avgDuration: window.totalDuration / window.viewCount
  })
);

// Cleanup old windows
const now = Date.now();
for (const [key, win] of windows) {
  const winStart = new Date(win.windowStart).getTime();
  if (now - winStart > 3600000) { // 1 hour old
    windows.delete(key);
  }
}
};

// Stage 3: Real-time dashboards
async streamToDashboard() {

```

```

const consumer = this.kafka.consumer({ groupId: 'dashboard' });
await consumer.connect();
await consumer.subscribe({ topic: 'events' });

// Track metrics
let eventCount = 0;
let lastUpdate = Date.now();

await consumer.run({
  eachMessage: async ({ message }) => {
    eventCount++;

    // Update dashboard every second
    const now = Date.now();
    if (now - lastUpdate > 1000) {
      const metrics = {
        eventsPerSecond: eventCount,
        timestamp: new Date().toISOString()
      };
      this.broadcastToDashboards(metrics);
      eventCount = 0;
      lastUpdate = now;
    }
  }
});

// Stage 4: Anomaly detection
async detectAnomalies() {
  const consumer = this.kafka.consumer({ groupId: 'anomaly-detector' });
  await consumer.connect();
  await consumer.subscribe({ topic: 'events' });

  // Baseline metrics
  const baselines = new Map();

  await consumer.run({
    eachMessage: async ({ message }) => {
      const event = JSON.parse(message.value);

      // Get baseline for this page
      if (!baselines.has(event.pageId)) {
        // Initialize from historical data
        const historical = await this.getHistoricalAverage(event.pageId);
        baselines.set(event.pageId, historical);
      }
    }
  });
}

```

```
baselines.set(event.pageId, historical);
}

const baseline = baselines.get(event.pageId);

// Get current rate (from Redis)
const currentRate = await this.getCurrentRate(event.pageId);

// Detect spike (3x baseline)
if (currentRate > baseline * 3) {
  console.log(`⚠️ Anomaly: Page ${event.pageId} has ${currentRate} views/min (baseline: ${baseline})`);

  await this.alertOps({
    type: 'traffic_spike',
    pageId: event.pageId,
    expected: baseline,
    actual: currentRate,
    severity: currentRate > baseline * 10 ? 'critical' : 'warning'
  });
}

// Update baseline (moving average)
baselines.set(
  event.pageId,
  baseline * 0.95 + currentRate * 0.05
);
}

});

}

async getCurrentRate(pageId) {
  // Get recent window from Redis
  const keys = await this.redis.keys(`window:${pageId}:*`);

  if (keys.length === 0) return 0;

  let totalViews = 0;
  for (const key of keys) {
    const data = JSON.parse(await this.redis.get(key));
    totalViews += data.viewCount;
  }

  return totalViews / keys.length; // Average per window
}

async getHistoricalAverage(pageId) {
  // Query batch data for baseline
```

```

// Simplified: return reasonable default
return 100; // 100 views/minute baseline
}

broadcastToDashboards(metrics) {
    // Send to all WebSocket clients
    this.websockets.forEach(ws => {
        ws.send(JSON.stringify(metrics));
    });
}

async alertOps(alert) {
    console.log('Sending alert:', alert);
    // Send to PagerDuty, Slack, etc.
}
}

// Start all stages
const analytics = new RealTimeAnalytics();
analytics.collectEvents();
analytics.aggregateRealTime();
analytics.streamToDashboard();
analytics.detectAnomalies();

```

Real-Time Analytics Patterns

Pattern 1: Windowed Aggregation

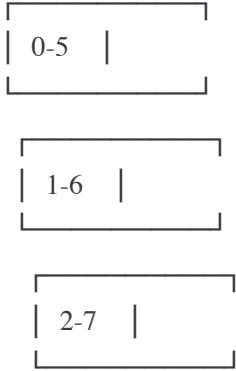
Time-based windows:

Tumbling Window (non-overlapping):



Each event in exactly one window

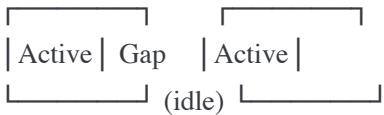
Sliding Window (overlapping):



Windows slide by 1 minute

Event can be in multiple windows

Session Window (activity-based):



Window ends after 30 min inactivity

Implementation:

javascript

```

class WindowedAggregation {
  constructor(windowSizeMs, slideMs = null) {
    this.windowSize = windowSizeMs;
    this.slide = slideMs || windowSizeMs; // Default: tumbling
    this.windows = new Map();
  }

  processEvent(event) {
    const timestamp = new Date(event.timestamp).getTime();

    // Find all windows this event belongs to
    const windowStarts = this.getWindowStarts(timestamp);

    windowStarts.forEach(windowStart => {
      const windowKey = `${event.key}:${windowStart}`;

      if (!this.windows.has(windowKey)) {
        this.windows.set(windowKey, {
          key: event.key,
          windowStart,
          windowEnd: windowStart + this.windowSize,
          count: 0,
          sum: 0,
          events: []
        });
      }

      const window = this.windows.get(windowKey);
      window.count++;
      window.sum += event.value;
      window.events.push(event);
    });

    // Emit completed windows
    this.emitCompletedWindows(timestamp);
  }

  getWindowStarts(timestamp) {
    if (this.slide === this.windowSize) {
      // Tumbling window
      const windowStart = Math.floor(timestamp / this.windowSize) * this.windowSize;
      return [windowStart];
    } else {
      // Sliding window
      const starts = [];
      let start = Math.floor(timestamp / this.slide) * this.slide;
      for (let i = 0; i < this.windowSize; i++) {
        starts.push(start);
        start += this.slide;
      }
    }
  }
}

```

```
while (start + this.windowSize > timestamp) {
  starts.push(start);
  start -= this.slide;
}

return starts;
}

emitCompletedWindows(currentTimestamp) {
  const completed = [];

  for (const [windowKey, window] of this.windows) {
    // Window is complete if current time > window end
    if (currentTimestamp > window.windowEnd) {
      completed.push(windowKey);

      // Emit result
      console.log('Window completed:', {
        key: window.key,
        start: new Date(window.windowStart).toISOString(),
        end: new Date(window.windowEnd).toISOString(),
        count: window.count,
        average: window.sum / window.count
      });
    }
  }

  // Cleanup completed windows
  completed.forEach(key => this.windows.delete(key));
}

}

// Usage
const aggregator = new WindowedAggregation(
  5 * 60 * 1000, // 5-minute windows
  1 * 60 * 1000 // Slide by 1 minute (overlapping)
);

// Process events
aggregator.processEvent({
  key: 'page-123',
  value: 1,
  timestamp: '2024-01-20T10:00:00Z'
});
```

```
aggregator.processEvent({  
    key: 'page-123',  
    value: 1,  
    timestamp: '2024-01-20T10:01:00Z'  
});  
  
// Event at 10:00:00 is in windows:  
// - [09:56:00 - 10:01:00]  
// - [09:57:00 - 10:02:00]  
// - [09:58:00 - 10:03:00]  
// - [09:59:00 - 10:04:00]  
// - [10:00:00 - 10:05:00]
```

Pattern 2: Stream Joins

Joining two event streams:

Stream A (Orders):	Stream B (Payments):
Order 1 (10:00:00)	Payment 1 (10:00:05)
Order 2 (10:00:10)	Payment 2 (10:00:15)
Order 3 (10:00:20)	Payment 3 (10:00:25)

Goal: Match orders with payments

Challenges:

- Events arrive out of order
- Time difference between order and payment
- Need to buffer events

Solution: Time-windowed join

Wait up to 1 minute for matching event

Implementation:

javascript

```
class StreamJoiner {
  constructor() {
    this.orderBuffer = new Map();
    this.paymentBuffer = new Map();
    this.maxWaitTime = 60000; // 1 minute
  }

  async processOrder(order) {
    const orderId = order.orderId;

    // Check if payment already arrived
    if (this.paymentBuffer.has(orderId)) {
      const payment = this.paymentBuffer.get(orderId);
      this.paymentBuffer.delete(orderId);

      // Emit joined event
      this.emitJoined(order, payment);
      return;
    }

    // Buffer order, wait for payment
    this.orderBuffer.set(orderId, {
      order,
      arrivedAt: Date.now()
    });

    // Set timeout to cleanup
    setTimeout(() => {
      if (this.orderBuffer.has(orderId)) {
        console.log(`Order ${orderId} timeout (no payment received)`);
        this.orderBuffer.delete(orderId);
      }
    }, this.maxWaitTime);
  }

  async processPayment(payment) {
    const orderId = payment.orderId;

    // Check if order already arrived
    if (this.orderBuffer.has(orderId)) {
      const { order } = this.orderBuffer.get(orderId);
      this.orderBuffer.delete(orderId);

      // Emit joined event
      this.emitJoined(order, payment);
      return;
    }
  }
}
```

```

}

// Buffer payment, wait for order
this.paymentBuffer.set(orderId, {
  payment,
  arrivedAt: Date.now()
});

setTimeout(() => {
  if (this.paymentBuffer.has(orderId)) {
    console.log(`Payment ${orderId} timeout (no order received)`);
    this.paymentBuffer.delete(orderId);
  }
}, this.maxWaitTime);
}

emitJoined(order, payment) {
  const joined = {
    orderId: order.orderId,
    userId: order.userId,
    orderTotal: order.total,
    paymentAmount: payment.amount,
    paymentStatus: payment.status,
    orderTimestamp: order.timestamp,
    paymentTimestamp: payment.timestamp,
    timeDiff: new Date(payment.timestamp) - new Date(order.timestamp)
  };
  console.log('Joined event:', joined);

  // Publish to output stream
  kafka.publish('order-payment-joined', joined);
}
}

const joiner = new StreamJoiner();

// Process streams
kafka.subscribe('orders', (order) => joiner.processOrder(order));
kafka.subscribe('payments', (payment) => joiner.processPayment(payment));

```

Pattern 3: Real-Time Dashboard

javascript

```
// Live dashboard with WebSocket updates

const WebSocket = require('ws');
const wss = new WebSocket.Server({ port: 8080 });

class LiveDashboard {
  constructor() {
    this.kafka = new Kafka({ brokers: ['kafka:9092'] });
    this.currentMetrics = {
      eventsPerSecond: 0,
      activeUsers: new Set(),
      topPages: [],
      errorRate: 0
    };

    this.clients = new Set();
  }

  async start() {
    // Accept WebSocket connections
    wss.on('connection', (ws) => {
      console.log('Dashboard client connected');
      this.clients.add(ws);

      // Send current metrics immediately
      ws.send(JSON.stringify(this.currentMetrics));

      ws.on('close', () => {
        this.clients.delete(ws);
      });
    });
  }

  // Process events
  const consumer = this.kafka.consumer({ groupId: 'dashboard' });
  await consumer.connect();
  await consumer.subscribe({ topic: 'events' });

  let eventCount = 0;
  let lastUpdate = Date.now();

  await consumer.run({
    eachMessage: async ({ message }) => {
      const event = JSON.parse(message.value);
      eventCount++;
    }
  });

  // Update metrics
}
```

```

this.currentMetrics.activeUsers.add(event.userId);

// Every second, calculate and broadcast
const now = Date.now();
if (now - lastUpdate > 1000) {
  this.currentMetrics.eventsPerSecond = eventCount;
  this.currentMetrics.activeUsersCount = this.currentMetrics.activeUsers.size;

  // Broadcast to all connected clients
  this.broadcast(this.currentMetrics);

  // Reset
  eventCount = 0;
  this.currentMetrics.activeUsers.clear();
  lastUpdate = now;
}

});

}

};

}

broadcast(data) {
  const message = JSON.stringify(data);

  this.clients.forEach(client => {
    if (client.readyState === WebSocket.OPEN) {
      client.send(message);
    }
  });
}

const dashboard = new LiveDashboard();
dashboard.start();

// Browser client:
// const ws = new WebSocket('ws://localhost:8080');
// ws.onmessage = (event) => {
//   const metrics = JSON.parse(event.data);
//   updateDashboard(metrics);
// };

```

Key Takeaways

1. Batch vs Stream:

- Batch: Large volumes, periodic, high latency

- Stream: Continuous, real-time, low latency
- Use both in Lambda architecture

2. MapReduce:

- Map: Transform data in parallel
- Reduce: Aggregate results
- Shuffle: Group by key
- Foundation of big data processing

3. Hadoop Ecosystem:

- HDFS: Distributed storage
- YARN: Resource management
- MapReduce: Processing framework
- Hive: SQL on Hadoop

4. Apache Spark:

- In-memory processing (100x faster)
- Unified API (batch + stream)
- RDD, DataFrame, SQL
- Machine learning (MLlib)

5. Lambda Architecture:

- Batch layer: Complete, accurate
- Speed layer: Real-time, approximate
- Serving layer: Merge both
- Complex but comprehensive

6. Kappa Architecture:

- Stream processing only
- Simpler than Lambda
- Replay for reprocessing
- Modern approach

7. Real-Time Analytics:

- Windowed aggregation
- Stream joins

- Anomaly detection
- Live dashboards

Practice Problems

1. Design a real-time analytics system for Twitter (trending topics, user engagement)
2. Calculate: 1 PB of data, 1000 machines, process at 100 GB/sec. How long to process?
3. Compare Lambda vs Kappa for Netflix viewing analytics. Which would you choose?
4. Design a fraud detection system processing 100,000 transactions/second

Next Chapter Preview

Chapter 18: Search Systems

- Full-text search
- Inverted indexes
- Elasticsearch
- Ranking and relevance
- Search optimization

Ready to continue with more chapters or practice system design problems?