


Chapter 19: Distributed Caching and Data Stores


Introduction: The Challenge of Distributed Caching

When you scale horizontally, caching becomes complex.

Single Server (Simple):



Multiple Servers (Complex):



Problems:

1. Cache inconsistency (each server has different data)
2. Cache invalidation (how to update all caches?)
3. Memory waste (same data cached 3 times)
4. Hot spots (one server gets all traffic for popular item)

1. Distributed Cache Coherence

The Consistency Problem

Scenario: User updates profile

Time Server 1 Cache Server 2 Cache Database

10:00 name: "John" name: "John" name: "John"
10:01 (User updates via Server 1)
name: "John Doe" name: "John" X name: "John Doe"
Stale data!


10:02 User visits site, hits Server 2
Sees old name: "John" X

This is CACHE COHERENCE problem!

Solution 1: Centralized Cache (Shared Cache)

Use Redis/Memcached as shared cache.

All Servers Share One Cache:



Benefits:

- ✓ Single source of truth
- ✓ No inconsistency
- ✓ Memory efficient
- ✓ Easy invalidation

Drawbacks:

- X Network latency (vs local cache)
- X Redis becomes bottleneck
- X Single point of failure (mitigated with clustering)

Implementation:

javascript

```
const redis = require('redis');
const client = redis.createClient({
  host: 'redis-cluster',
  port: 6379
});

class SharedCacheService {
  constructor(redisClient, database) {
    this.cache = redisClient;
    this.db = database;
  }

  async get(key) {
    // Try cache first
    const cached = await this.cache.get(key);

    if (cached) {
      console.log(`Cache HIT: ${key}`);
      return JSON.parse(cached);
    }

    console.log(`Cache MISS: ${key}`);

    // Fetch from database
    const data = await this.db.get(key);

    if (data) {
      // Store in cache (all servers share this cache)
      await this.cache.setex(key, 3600, JSON.stringify(data));
    }

    return data;
  }

  async set(key, value) {
    // Update database
    await this.db.set(key, value);

    // Update cache (visible to all servers immediately)
    await this.cache.setex(key, 3600, JSON.stringify(value));

    console.log(`Updated cache and DB: ${key}`);
  }

  async invalidate(key) {
    // Delete from cache (affects all servers)
  }
}
```

```
await this.cache.del(key);
console.log(`Invalidated cache: ${key}`);
}

}


// Usage across multiple servers
// Server 1
const cache1 = new SharedCacheService(client, db);
await cache1.set('user:123', { name: 'John Doe' });

// Server 2 (sees the same cache)
const cache2 = new SharedCacheService(client, db);
const user = await cache2.get('user:123'); // ✓ Gets updated value!
```


Solution 2: Cache Invalidation Messages

Broadcast invalidation to all servers.


User updates profile on Server 1:



| Publish to message bus



| Subscribe



✓ Delete Delete
 user:123 user:123
 from cache from cache

All servers now have consistent state!

Implementation:

javascript

```
const redis = require('redis');
const EventEmitter = require('events');

class DistributedCacheWithInvalidation extends EventEmitter {
  constructor(serverId) {
    super();
    this.serverId = serverId;
    this.localCache = new Map();

    // Redis for pub/sub
    this.publisher = redis.createClient();
    this.subscriber = redis.createClient();

    // Subscribe to invalidation messages
    this.subscriber.subscribe('cache-invalidation');

    this.subscriber.on('message', (channel, message) => {
      if (channel === 'cache-invalidation') {
        const { key, sourceServer } = JSON.parse(message);

        // Don't invalidate if we sent the message
        if (sourceServer !== this.serverId) {
          console.log(`[${this.serverId}] Received invalidation for ${key}`);
          this.localCache.delete(key);
        }
      }
    });
  }

  async get(key) {
    // Check local cache
    if (this.localCache.has(key)) {
      console.log(`[${this.serverId}] Local cache HIT: ${key}`);
      return this.localCache.get(key);
    }

    console.log(`[${this.serverId}] Local cache MISS: ${key}`);

    // Fetch from database
    const data = await database.get(key);

    if (data) {
      this.localCache.set(key, data);
    }

    return data;
  }
}
```

```
}

async set(key, value) {
    // Update database
    await database.set(key, value);

    // Update local cache
    this.localCache.set(key, value);

    // Broadcast invalidation to other servers
    await this.publisher.publish('cache-invalidation', JSON.stringify({
        key: key,
        sourceServer: this.serverId,
        timestamp: Date.now()
    }));
}

console.log(`[${this.serverId}] Updated and broadcasted invalidation`);

}

async invalidate(key) {
    // Delete from local cache
    this.localCache.delete(key);

    // Broadcast to other servers
    await this.publisher.publish('cache-invalidation', JSON.stringify({
        key: key,
        sourceServer: this.serverId,
        timestamp: Date.now()
    }));
}

}

// Usage on different servers
const server1Cache = new DistributedCacheWithInvalidation('server-1');
const server2Cache = new DistributedCacheWithInvalidation('server-2');
const server3Cache = new DistributedCacheWithInvalidation('server-3');

// Server 1 updates user
await server1Cache.set('user:123', { name: 'John Doe' });
// → Broadcasts invalidation
// → Server 2 and 3 delete from their local caches


// Server 2 reads user (after invalidation)
const user = await server2Cache.get('user:123');
// → Local cache empty (invalidated)
```

// → Fetches from database
// → Gets updated value ✓

Solution 3: Consistent Hashing (Distribute Keys)

Each key always goes to same cache server.

Consistent Hashing Ring:



Hash key to find server:

- hash("user:123") = 45° → Server B
 - hash("user:456") = 200° → Server C
 - hash("user:789") = 310° → Server A

Same key always goes to same server!

No duplication, no inconsistency

Benefits:

- ✓ No duplicate data
 - ✓ Consistent location
 - ✓ Scales horizontally

Drawbacks:

- ✗ Uneven distribution (hotspots)
 - ✗ Adding server moves data

Implementation (from Chapter 7):

javascript

```
const crypto = require('crypto');

class ConsistentHashCache {
  constructor(servers) {
    this.servers = new Map();
    this.ring = [];
    this.virtualNodes = 150; // Per server

    servers.forEach(server => this.addServer(server));
  }

  hash(key) {
    return parseInt(
      crypto.createHash('md5').update(key).digest('hex').substring(0, 8),
      16
    );
  }

  addServer(server) {
    // Add virtual nodes for better distribution
    for (let i = 0; i < this.virtualNodes; i++) {
      const hash = this.hash(`${server.id}:${i}`);
      this.ring.push({ hash, server });
    }

    // Sort ring by hash
    this.ring.sort((a, b) => a.hash - b.hash);

    this.servers.set(server.id, server);
    console.log(`Added server ${server.id} to ring`);
  }

  getServer(key) {
    const keyHash = this.hash(key);

    // Find first server clockwise from key
    for (let i = 0; i < this.ring.length; i++) {
      if (this.ring[i].hash >= keyHash) {
        return this.ring[i].server;
      }
    }

    // Wrap around
    return this.ring[0].server;
  }
}
```

```

async get(key) {
  const server = this.getServer(key);
  console.log(`Key ${key} → Server ${server.id}`);

  // Get from specific server
  return await server.cache.get(key);
}

async set(key, value) {
  const server = this.getServer(key);
  console.log(`Key ${key} → Server ${server.id}`);

  // Set on specific server
  await server.cache.set(key, value);
}

// Create cache cluster
const cacheCluster = new ConsistentHashCache([
  { id: 'cache-1', cache: redisClient1 },
  { id: 'cache-2', cache: redisClient2 },
  { id: 'cache-3', cache: redisClient3 }
]);

// All requests for same key go to same server
await cacheCluster.set('user:123', userData);
await cacheCluster.get('user:123'); // Always cache-1 (example)


// No cache inconsistency!
// No duplicate storage!


```

2. Write-Through and Write-Behind at Scale

Write-Through Caching (Distributed)

Architecture:





Write Flow:

1. Write to cache
2. Write to database (synchronous)
3. Both must succeed

Implementation:

javascript

```
class DistributedWriteThroughCache {
  constructor(redisCluster, database) {
    this.cache = redisCluster;
    this.db = database;
  }

  async get(key) {
    try {
      // Try cache first
      const cached = await this.cache.get(key);

      if (cached) {
        console.log(`Cache HIT: ${key}`);
        return JSON.parse(cached);
      }

      console.log(`Cache MISS: ${key}`);

      // Fetch from database
      const data = await this.db.query(
        'SELECT * FROM users WHERE id = ?',
        [key]
      );

      if (data) {
        // Populate cache
        await this.cache.setex(key, 3600, JSON.stringify(data));
      }
    }

    return data;
  } catch (error) {
    console.error('Cache error, falling back to database:', error);
    // Degrade gracefully: bypass cache
    return await this.db.query('SELECT * FROM users WHERE id = ?', [key]);
  }
}

async set(key, value) {
  try {
    // Write-through: Update BOTH cache and database
    const dbPromise = this.db.query(
      'UPDATE users SET data = ? WHERE id = ?',
      [JSON.stringify(value), key]
    );
  }
}
```

```

const cachePromise = this.cache.setex(
  key,
  3600,
  JSON.stringify(value)
);

// Wait for both (atomic)
await Promise.all([dbPromise, cachePromise]);

console.log(`Write-through complete: ${key}`);
}

} catch (error) {
  console.error('Write-through failed:', error);
}

// Ensure consistency: invalidate cache on error
await this.cache.del(key);

throw error;
}
}
}

// Usage across multiple servers
const cache = new DistributedWriteThroughCache(redisCluster, database);

// Server 1 writes
await cache.set('user:123', { name: 'John Doe', email: 'john@example.com' });


// Server 2 reads (gets updated value from shared cache)
const user = await cache.get('user:123');
console.log(user); // ✓{ name: 'John Doe', ... }

```

Write-Behind Caching (Distributed)

Asynchronous database writes for better performance.

Architecture with Write-Behind:



Write Flow:

1. Write to cache (fast!)
2. Return success to user
3. Queue database write
4. Background worker processes queue

Implementation:

javascript

```
const { Kafka } = require('kafkajs');

class DistributedWriteBehindCache {
  constructor(redisCluster, writeQueue) {
    this.cache = redisCluster;
    this.queue = writeQueue;
  }

  async get(key) {
    // Same as write-through
    const cached = await this.cache.get(key);

    if (cached) {
      return JSON.parse(cached);
    }

    const data = await database.get(key);

    if (data) {
      await this.cache.setex(key, 3600, JSON.stringify(data));
    }

    return data;
  }

  async set(key, value) {
    // Write to cache immediately (fast!)
    await this.cache.setex(key, 3600, JSON.stringify(value));

    console.log(`Write-behind: Cached ${key} (0.5ms)`);

    // Queue database write (async, don't wait)
    await this.queue.send({
      topic: 'cache-writes',
      messages: [
        {
          key: key,
          value: JSON.stringify({
            operation: 'set',
            key: key,
            value: value,
            timestamp: Date.now()
          })
        }
      ]
    });

    console.log(`Write-behind: Queued DB write for ${key}`);
  }
}
```

```
// Return immediately (user doesn't wait for DB!)
return { success: true };
}

}

// Background Worker (processes queue)
class WriteBehindWorker {
  constructor(database) {
    this.db = database;
    this.kafka = new Kafka({ brokers: ['kafka:9092'] });
    this.batchSize = 100;
    this.batchTimeout = 5000; // 5 seconds
  }

  async start() {
    const consumer = this.kafka.consumer({ groupId: 'cache-write-workers' });
    await consumer.connect();
    await consumer.subscribe({ topic: 'cache-writes' });

    let batch = [];
    let lastFlush = Date.now();

    await consumer.run({
      eachMessage: async ({ message }) => {
        const write = JSON.parse(message.value);
        batch.push(write);

        // Flush if batch full or timeout
        if (batch.length >= this.batchSize ||
            Date.now() - lastFlush > this.batchTimeout) {
          await this.flushBatch(batch);
          batch = [];
          lastFlush = Date.now();
        }
      }
    });

    async flushBatch(batch) {
      if (batch.length === 0) return;

      console.log(`Flushing batch of ${batch.length} writes to database...`);

      try {
        // Begin transaction
        await this.db.query('BEGIN');
        
```

```

// Execute all writes in batch
for (const write of batch) {
    await this.db.query(
        'UPDATE users SET data = ? WHERE id = ?',
        [JSON.stringify(write.value), write.key]
    );
}

// Commit
await this.db.query('COMMIT');

console.log(`✓ Batch committed (${batch.length} writes)`);

} catch (error) {
    console.error('Batch write failed:', error);
    await this.db.query('ROLLBACK');

    // Re-queue failed writes
    // In production: send to DLQ
}

}

}

// Start worker
const worker = new WriteBehindWorker(database);
worker.start();

// Usage
const cache = new DistributedWriteBehindCache(redisCluster, kafka.producer());

// Fast writes!
await cache.set('user:123', { name: 'John' }); // Returns in 1ms
await cache.set('user:456', { name: 'Jane' }); // Returns in 1ms

// Database updated in background (batched)
// User doesn't wait!

```

Write-Behind Challenges

Challenges:

1. DATA LOSS RISK

Cache crashes before DB write
→ Data lost!

Solution: Persistent cache (Redis with AOF)

+ Message queue durability

2. WRITE ORDERING

Updates must be applied in order

Solution: Partition queue by key

Same key → same partition → ordered

3. READ-YOUR-WRITES

User writes, immediately reads

Sees old data from database!

Solution: Read from cache after write

4. CONSISTENCY

Cache and DB temporarily inconsistent

Solution: Acceptable for non-critical data only

Read-Your-Writes Implementation:

javascript

```
class WriteBehindWithConsistency {
  constructor(cache, queue, db) {
    this.cache = cache;
    this.queue = queue;
    this.db = db;
    this.recentWrites = new Map(); // Track recent writes per session
  }

  async set(key, value, sessionId) {
    // Write to cache
    await this.cache.setex(key, 3600, JSON.stringify(value));

    // Track write for this session
    if (!this.recentWrites.has(sessionId)) {
      this.recentWrites.set(sessionId, new Set());
    }
    this.recentWrites.get(sessionId).add(key);

    // Queue DB write
    await this.queue.send({ key, value });

    // Cleanup after 60 seconds
    setTimeout(() => {
      this.recentWrites.get(sessionId)?.delete(key);
    }, 60000);
  }

  async get(key, sessionId) {
    // Check if this session recently wrote this key
    if (this.recentWrites.get(sessionId)?._has(key)) {
      // Read from cache (has latest value)
      console.log('Read-your-writes: Using cache');
      const cached = await this.cache.get(key);
      return JSON.parse(cached);
    }

    // Normal flow: try cache, then DB
    const cached = await this.cache.get(key);
    if (cached) {
      return JSON.parse(cached);
    }

    return await this.db.get(key);
  }
}
```

```

// Usage
const cache = new WriteBehindWithConsistency(redis, kafka, db);

// User writes
await cache.set('user:123', { name: 'Updated' }, 'session-abc');

// Immediately reads (same session)
const user = await cache.get('user:123', 'session-abc');
// → Reads from cache (guaranteed to see own write) ✓

// Different user reads
const user2 = await cache.get('user:123', 'session-xyz');
// → May read from DB (might not see update yet)
// → Eventually consistent


```

3. Session Management in Distributed Systems


The Session Problem

Problem: HTTP is stateless

User logs in on Server 1:




Next request hits Server 2:



User thinks: "I just logged in! Why ask again?"

Solution 1: Sticky Sessions (Session Affinity)

Load Balancer pins user to same server:



Pros:

- ✓ Simple
- ✓ Local session (fast)

Cons:

- ✗ Uneven load (if User A very active)
- ✗ If Server 1 crashes, User A loses session
- ✗ Can't scale horizontally easily

Implementation:

javascript

```
// Load balancer with sticky sessions
class StickyLoadBalancer {
    constructor(servers) {
        this.servers = servers;
        this.userServerMap = new Map();
    }

    getServer(userId) {
        // Check if user already mapped
        if (this.userServerMap.has(userId)) {
            const server = this.userServerMap.get(userId);

            // Check if server still healthy
            if (server.healthy) {
                return server;
            }
        }

        // Server unhealthy, remap user
        console.log(`Server ${server.id} unhealthy, remapping user ${userId}`);
        this.userServerMap.delete(userId);
    }

    // Map to least loaded healthy server
    const healthyServers = this.servers.filter(s => s.healthy);
    const server = healthyServers.reduce((min, s) =>
        s.activeUsers < min.activeUsers ? s : min
    );

    this.userServerMap.set(userId, server);
    server.activeUsers++;

    console.log(`Mapped user ${userId} to server ${server.id}`);
    return server;
}


route(request) {
    const userId = request.userId;
    const server = this.getServer(userId);

    // Route to specific server
    return server.handle(request);
}
```

Solution 2: Shared Session Store (Recommended)

Store sessions in Redis (centralized).

All Servers Share Session Store:



Benefits:

- ✓ Any server can handle any request
- ✓ True horizontal scaling
- ✓ Server crash doesn't lose sessions
- ✓ Easy to add/remove servers

Drawbacks:

- ✗ Network call for every session access
- ✗ Redis is critical dependency

Implementation:

javascript

```
const session = require('express-session');
const RedisStore = require('connect-redis')(session);
const redis = require('redis');

const redisClient = redis.createClient({
  host: 'redis-cluster',
  port: 6379
});

// Express session middleware
app.use(session({
  store: new RedisStore({ client: redisClient }),
  secret: 'your-secret-key',
  resave: false,
  saveUninitialized: false,
  cookie: {
    maxAge: 24 * 60 * 60 * 1000, // 24 hours
    secure: true,
    httpOnly: true
  }
}));

// Login endpoint
app.post('/login', async (req, res) => {
  const { username, password } = req.body;

  const user = await authenticate(username, password);

  if (user) {
    // Store in Redis session (shared across all servers)
    req.session.userId = user.id;
    req.session.username = user.username;
    req.session.role = user.role;

    console.log('Session created in Redis:', req.session.id);

    res.json({ success: true });
  } else {
    res.status(401).json({ error: 'Invalid credentials' });
  }
});

// Protected endpoint
app.get('/dashboard', (req, res) => {
  if (!req.session.userId) {
    return res.status(401).json({ error: 'Not logged in' });
  }
});
```

```
}

// Session available on ANY server!
res.json({
  userId: req.session.userId,
  username: req.session.username
});
});

// Logout
app.post('/logout', (req, res) => {
  req.session.destroy((err) => {
    if (err) {
      console.error('Failed to destroy session:', err);
    }
    res.json({ success: true });
  });
});
```

Solution 3: JWT (Stateless Sessions)

No server-side storage needed!

JWT Flow:

1. Login:

User → Server

Server validates

Server creates JWT: {userId: 123, exp: ...}

Server signs: JWT + secret = signed token

Server → User (token)

2. Subsequent Requests:

User → Server (token in header)


Server verifies signature

Server decodes: {userId: 123}

Server processes request

No session storage needed!

Token contains all info!



No Redis needed!

Any server can validate any token!

Implementation:

javascript

```
const jwt = require('jsonwebtoken');

const SECRET_KEY = 'your-secret-key';

// Login (create JWT)
app.post('/login', async (req, res) => {
  const { username, password } = req.body;

  const user = await authenticate(username, password);

  if (user) {
    // Create JWT token
    const token = jwt.sign(
      {
        userId: user.id,
        username: user.username,
        role: user.role
      },
      SECRET_KEY,
      {
        expiresIn: '24h',
        issuer: 'my-app'
      }
    );

    res.json({ token });
  } else {
    res.status(401).json({ error: 'Invalid credentials' });
  }
});

// Authentication middleware
function authenticateJWT(req, res, next) {
  const token = req.headers.authorization?.replace('Bearer ', '');

  if (!token) {
    return res.status(401).json({ error: 'No token provided' });
  }

  try {
    // Verify and decode token
    const decoded = jwt.verify(token, SECRET_KEY);

    req.user = decoded;
    next();
  }
}
```

```

} catch (error) {
  res.status(401).json({ error: 'Invalid token' });
}

// Protected endpoint
app.get('/dashboard', authenticateJWT, (req, res) => {
  // User info from JWT (no database/cache lookup needed!)
  res.json({
    userId: req.user.userId,
    username: req.user.username,
    role: req.user.role
  });
});

// No session storage, completely stateless!

```

Session Management Comparison

Method	Storage	Scaling	Security	
Sticky Sessions (Server memory)	Local	Poor	Low	
Shared Store (Redis)	Redis	Excellent	Good	
JWT (Stateless)	None	Excellent	Medium	

Recommendation:


- Small apps: Sticky sessions (simple)
- Medium apps: Redis sessions (scalable)
- APIs: JWT (stateless, microservices-friendly)
- Hybrid: Redis + JWT (best of both)

4. Data Locality and Hotspot Management

Data Locality


Concept: Keep data close to computation.

Poor Locality:




Every cache access: 100ms latency!

Good Locality:




↓ Sync with




100x faster!

Multi-Tier Caching

Caching Hierarchy (Bring Data Closer):





Lookup Flow:

1. Check L1 (local) → 90% hit rate, 0.1ms
2. Check L2 (Redis) → 9% hit rate, 1ms
3. Check L3 (DB) → 1% hit rate, 10ms

$$\begin{aligned}\text{Average latency: } & (0.9 \times 0.1) + (0.09 \times 1) + (0.01 \times 10) \\ & = 0.09 + 0.09 + 0.1 = 0.28\text{ms}\end{aligned}$$

vs Redis only: 1ms

vs Database only: 10ms

35x faster than Redis, 35x faster than DB!

Implementation:

```
javascript
```

```
const NodeCache = require('node-cache');
const redis = require('redis');

class MultiTierCache {
  constructor() {
    // L1: Local in-memory cache (per server instance)
    this.l1Cache = new NodeCache({
      stdTTL: 300, // 5 minutes
      checkperiod: 60,
      maxKeys: 10000 // Limit memory usage
    });

    // L2: Shared Redis cluster
    this.l2Cache = redis.createClient({
      host: 'redis-cluster',
      port: 6379
    });

    // Statistics
    this.stats = {
      l1Hits: 0,
      l2Hits: 0,
      dbHits: 0,
      totalRequests: 0
    };
  }

  async get(key) {
    this.stats.totalRequests++;

    // L1: Check local cache
    const l1Value = this.l1Cache.get(key);
    if (l1Value !== undefined) {
      this.stats.l1Hits++;
      console.log(`L1 HIT: ${key} (0.1ms}`);
      return l1Value;
    }

    // L2: Check Redis
    const l2Value = await this.l2Cache.get(key);
    if (l2Value) {
      this.stats.l2Hits++;
      console.log(`L2 HIT: ${key} (1ms)`);

      const parsed = JSON.parse(l2Value);
    }
  }
}
```

```
// Promote to L1
this.l1Cache.set(key, parsed);

return parsed;
}

// L3: Database
this.stats.dbHits++;
console.log(`DB HIT: ${key} (10ms)`);

const dbValue = await database.get(key);

if (dbValue) {
    // Populate both caches
    this.l1Cache.set(key, dbValue);
    await this.l2Cache.setex(key, 3600, JSON.stringify(dbValue));
}

return dbValue;
}

async set(key, value) {
    // Write to database
    await database.set(key, value);

    // Update L2 cache
    await this.l2Cache.setex(key, 3600, JSON.stringify(value));

    // Update L1 cache
    this.l1Cache.set(key, value);

    console.log(`Updated all cache tiers: ${key}`);
}

async invalidate(key) {
    // Invalidate all tiers
    this.l1Cache.del(key);
    await this.l2Cache.del(key);

    console.log(`Invalidated: ${key}`);
}

getStats() {
    const total = this.stats.totalRequests;

    return {
        totalRequests: total,
```

```

l1HitRate: ((this.stats.l1Hits / total) * 100).toFixed(2) + '%',
l2HitRate: ((this.stats.l2Hits / total) * 100).toFixed(2) + '%',
dbHitRate: ((this.stats.dbHits / total) * 100).toFixed(2) + '%',
avgLatency: (
  (this.stats.l1Hits * 0.1) +
  (this.stats.l2Hits * 1) +
  (this.stats.dbHits * 10)
) / total
};

}
}

```

// Usage

```

const cache = new MultiTierCache();

// Simulate traffic
for (let i = 0; i < 1000; i++) {
  // 80% requests for popular items (cache hits)
  const userId = Math.random() < 0.8
    ? Math.floor(Math.random() * 10) // Popular users
    : Math.floor(Math.random() * 1000); // Random users

  await cache.get(`user:${userId}`);
}

```

```
console.log('Cache Statistics:', cache.getStats());
```

// Output:

```

// Cache Statistics: {
//   totalRequests: 1000,
//   l1HitRate: '89.2%', (most requests)
//   l2HitRate: '8.5%', (some misses)
//   dbHitRate: '2.3%', (rare)
//   avgLatency: 0.32ms (very fast!)
// }

```

Hotspot Management


Problem: Popular keys overwhelm single cache server.

Problem: Celebrity Tweet


Celebrity tweets (1M followers):

Tweet ID: tweet-123

Without hotspot handling:



Solution: Hotspot Replication



Hotspot Detection and Mitigation

javascript

```
class HotspotDetector {
  constructor(threshold = 1000) {
    this.accessCounts = new Map();
    this.hotspotThreshold = threshold; // requests/sec
    this.hotKeys = new Set();
    this.window = 1000; // 1 second window
  }

  recordAccess(key) {
    const now = Date.now();

    if (!this.accessCounts.has(key)) {
      this.accessCounts.set(key, []);
    }

    const timestamps = this.accessCounts.get(key);
    timestamps.push(now);

    // Remove old timestamps (outside window)
    const cutoff = now - this.window;
    const filtered = timestamps.filter(ts => ts > cutoff);
    this.accessCounts.set(key, filtered);

    // Check if hotspot
    if (filtered.length > this.hotspotThreshold) {
      if (!this.hotKeys.has(key)) {
        console.log(`🔥 HOTSPOT DETECTED: ${key} (${filtered.length} req/s}`);
        this.hotKeys.add(key);
        this.emit('hotspot', key);
      }
    } else if (filtered.length < this.hotspotThreshold * 0.5) {
      // Cool down
      if (this.hotKeys.has(key)) {
        console.log(`❄️ Hotspot cooled: ${key}`);
        this.hotKeys.delete(key);
        this.emit('cooldown', key);
      }
    }
  }

  isHotspot(key) {
    return this.hotKeys.has(key);
  }
}

class CacheWithHotspotHandling {
```

```
constructor(cacheServers) {
  this.cacheServers = cacheServers;
  this.hotspotDetector = new HotspotDetector(1000);
  this.replicatedKeys = new Set();

  // Listen for hotspot events
  this.hotspotDetector.on('hotspot', (key) => {
    this.replicateKey(key);
  });

  this.hotspotDetector.on('cooldown', (key) => {
    this.unreplicateKey(key);
  });
}

async get(key) {
  // Record access for hotspot detection
  this.hotspotDetector.recordAccess(key);

  if (this.replicatedKeys.has(key)) {
    // Hotspot: Load balance across all servers
    const server = this.cacheServers[
      Math.floor(Math.random() * this.cacheServers.length)
    ];

    console.log(`Hotspot ${key} → Random server ${server.id}`);
    return await server.get(key);
  } else {
    // Normal: Consistent hashing
    const server = this.getServerForKey(key);
    console.log(`Normal ${key} → Server ${server.id}`);
    return await server.get(key);
  }
}

async replicateKey(key) {
  console.log(`Replicating hotspot key: ${key}`);

  // Get value from primary server
  const primaryServer = this.getServerForKey(key);
  const value = await primaryServer.get(key);

  // Replicate to all servers
  for (const server of this.cacheServers) {
    if (server !== primaryServer) {
      await server.set(key, value);
    }
  }
}
```

```

    }

}

this.replicatedKeys.add(key);
console.log(`✓ Key ${key} replicated to all ${this.cacheServers.length} servers`);
}

async unreplicateKey(key) {
  console.log(`Removing replication for cooled key: ${key}`);

  // Remove from all servers except primary
  const primaryServer = this.getServerForKey(key);

  for (const server of this.cacheServers) {
    if (server !== primaryServer) {
      await server.delete(key);
    }
  }

  this.replicatedKeys.delete(key);
}

getServerForKey(key) {
  // Consistent hashing (simplified)
  const hash = this.hash(key);
  const index = hash % this.cacheServers.length;
  return this.cacheServers[index];
}

hash(key) {
  let hash = 0;
  for (let i = 0; i < key.length; i++) {
    hash = ((hash << 5) - hash) + key.charCodeAt(i);
    hash = hash & hash;
  }
  return Math.abs(hash);
}

// Simulation
const cache = new CacheWithHotspotHandling([
  { id: 'cache-1', get: async (k) => `value-${k}` , set: async (k,v) => {}, delete: async (k) => {} },
  { id: 'cache-2', get: async (k) => `value-${k}` , set: async (k,v) => {}, delete: async (k) => {} },
  { id: 'cache-3', get: async (k) => `value-${k}` , set: async (k,v) => {}, delete: async (k) => {} }
]);

// Simulate viral content

```

```
console.log('Simulating viral tweet...');


for (let i = 0; i < 2000; i++) {
    await cache.get('tweet:viral-123'); // Same key, high volume
}

// Output:
// Normal tweet:viral-123 → Server cache-2 (first 1000)
// 🔥 HOTSPOT DETECTED: tweet:viral-123 (1000 req/s)
// Replicating hotspot key: tweet:viral-123
// ✓ Key tweet:viral-123 replicated to all 3 servers
// Hotspot tweet:viral-123 → Random server cache-1 (load balanced!)
// Hotspot tweet:viral-123 → Random server cache-3 (load balanced!)
```

Geographic Data Locality

Keep data close to users geographically.

Global Architecture:



Lookup:

1. US user → Check US cache (5ms)
2. If miss → Check global cache (50ms)
3. If miss → Database (100ms)

EU user → Check EU cache (5ms)

Asia user → Check Asia cache (5ms)

Benefit: 20x lower latency than cross-region

Implementation:

javascript

```
class GeographicCache {
  constructor(region) {
    this.region = region;

    // Local region cache (fast)
    this.localCache = redis.createClient({
      host: `cache-${region}`,
      port: 6379
    });

    // Global cache (slower, fallback)
    this.globalCache = redis.createClient({
      host: 'cache-global',
      port: 6379
    });
  }

  async get(key) {
    const start = Date.now();

    try {
      // Try local cache first
      const local = await this.localCache.get(key);
      if (local) {
        console.log(`Local cache HIT (${this.region}): ${Date.now() - start}ms`);
        return JSON.parse(local);
      }
    }

    // Try global cache
    const global = await this.globalCache.get(key);
    if (global) {
      console.log(`Global cache HIT: ${Date.now() - start}ms`);

      const parsed = JSON.parse(global);

      // Populate local cache
      await this.localCache.setex(key, 3600, global);

      return parsed;
    }

    // Database fallback
    const data = await database.get(key);
    console.log(`Database HIT: ${Date.now() - start}ms`);

    if (data) {
```

```
// Populate both caches
const serialized = JSON.stringify(data);
await this.localCache.setex(key, 3600, serialized);
await this.globalCache.setex(key, 3600, serialized);
}

return data;

} catch (error) {
  console.error('Cache error:', error);
  return await database.get(key);
}
}

async set(key, value) {
// Write to database
await database.set(key, value);

const serialized = JSON.stringify(value);

// Update local cache
await this.localCache.setex(key, 3600, serialized);

// Update global cache (async, don't wait)
this.globalCache.setex(key, 3600, serialized).catch(err => {
  console.error('Failed to update global cache:', err);
});

}

async invalidate(key) {
// Invalidate local
await this.localCache.del(key);

// Invalidate global
await this.globalCache.del(key);

// Optionally: Broadcast to other regions
await this.broadcastInvalidation(key);
}

async broadcastInvalidation(key) {
// Publish invalidation message
await pubsub.publish('cache-validation', JSON.stringify({
key,
region: this.region,
timestamp: Date.now()
}));
```

```

}

}

// Different regions
const usCache = new GeographicCache('us-east');
const euCache = new GeographicCache('eu-west');
const apCache = new GeographicCache('ap-southeast');

// US user
const user1 = await usCache.get('user:123'); // 5ms (local)

// EU user (same data)
const user2 = await euCache.get('user:123'); // 50ms (global cache)
// But second request: 5ms (now in local cache)

```

Cache Stampede Prevention

Problem: Cache expires, many requests hit database simultaneously.

Cache Stampede Scenario:

Popular item cached (1M requests/hour)

Cache expires at 10:00:00

10:00:00.000 - Request 1 arrives

 Cache miss → Query DB

10:00:00.001 - Request 2 arrives

 Cache miss → Query DB

10:00:00.002 - Request 3 arrives

 Cache miss → Query DB

...

10:00:00.100 - Request 100 arrives

 Cache miss → Query DB

100 simultaneous DB queries!

Database overloaded! 

Solution: Mutex/Lock Pattern

javascript

```
class CacheWithStampedeProtection {
  constructor(cache, db) {
    this.cache = cache;
    this.db = db;
    this.locks = new Map();
  }

  async get(key) {
    // Check cache
    const cached = await this.cache.get(key);
    if (cached) {
      return JSON.parse(cached);
    }

    // Cache miss - acquire lock
    const lock = await this.acquireLock(key);

    if (lock.acquired) {
      try {
        // Double-check cache (another thread might have loaded it)
        const recheck = await this.cache.get(key);
        if (recheck) {
          return JSON.parse(recheck);
        }
      }
    }

    // Load from database (only one thread does this!)
    console.log(`[LOCK HOLDER] Loading ${key} from database...`);
    const data = await this.db.get(key);

    // Populate cache
    await this.cache.setex(key, 3600, JSON.stringify(data));

    return data;
  } finally {
    // Release lock
    await this.releaseLock(key, lock.token);
  }
}

} else {
  // Another thread is loading, wait and retry
  console.log(`[WAITING] Another thread loading ${key}...`);

  // Wait a bit
  await new Promise(resolve => setTimeout(resolve, 100));
}
```

```

// Retry (cache should be populated now)
return await this.get(key);
}

}

async acquireLock(key, timeout = 10000) {
  const lockKey = `lock:${key}`;
  const token = Math.random().toString(36);

  // Try to acquire lock (SET NX - set if not exists)
  const acquired = await this.cache.set(
    lockKey,
    token,
    'PX', timeout, // Expire after timeout
    'NX' // Only if not exists
  );

  return {
    acquired: acquired === 'OK',
    token: token
  };
}

async releaseLock(key, token) {
  const lockKey = `lock:${key}`;

  // Lua script for atomic check-and-delete
  const script = `
if redis.call("get", KEYS[1]) == ARGV[1] then
  return redis.call("del", KEYS[1])
else
  return 0
end
`;

  await this.cache.eval(script, 1, lockKey, token);
}

}

// Usage
const cache = new CacheWithStampedeProtection(redisClient, database);

// Simulate 100 simultaneous requests
const promises = [];
for (let i = 0; i < 100; i++) {
  promises.push(cache.get('popular-item-123'));
}

```

```

await Promise.all(promises);

// Output:
// [LOCK HOLDER] Loading popular-item-123 from database...
// [WAITING] Another thread loading popular-item-123...
// [WAITING] Another thread loading popular-item-123...
// ... (99 threads wait)
// All threads get the value
// Only ONE database query! ✓


```

Real-World Examples

Facebook's TAO (The Associations and Objects)

Facebook's Distributed Cache:

Architecture:



Write Flow:

1. Write to leader cache
2. Write to MySQL
3. Invalidate followers
4. Followers refill from leader

Read Flow:

1. Check follower cache (99% hit rate, 1ms)
2. If miss, check leader cache (1ms)
3. If miss, query MySQL (10ms)

Scale: Billions of requests/second

Twitter's Manhattan

Twitter's Distributed Key-Value Store:

Features:

- Strong consistency within datacenter
- Eventual consistency across datacenters
- Multi-tier caching
- Automatic hotspot detection

Hotspot Handling:

1. Detect hot keys (>10K req/s)
2. Replicate to multiple cache servers
3. Load balance across replicas
4. Automatically de-replicate when cool

Example: Trending tweet

- Normal tweet: 100 req/s → 1 cache server
- Viral tweet: 100K req/s → Replicated to 10 servers
- Each handles 10K req/s ✓

Key Takeaways

1. Distributed Cache Coherence:

- Shared cache (Redis cluster) - simple, consistent
- Invalidation messages - local cache + notifications
- Consistent hashing - partition data
- Multi-tier - balance speed and consistency

2. Write-Through vs Write-Behind:

- Write-through: Sync, slower, consistent
- Write-behind: Async, faster, eventual consistency

- Batching for efficiency
- Choose based on requirements

3. Session Management:

- Sticky sessions - simple, poor scaling
- Shared session store - scalable, Redis dependency
- JWT - stateless, no storage needed
- Hybrid approaches

4. Data Locality:

- Multi-tier caching (L1 local, L2 shared)
- Geographic distribution
- Bring data close to computation
- 10-100x latency improvement

5. Hotspot Management:

- Detect hot keys
- Replicate to multiple servers
- Load balance across replicas
- Mutex for cache stampede
- Auto-scale for viral content

Practice Problems

1. Design a caching layer for Instagram (global users, viral posts). How would you handle hotspots?
2. Calculate: L1 cache (0.1ms, 90% hit), L2 cache (1ms, 9% hit), DB (10ms, 1% hit). Average latency?
3. Design session management for a banking app. Which approach and why?

Congratulations! 

You've completed **19 comprehensive chapters!** This is excellent progress covering all major system design topics.

Would you like to:

1. **Practice complete system design problems** (highly recommended now!)
2. Continue with remaining specialized chapters
3. Deep dive on specific topics

What interests you most?